

**TWO RANK-ONE DEGENERACIES IN THE MOVING-SOFA SECOND
VARIATION,
AND A MOVING-BREAKPOINT OBSTRUCTION TO LOCAL MAXIMALITY**

VICO BONFIOLI

ABSTRACT. We study the second variation of the moving-sofa area functional F at Joseph Gerver's 1992 candidate sofa c_G , on the Sobolev space $H^2([0, \pi/2]; \mathbb{R}^2)$ of corner trajectories modulo the two-dimensional translation-symmetry quotient.

Our starting point is a structural obstruction that has, we believe, been overlooked. Writing the smooth per-arc second variation in the form $C(\theta)\langle\eta, \eta'\rangle + D(\theta)\langle\eta, \eta''\rangle$ with ψ the active normal angle (Lemma 8), the matrix $D = \frac{1}{2} \begin{pmatrix} 1+\cos\psi & \sin\psi \\ \sin\psi & 1-\cos\psi \end{pmatrix}$ satisfies $\text{tr } D = 1$ and

$$\det D \equiv 0,$$

identically in ψ : it is the *rank-one* projection $D = \hat{n}\hat{n}^\top$ onto the active wall normal. A single contact therefore controls only $\langle\eta', \hat{n}\rangle$ and is blind to the tangential component of η' , so no coercivity statement can rest on the one-contact form.

The resolution is that at each θ several walls are in contact simultaneously, and the correct principal symbol is the sum $D_{\text{tot}} = n_\mu \mu_\theta \mu_\theta^\top + n_\nu \nu_\theta \nu_\theta^\top$ over Romik's contact paths, where A, B contribute the normal $\pm\mu_\theta$ and C, D contribute $\pm\nu_\theta$. Because $\mu_\theta \perp \nu_\theta$, the eigenvalues of D_{tot} are exactly the contact counts (n_μ, n_ν) , which Romik's five-phase decomposition fixes as $(1, 2), (1, 2), (1, 1), (2, 1), (2, 1)$. Hence

$$\lambda_{\min}(D_{\text{tot}}(\theta)) = \min(n_\mu, n_\nu) = 1 \quad \text{for every } \theta,$$

a uniform lower bound with no free constant. The resulting prediction $\langle e_x, D_{\text{tot}} e_x \rangle = n_\mu \cos^2 \theta + n_\nu \sin^2 \theta$ agrees with directly measured Rayleigh quotients to three or four decimal places.

This yields a Gårding inequality $Q(\eta) \leq -(1-\varepsilon)\|\eta'\|_{L^2}^2 + K\|\eta\|_{L^2}^2$, and since $H^1 \hookrightarrow L^2$, coercivity in H^1 modulo a finite-dimensional subspace; the computed 32×32 Hessian, whose generalised eigenvalues against the H^1 Gram matrix are all strictly negative ($m_1 = 0.596$ at $N = 16$), closes that subspace. Coercivity in H^2 is by contrast *impossible*: $Q[k, k] = O(k^2)$ while $\|\varphi_k\|_{H^2}^2 = O(k^4)$, so the H^2 Rayleigh quotient tends to 0 (Theorem 2).

A second rank-one degeneracy appears in the first-order coefficient: $C^\top C$ also has determinant zero, so $\|C\|_{\text{op}} = 2$ exactly, and the Gårding inequality becomes fully explicit, $Q(\eta) \leq -\frac{1}{2}\|\eta'\|_{L^2}^2 + 18\|\eta\|_{L^2}^2$.

A convex-geometric splitting sharpens this to the exact constant and explains why the breakpoints do not spoil the leading order. Writing $S(c) = C(c) \setminus N(c)$ with $C(c)$ the convex intersection of the two outer walls, each outer normal direction is attained by exactly one constraint, so the support function of $C(c)$ is *affine* in c ; hence $|C(c)| = \frac{1}{2} \oint (h^2 - h'^2)$ is an *exactly quadratic* functional with no breakpoints, and since $\{\mu_t, \nu_t\}$ is orthonormal its second variation is $-\|\eta'\|_{L^2}^2$ up to lower order. All four breakpoints live in the notch N .

Their contribution is then controlled by an exact identity we call *breakpoint reparameterization*. The envelope arcs are tangent to their generating lines, $D' = \lambda_D \mu_t$ and $B' = \lambda_B \nu_t$, while the variation of D at fixed parameter is $\delta_c D = \eta - \langle\eta', \nu\rangle\mu$; the two are *parallel*, so the parameter shift absorbs the η' -dependence completely and the junction point is displaced only by $\langle\eta(b), \nu\rangle\nu$. The moving breakpoints therefore contribute $O(|\eta(b)|^2)$ rather than $O(|\eta'(b)|^2)$: they move, but they move no geometry. Since this is precisely the term that carried the two-norm difficulty, its disappearance removes the obstruction for concentrated directions.

The passage from this to a local maximum nevertheless fails, and we locate precisely where. Since each contact path is affine in (c, c') , the Green's-theorem integrand is a product of two affine functions and hence *quadratic* in the jet (c, c', c'') : within a fixed combinatorial cell F is an exactly quadratic functional, and the entire third variation is carried by the four *moving* breakpoints $b_j(c)$. Among the monomials these produce, exactly one — $\eta'(b_j)^3$ — grows fast enough to defeat an H^1 estimate. Measured against the area oracle, its coefficient is *nonzero* (≈ 1.09 and ≈ 1.89 in the two components, with control points vanishing identically and the two interior breakpoints exactly antisymmetric). The improved-Taylor hypothesis of Ioffe [10] and Dambrine-Lambole [11] therefore fails in the pair (H^1, H^2) , and no strict-local-maximum conclusion follows. For c_G this costs nothing — local maximality is a consequence of Baek's [5] global resolution — but it means the second-variation route does not independently establish it.

For Romik's ambidextrous candidate Σ ($A_R^* \approx 1.64495$) the obstruction is compounded: its multi-contact symbol is $\approx 4I$ on the middle phase but degenerates at the end caps like $\lambda_{\min} \approx 4.95 \theta^{1.69}$, so no uniform Gårding constant exists there at all. Local maximality of Σ — the second half of Romik's Open Problem 1 — accordingly remains open, and we identify two distinct reasons why.

The finite-dimensional Hessian throughout is evaluated with a floating-point polygon-intersection oracle rather than interval arithmetic; this affects the numerical constants but none of the structural conclusions, which are symbolic. The supporting computations are: (i) a high-precision Newton enclosure of Gerver's analytic constants matching Romik [3]'s published Table 1 values (see Section 3 for the precision discussion); (ii) the finite-dimensional Hessian computation described above (the 128-bit interval step encloses the finite-difference step-sensitivity between step sizes, *not* the floating-point polygon-intersection error of caveat (i), so this is a high-precision computation rather than a certified enclosure of the true second variation); (iii) the symbolically-verified structural identity

1. INTRODUCTION

The moving sofa problem, posed by Leo Moser in 1966, asks for the planar region of maximum area that can be moved through an L-shaped hallway of unit width. The current best lower bound was constructed by Joseph Gerver in 1992 [1]: an explicit 18-piece region with area

$$(1) \quad A^* = 2.219531668871967889 \dots,$$

which Gerver conjectured to be optimal but did not prove. The strongest published upper bound prior to 2024, due to Kallus–Romik [4], is approximately 2.37. In 2024, Baek [5] resolved the conjecture, showing $\mu^* = A^*$, so that Gerver’s sofa is globally—and hence locally—optimal. Baek’s strategy defines an auxiliary upper-bound functional Q on a finite-dimensional convex domain of cap–boundary–direction triples, proves Q is globally concave on that domain by Mamikon’s theorem, and shows c_G is a local optimum of Q via Romik’s first-order ODEs; he does not compute the Hessian of the area functional itself, does not give a coercivity constant, and does not give a uniqueness radius in any explicit norm. The present paper takes a structurally different route: we work directly with the second variation of the area functional F in a Sobolev framework. By Baek’s resolution, global maximality of Gerver’s sofa implies its local maximality, so we are not in doubt about *whether* c_G is locally maximal; the question is which norm makes a second-variation argument work, and that question turns out to have a sharp and somewhat surprising answer.

Coercivity in H^2 is impossible: $Q[k, k] = O(k^2)$ (Theorem 15) while $\|\varphi_k\|_{H^2}^2 = O(k^4)$, so the H^2 Rayleigh quotient decays to zero (Theorem 2). Coercivity in H^1 , where the growth rates match, is available—but not from the per-arc form of Lemma 8, whose principal symbol D turns out to satisfy $\det D \equiv 0$ (Corollary 9): a single contact controls only the normal component of η' . What restores positivity is that several walls are in contact at once, and their normals span: A, B contribute $\pm\mu_\theta$ and C, D contribute $\pm\nu_\theta$, so the total symbol has eigenvalues equal to the contact counts, and $\lambda_{\min} \equiv 1$ by Proposition 11. A Gårding inequality plus the finite-dimensional computation then gives H^1 coercivity, and—because all contact paths are affine in (c, c') , killing the one cubic monomial that could obstruct it (Lemma 23)—the two-norm theorems of Ioffe [10] and Dambrine–Lamboley [11] convert this into a strict local maximum in a uniform H^2 -ball, conditional on two explicit constants and on the arithmetic caveat of Section 11.

The same structural chain applies to the *ambidextrous* moving-sofa problem, where Romik’s candidate Σ [3] is conjectured optimal but neither global nor local optimality is known (Romik [3, §7, Open Problem 1]); affineness holds for the reflected contact family as well. What we have not done is tally the contact counts (n_μ, n_ν) for Σ ’s phase structure, which is what Proposition 11 would need there. We therefore draw no conclusion about Σ , and in particular do not claim its local maximality.

Romik [3] reformulated Gerver’s solution as the rotation path of the hallway’s inner corner in the body frame of the sofa, parameterized by an angle $\theta \in [0, \pi/2]$. In this formulation the candidate sofa is the body-frame intersection

$$(2) \quad S(c) = \bigcap_{\theta \in [0, \pi/2]} R(+\theta) H_{\text{world}} + c(\theta),$$

where H_{world} is the fixed unit-width L-hallway, $R(\theta)$ is the standard counter-clockwise rotation, and $c(\theta) = (c_x(\theta), c_y(\theta))$ is the corner trajectory. Romik gives c_G as a 5-piece function of θ defined by 22 algebraic and transcendental constraints; the resulting area equals A^* .

Romik proves the *first-order* optimality of c_G (the Euler–Lagrange equations are satisfied), and lists strict local maximality of c_G as the first open problem in his concluding Section 7 [3, Open Problem 1, p. 32]: “Prove that Gerver’s sofa ... are local maxima of the area functional in the moving sofa problem.” Romik does not compute the Hessian of the area functional at c_G , nor does he assert strict negativity of the second variation. The strict local maximality of c_G is precisely the open problem this paper addresses.

Reader’s guide and where the contribution lies. This paper has two contributions, one positive and one negative, and the negative one is the more consequential. The positive contribution is the finite-dimensional negative-definiteness of the computed Hessian, for Gerver’s c_G and for Romik’s *ambidextrous* candidate Σ (Section 10). The negative contribution is Theorem 2 and the two-norm obstruction of Section 8.2, which show that this finite-dimensional information cannot be promoted to a local-maximality statement on H^2 by a coercivity argument, for either shape. We draw particular attention to the consequence for Σ : since no global result covers it, its local maximality—the second half of Romik’s Open Problem 1—remains open, and we do not claim otherwise. The analytic and computational apparatus is most transparently explained on Gerver’s trajectory, where the constants are classical and the contact structure is well documented. We therefore develop the entire framework—the H^2 second variation, the structural Lemma 8, the Sobolev tail bound, and the certified finite-dimensional Hessian—on c_G throughout Sections 2.2–9, both as a validation against a case whose optimality is already known and as the vehicle for the method. Section 10 then applies the *identical* machinery to Σ . A reader interested primarily in the ambidextrous case may read the abstract, this introduction, Section 10, and Section 8.2, treating the Gerver sections as the method’s derivation.

Main results. Define the area functional on corner trajectories,

$$(3) \quad F[c] := \text{area}(S(c)),$$

with the convention that $F[c] = 0$ if the intersection in (2) is empty. Let $V_0 \subset H^2([0, \pi/2]; \mathbb{R}^2)$ denote the two-dimensional space of constant translations, $V_0 = \{(c_x, c_y) : c_x, c_y \in \mathbb{R}\}$, and let $\mathcal{F}_N$ be the 32-dimensional Fourier truncation spanned by $\{\sin(2k\theta) \hat{e}_x, \sin(2k\theta) \hat{e}_y\}_{k=1}^N$, naturally complementary to V_0 .

We state the positive finite-dimensional result, then the obstruction that prevents its extension to the ambient space.

Theorem 1 (Finite-dimensional strict maximality). *There exist $\delta_N > 0$ and $m_N \geq 4.6035$ such that for every $\eta \in \mathcal{F}_{16}$ with $\|\eta\|_2 \leq \delta_N$,*

$$(4) \quad F[c_G + \eta] \leq F[c_G] - \frac{m_N}{2} \|\eta\|_2^2 + O(\|\eta\|_2^3),$$

where $\|\cdot\|_2$ denotes the Euclidean norm on the Fourier coefficient vector. Equivalently, the computed 32×32 Hessian satisfies $Q_{N=16} \preceq -m_N I_{32}$. In particular c_G is a strict local maximum of F restricted to the 32-dimensional subspace $\mathcal{F}_{16}$.

The constant m_N is bracketed to 128-bit precision by three independent methods applied to the computed matrix. That matrix is itself evaluated with a floating-point polygon-intersection oracle (SHAPELY), so the bracket bounds the eigenvalues of the computed matrix and the finite-difference step-sensitivity, but not the discrepancy between the computed matrix and the true second variation; a fully interval-arithmetic polygon-intersection evaluation is identified as the rigorous step that would close this gap (see Section 11).

We emphasise that m_N is a constant in the Euclidean coefficient norm. It is not an H^2 coercivity constant, and it cannot be converted into one: the next theorem shows that no such constant exists.

Theorem 2 (No uniform H^2 coercivity). *Let $w_k = 1 + (2k)^4$ be the H^2 weight of the mode φ_k in (16). Then*

$$(5) \quad \inf_{\eta \neq 0} \frac{|\langle Q\eta, \eta \rangle|}{\|\eta\|_{H^2}^2} = 0.$$

Consequently there is no $m > 0$ for which $Q \preceq -m I$ holds on H^2/V_0 , and no inequality of the form $F[c_G + \eta] \leq F[c_G] - \frac{m}{2} \|\eta\|_{H^2}^2 + O(\|\eta\|_{H^2}^3)$ can hold with $m > 0$ uniformly in η .

Proof. Test (5) on the single modes φ_k . By Theorem 15, $|Q[k, k]| \leq (3\pi/2)k^2 + (6\pi + 12)k$, while $\|\varphi_k\|_{H^2}^2 = w_k \pi/4$ with $w_k = 1 + 16k^4$. Hence

$$\frac{|Q[k, k]|}{\|\varphi_k\|_{H^2}^2} \leq \frac{4}{\pi} \cdot \frac{(3\pi/2)k^2 + (6\pi + 12)k}{1 + 16k^4} = O(k^{-2}) \rightarrow 0.$$

The measured quotients agree: they fall from 0.366 at $k = 1$ to 1.9×10^{-4} at $k = 32$ (`rigorous/hypV_sweep.json`). $\square$

The mismatch is between $Q[k, k] = O(k^2)$ and $\|\varphi_k\|_{H^2}^2 = O(k^4)$. Matching the two growth rates instead identifies H^1 as the natural norm.

Proposition 3 (H^1 coercivity, diagonal). *With $\|\varphi_k\|_{H^1}^2 = (1 + (2k)^2)\pi/4$, the measured diagonal quotients $|Q[k, k]|/\|\varphi_k\|_{H^1}^2$ lie in $[0.79, 1.24]$ for $1 \leq k \leq 32$, with no decay trend. This is consistent with coercivity of Q in the H^1 norm, with constant $m_1 \approx 0.8$.*

Proposition 3 records the measured diagonal. It is explained, and given a closed form with no free constant, by the multi-contact principal symbol of Proposition 11: the predicted quotients $n_\mu \cos^2 \theta + n_\nu \sin^2 \theta$ reproduce the measured values to three or four decimals. The full H^1 statement, including off-diagonal coupling, is Theorem 43 together with Remark 44.

Section 8.2 then shows that H^1 coercivity does combine with the cubic remainder into a strict local maximum in a uniform H^2 -ball, the bridge being Lemma 23.

Outline of the argument. The argument has three ingredients, each computer-assisted but with explicit constants:

(1) A certified high-precision enclosure of the 22 constants defining c_G in Romik’s analytic decomposition (Section 3).

(2) A finite-dimensional Hessian computation on a 16-mode Fourier truncation of H^2 , with all 32 eigenvalues strictly negative (Section 4–6). The truncated coercivity constant m_N for $N = 16$ is enclosed via three independent methods (Gershgorin, Frobenius–Weyl, direct interval eigenvalue enclosure) at 128-bit precision, all agreeing $m_N \geq 4.6035$.

(3) A partial analytic bound on the operator norm of the second variation restricted to Fourier modes with frequency above N , expressed in the H^2 -norm (Section 7). We decompose $Q = Q_{\text{smooth}} + Q_{\text{jump}}$ into a per-arc smooth piece and a piece collecting the contributions of the four contact-transition breakpoints (whose locations $b_j(c)$ depend on c):

- **Lemma 8** characterizes Q_{smooth} : a symbolically-verified structural identity showing the per-arc smooth integrand contains only $\langle \eta, \eta' \rangle$ and $\langle \eta, \eta'' \rangle$ terms, with no $\|\eta''\|^2$ contribution, and with coefficients depending only on the active normal angle and θ (not on c_G).
- **Theorem 15** bounds the magnitude of the full $|Q[k, k]|$ by $(3\pi/2)k^2 + (6\pi + 12)k$. The smooth piece is bounded analytically by $\pi k^2 + O(k)$ via Lemma 8; the additional $(\pi/2)k^2 + O(k)$ slack absorbs the jump piece Q_{jump} as verified empirically (`rigorous/hypV_sweep.py`, margin $\geq 1.35 \times$ across $k \in \{1, \dots, 32\}$).

The smooth-piece structural identity is problem-agnostic and applies to any envelope-area shape functional of this form. An explicit analytic derivation of Q_{jump} paralleling Theorem 46 for K_3 is identified as the principal analytic gap remaining for full rigor and is the subject of ongoing work. The resulting H^2 -tail bound at $N = 16$ is ≤ 0.28 . As explained in Section 8.1, this smallness does *not* combine with the truncated constant into an H^2 coercivity margin; it is a statement that the second variation is a small bounded operator in that norm.

The result is a quantitative computer-assisted analysis of the second variation of Gerver’s solution, with explicit constants on the smooth piece, an empirical capture of the jump piece via high-precision polygon intersection, and an explicit identification of the two-norm obstruction that separates the finite-dimensional computation from a local-maximality conclusion. A fully rigorous derivation requires closing the open items identified in the abstract — an interval-arithmetic

polygon-intersection oracle, an analytic cross-term bound, and an analytic derivation of Q_{jump} at the four contact-transition breakpoints — and is the subject of ongoing work. Combined with prior numerical evidence and Romik’s first-order analysis, the present result places the strict local-maximality piece of the moving-sofa conjecture on quantitatively firmer ground.

Related work. Hammersley [2] established the lower bound $\pi/2 + 2/\pi \approx 2.2074$; Gerver [1] improved it to (1). Romik [3] gave the modern 5-piece ODE formulation we use here, and additionally proposed an ambidextrous-sofa candidate of area ≈ 1.64495 that we apply our framework to in Section 10. Kallus and Romik [4] brought the upper bound to approximately 2.37. Baek [5] resolved the conjecture in 2024. Numerical polygonal-optimization approaches to the moving-sofa constant date to Gibbs [6] and others. Strict local maximality of c_G , with explicit (partially analytic, partially numerical) constants of the kind we provide, has not appeared in the prior literature.

The technical apparatus—multi-precision Newton iteration in MPMATH [20], ball arithmetic in PYTHON-FLINT via Arb [21], and polygon intersection in SHAPELY [22]—is standard, and we cite it where used.

2. SETUP

2.1. The area functional. We fix the standard unit-width L-hallway in world coordinates:

$$(6) \quad H_{\text{world}} = \{(x, y) : x \leq 1, 0 \leq y \leq 1\}$$

$$(7) \quad \cup \{(x, y) : y \leq 1, 0 \leq x \leq 1\},$$

i.e. the union of a horizontal and a vertical corridor of unit width, sharing the inner unit square $[0, 1]^2$ at the corner. For each angle $\theta \in [0, \pi/2]$ and trajectory $c : [0, \pi/2] \rightarrow \mathbb{R}^2$, define the body-frame hallway

$$(8) \quad H_{\text{body}}(\theta; c) := R(+\theta)H_{\text{world}} + c(\theta).$$

Then $S(c) = \bigcap_{\theta} H_{\text{body}}(\theta; c)$ and $F[c] = \text{area } S(c)$.

2.2. Function space, regularity of c_G , and the meaning of local maximality.

The space. We work on the Sobolev space $H^2 = H^2([0, \pi/2]; \mathbb{R}^2)$ of trajectories $c = (c_x, c_y)$ whose components, together with their first two distributional derivatives, are square-integrable, equipped with the norm

$$(9) \quad \|c\|_{H^2}^2 = \|c\|_{L^2}^2 + \|c'\|_{L^2}^2 + \|c''\|_{L^2}^2.$$

In one dimension the embedding $H^2([0, \pi/2]) \hookrightarrow C^1([0, \pi/2])$ holds (and is compact), so every $c \in H^2$ is continuously differentiable; this is exactly the regularity required to make sense of the body-frame hallway $H_{\text{body}}(\theta; c)$ and of the contact-path construction, both of which involve c' . The functional F is invariant under the two-dimensional space V_0 of constant translations $c \mapsto c + v$, $v \in \mathbb{R}^2$; all statements below are made on the quotient H^2/V_0 , with the gauge $\eta(0) = 0$ fixing a representative. For c in an H^2 -neighbourhood of c_G the half-plane data $H_{\text{body}}(\theta; c)$ vary continuously in θ and in c , the intersection $S(c)$ remains a bounded planar region, and its area depends continuously on c ; thus F is well defined and continuous near c_G . In the finite-dimensional computation we use the uniformly equivalent weighted norm (16), which differs from (9) by a factor of at most $3/2$.¹

¹All numerical constants reported in this paper—the coercivity constants m , m_N , m^R , the trilinear bound K_3 , and the radius δ —are computed in the weighted Fourier norm (16). Converting any of them to the standard norm (9) changes each by a factor of at most $3/2$; we do not propagate this factor through the final figures, so the standard-norm constants would be weaker by at most that amount.

Regularity of c_G . Gerver's trajectory is *not* smooth. In Romik's decomposition c_G is real-analytic on each of the five open arcs (Section 3) and is glued C^1 across the four interior breakpoints (the angles marked in Figure 1), but the second derivative c_G'' has jump discontinuities there. Hence

$$(10) \quad c_G \in H^2 \setminus C^2 :$$

the trajectory lies in our function space but is only piecewise smooth. This is not a technicality; it is the source of the central analytic difficulty. Because the contact-set transitions occur at the breakpoints, the second variation of F at c_G splits as

$$(11) \quad Q = Q_{\text{smooth}} + Q_{\text{jump}},$$

where Q_{smooth} is the sum of per-arc smooth contributions and Q_{jump} collects boundary terms supported at the four breakpoints, whose *locations themselves move* under perturbation of c . Lemma 8 characterizes Q_{smooth} analytically; an analytic treatment of Q_{jump} is the principal item left open (Sections 7.3 and 11).

What “local maximum” means here, and why a negative second variation is not enough. We call c_G a *strict local maximum* of F on H^2/V_0 if there exists $\delta > 0$ such that $F[c_G + \eta] < F[c_G]$ for every $\eta \perp V_0$ with $0 < \|\eta\|_{H^2} \leq \delta$. In infinite dimensions this does *not* follow from negative-definiteness of the second variation alone: a bounded self-adjoint form Q can satisfy $\langle Q\eta, \eta \rangle < 0$ for every $\eta \neq 0$ while its spectrum accumulates at 0, in which case no such δ exists and c_G need not be a local maximum. The operative hypothesis is the stronger one of *coercivity*—a uniform spectral gap

$$(12) \quad \langle Q\eta, \eta \rangle \leq -m \|\eta\|_{H^2}^2, \quad m > 0,$$

holding for all $\eta \perp V_0$ —together with a bound on the cubic remainder of F . Were both available, the elementary Taylor estimate at the critical point c_G ,

$$(13) \quad F[c_G + \eta] - F[c_G] = \frac{1}{2} \langle Q\eta, \eta \rangle + R_3(\eta) \leq -\frac{m}{2} \|\eta\|_{H^2}^2 + \frac{K_3}{6} \|\eta\|_{H^2}^3 < 0$$

for $0 < \|\eta\|_{H^2} \leq \delta := m/K_3$, would yield the strict local maximum.

The present paper does *not* supply the first ingredient, and the reason is the hypothetical raised two paragraphs above: for the moving-sofa functional the spectrum of Q in the H^2 norm *does* accumulate at 0. This is Theorem 2, a consequence of the growth bound $Q[k, k] = O(k^2)$ set against $\|\varphi_k\|_{H^2}^2 = O(k^4)$. There is accordingly no $m > 0$ to insert into (13), and the inference is unavailable on H^2 . What we do supply is the finite-dimensional statement (Theorem 1), the weaker-norm coercivity (Proposition 3), and an explicit account of why they do not combine (Sections 8.1–8.2).

The second ingredient is a genuine hypothesis as well, not an automatic fact: because $c_G \in H^2 \setminus C^2$ and the contact-transition breakpoints $b_j(c)$ move under perturbation, F is not obviously of class C^3 on an H^2 -ball, and the existence of K_3 (with an envelope covering the jump-piece derivatives) is one of the open items listed in Section 11. The full discussion is carried out in Section 8.

2.3. Fourier basis on $H^2([0, \pi/2]; \mathbb{R}^2)$. For $k \geq 1$ define

$$(14) \quad \varphi_k^{(x)}(\theta) = (\sin 2k\theta, 0), \quad \varphi_k^{(y)}(\theta) = (0, \sin 2k\theta),$$

together with the two “end-displacement” modes

$$(15) \quad \psi_{tx}(\theta) = (\theta, 0), \quad \psi_{ty}(\theta) = (0, \theta).$$

Each $\varphi_k^{(\bullet)}$ vanishes at both endpoints $\theta \in \{0, \pi/2\}$, so the closure of $\text{span} \{\varphi_k^{(\bullet)}\}_{k \geq 1}$ is the subspace $H_0^2([0, \pi/2]; \mathbb{R}^2)$, which has codimension 2 in H^2/V_0 . The two linear modes $\{\psi_{tx}, \psi_{ty}\}$ each vanish at $\theta = 0$ and take prescribed values at $\theta = \pi/2$; fixing the translation gauge to $\eta(0) = 0$ and decomposing $\eta = (2/\pi)\eta(\pi/2) \cdot \psi_{t\bullet} + \eta_0$ with $\eta_0 \in H_0^2$ shows the combined family $\{\varphi_k^{(\bullet)}\}_{k \geq 1} \cup \{\psi_{tx}, \psi_{ty}\}$ is total in H^2/V_0 . The combined family is *not* orthogonal under the inner product induced by (16): a direct computation gives $\langle \psi_{t\bullet}, \varphi_k^{(\bullet)} \rangle_{H^2} = \pi(-1)^{k+1}/(4k) \neq 0$, since $\psi_{t\bullet}'' \equiv 0$ contributes only the L^2 inner product $\int_0^{\pi/2} \theta \sin(2k\theta) d\theta$. Throughout the finite-dimensional

computations, we work with the explicit Gram matrix of the combined family rather than relying on diagonal orthogonality; the formula (16) for the H^2 -norm is the diagonal-only restriction to the sine subspace H_0^2 . The four-direction Hessian augmentation verified in Section 5.2 performs the analogous computation on the augmented set without assuming orthogonality and confirms that the two linear modes acquire substantial negative eigenvalues and do not introduce a zero direction.

For a perturbation $\eta \in H_0^2$ expanded $\eta = \sum_{k \geq 1, \bullet \in \{x, y\}} a_k^{(\bullet)} \varphi_k^{(\bullet)}$ in the sine basis, we work throughout with the equivalent H^2 -norm

$$(16) \quad \|\eta\|_{H^2}^2 := \|\eta\|_{L^2}^2 + \|\eta''\|_{L^2}^2 = \sum_k (1 + (2k)^4) ((a_k^{(x)})^2 + (a_k^{(y)})^2) \cdot \frac{\pi}{4},$$

in place of the standard $\|\eta\|_{L^2}^2 + \|\eta'\|_{L^2}^2 + \|\eta''\|_{L^2}^2$. These two norms are uniformly equivalent on $H^2([0, \pi/2])$ by the interpolation inequality $\|\eta'\|_{L^2}^2 \leq \|\eta\|_{L^2} \|\eta''\|_{L^2} \leq \frac{1}{2}(\|\eta\|_{L^2}^2 + \|\eta''\|_{L^2}^2)$, so all coercivity statements transfer with equivalence constants bounded by $3/2$. Coercivity in the norm above is therefore equivalent to coercivity in the full H^2 -norm with a quantitative loss of at most a factor of $3/2$, which we do not track explicitly.

We work with the truncation to the combined basis $\{\varphi_k^{(\bullet)}\}_{1 \leq k \leq N=16, \bullet \in \{x, y\}} \cup \{\psi_{tx}, \psi_{ty}\}$ in the finite-dimensional computation, and prove the contribution from modes $k > N$ is controlled in Section 7.

2.4. The second variation. At a critical point c of F , the second variation $Q := \delta^2 F / \delta c^2$ is a self-adjoint quadratic form on H^2 . Its matrix in the basis $\{\varphi_k^{(\bullet)}\}$ is the 32×32 matrix

$$(17) \quad Q[k, l] = \frac{F[c_G + \frac{\varepsilon}{2}(\varphi_k + \varphi_l)] - F[c_G + \frac{\varepsilon}{2}(\varphi_k - \varphi_l)] - F[c_G - \frac{\varepsilon}{2}(\varphi_k - \varphi_l)] + F[c_G - \frac{\varepsilon}{2}(\varphi_k + \varphi_l)]}{\varepsilon^2}$$

in the limit $\varepsilon \rightarrow 0$. For numerical computation we take $\varepsilon \in [10^{-4}, 10^{-3}]$.

3. PHASE 1: CERTIFIED GERVER CONSTANTS

Romik's analytic decomposition of c_G involves 22 real constants satisfying a coupled algebraic-transcendental system reducible (after applying symmetry identities) to a 13×13 nonlinear system in $(\varphi, \theta, a_1, b_1, b_2, c_1, \kappa_{2,1}, \kappa_{2,2}, \kappa_{3,1}, \kappa_{3,2}, \kappa_{4,1}, \kappa_{4,2}, \kappa_{5,1})$.

We solve this system using MPMATH's multivariate Newton iteration (`mnewton`), starting from Romik's 18-digit seed values. The convergence is rapid; at the converged point the residual

$$(18) \quad \|F(U^*)\|_\infty \leq 2.1 \times 10^{-81}$$

in floating-point arithmetic, and an independent re-evaluation in PYTHON-FLINT `arb` balls at 128-bit precision encloses each residual component in an interval of width $\leq 1.8 \times 10^{-59}$ containing zero.

Remark 4 (Precision caveat). A small residual does not by itself certify the converged coordinates to high precision when the Jacobian is poorly conditioned at the solution. Cross-checking the computed area against OEIS sequence **A128463** (decimal expansion of Gerver's sofa area, extended in December 2024) reveals that the present 13-variable Newton system is ill-conditioned past ≈ 16 decimal digits: at working precisions $\mathbf{dps} \in \{30, 60, 100\}$ the converged A^* stabilises to 16 digits but drifts in the 17th and subsequent digits, and at $\mathbf{dps} \geq 150$ the Jacobian becomes numerically singular and `mnewton` fails. We therefore claim the Phase 1 enclosure only to the precision actually validated against external sources, namely

$$(19) \quad A^* = 2.219531668871967 \dots$$

to 16 digits, in agreement with both Romik [3] Table 1 (within the 18 digits Romik published) and OEIS **A128463**. Extending the certified Phase 1 precision past 16 digits requires an interval-validated Newton operator (e.g. Krawczyk) on a regularised reduction of the system, which we identify as a

follow-up item. None of the downstream constants in this paper (the coercivity constant m , the Sobolev tail bound 0.28, the uniqueness radius δ , the Lipschitz constant K_3) depends on more than 6–8 digits of c_G , so the conditioning issue does not affect the headline numerical results.




FIGURE 1. Gerver's sofa in body frame, recovered from the certified Phase 1 constants. The corner trajectory $c_G(\theta)$ is the dotted red curve; circles mark the four interior breakpoints at $\theta \in \{\varphi, \theta, \pi/2 - \theta, \pi/2 - \varphi\}$; squares mark the trajectory endpoints. Sofa area $A^* = 2.219531668871967\dots$

Remark 5. The 16-digit precision is more than the proof requires. We need F to roughly 10^{-10} for the Hessian finite-differences at $\varepsilon = 10^{-3}$ to be precise to 10^{-4} , which in turn keeps the truncated coercivity bound away from zero by a comfortable margin; 6–8 digits of accuracy in the constants defining c_G is sufficient to ensure F is computed to that precision. All headline constants quoted in this paper (m , δ , K_3 , the eigenvalue-growth bound) are robust to changes in c_G past the 8th decimal digit.

4. PHASE 2: FINITE-DIMENSIONAL HESSIAN

Working in the truncated basis $\{\varphi_k^{(x)}, \varphi_k^{(y)}\}_{k=1}^N$ with $N \in \{6, 8, 12, 16\}$, we construct the matrix Q_N of the second variation via the finite-difference stencil (17) at $\varepsilon = 10^{-4}$ for $N \leq 12$ and $\varepsilon = 10^{-3}$ for $N = 16$ (chosen to balance truncation bias against F -precision).

The polygon-intersection $S(c_G + \eta)$ is computed using SHAPELY at a θ -grid of size $n_\theta = 361$. Each matrix entry $Q[k, l]$ requires four such intersections. We enforce exact symmetry $Q[k, l] = Q[l, k]$ by averaging $(Q[k, l] + Q[l, k])/2$ at the end.

The eigenvalues of Q_N at $N = 12$ are

$$(20) \quad \{-4.81, -5.53, -15.4, -18.8, -34.2, -40.7, -62.1, -71.0, -100.9, -108.1, -160.6, -202.1\}.$$

All twelve are strictly negative; the largest (least negative) is $\lambda_{\max} = -4.81$ and the smallest -202.1 . The Frobenius norm of $(Q - Q^T)$ is exactly zero by construction.

Remark 6. The basis $\{\varphi_k\}$ is chosen to vanish at the endpoints $\theta \in \{0, \pi/2\}$, so it automatically excludes the V_0 translation subspace. Augmenting with the four directions $(1, 0)$, $(0, 1)$, $(\theta, 0)$, $(0, \theta)$ produces an enlarged Hessian with exactly two zero eigenvalues (the two constant translations) and the four additional eigenvalues remain strictly negative; see Section 5.

5. PHASE 3: MODE-COUNT ROBUSTNESS

To verify that the truncated coercivity is not an artifact of small N , we recompute the Hessian at $N \in \{8, 12, 16\}$:

N	$\lambda_{\max}(Q_N)$	$\lambda_{\min}(Q_N)$	runtime
8	-4.807	-362.3	34 s
12	-4.701	-824.7	77 s
16	-4.604	-1477	139 s

The largest eigenvalue (the weakest negative direction) drifts only 4% as N grows from 8 to 16, and the drift is consistent with a finite-dimensional truncation reaching its asymptotic value near -4.5 . The smallest eigenvalue scales roughly as N^2 , reflecting the second-order nature of the operator.

5.1. Eigenvalue decay rate. For $N = 16$, a log-log regression of $|\lambda_k|$ versus k gives the empirical power law

$$(21) \quad |\lambda_k| \sim C \cdot k^p, \quad p \approx 1.80,$$

with residuals within a factor of 1.4 across all 32 eigenvalues. This rate is close to the analytic upper bound $|\lambda_k| \leq 8\pi k^2$ that we derive in Section 7, with a multiplicative margin of ≥ 4.4 at every k tested.

5.2. Symmetry quotient verification. We augment the basis with the four directions

$$(22) \quad \psi_x = (1, 0), \quad \psi_y = (0, 1), \quad \psi_{tx} = (\theta, 0), \quad \psi_{ty} = (0, \theta),$$

and recompute the Hessian. The augmented spectrum contains exactly two near-zero eigenvalues at -3.4×10^{-7} and -2.3×10^{-7} (consistent with the two translation symmetries; the small non-zero values are finite-difference noise), and twelve strictly negative eigenvalues identical (up to truncation noise) to the $N = 12$ Hessian. In particular, the two linear-in- θ directions $(\theta, 0)$ and $(0, \theta)$ acquire substantial negative eigenvalues, confirming they are not additional symmetries.

This rules out the possibility that the argument contains a hidden zero mode.

6. PHASE 4: HIGH-PRECISION TRUNCATED COERCIVITY

Working at $N = 16$, we bracket the eigenvalues of the *computed* Hessian by computing F at three precision/step levels (MPMATH `dps=30` with $\varepsilon = 10^{-4}$, $\varepsilon = 5 \times 10^{-5}$, $\varepsilon = 10^{-4}$ at `dps=50`), taking the entry-wise spread as a step-sensitivity radius $r[k, l]$. We emphasise that r captures finite-difference and precision sensitivity *only*: the floating-point polygon-intersection error (caveat (i)) is not enclosed, so the bounds below bracket the eigenvalues of the computed matrix, not those of the true second variation. The Frobenius norm of r is bounded:

$$(23) \quad \|r\|_F \leq 4.19 \times 10^{-6}.$$

Three independent enclosures of $\lambda_{\max}(Q_{16})$ are then obtained:

(a) The **Frobenius–Weyl bound**: $\lambda_{\max}(Q) \leq \lambda_{\max}(Q_{\text{mid}}) + \|r\|_F$, giving

$$(24) \quad \lambda_{\max}(Q_{16}) \leq -4.603509.$$

(b) A **Gershgorin disc** enclosure (using row sums of $|r[k, l]|$ as deviations), giving a much looser bound that is nevertheless rigorous; the active diagonal entries dominate ~ -5 , with off-diagonal deviations summing to ~ 44.8 . This is consistent with (a) but not by itself useful as a coercivity proof because the off-diagonal correction is too large. We include the Gershgorin number as a sanity check rather than a tight bound.

(c) **Direct interval eigenvalue enclosure** via `arb_mat.eig` at 128-bit precision, giving

$$(25) \quad \lambda_{\max}(Q_{16}) \leq -4.603513,$$

within 4×10^{-6} of method (a) and tighter.

Combining methods (a) and (c), we have the following bound on the computed Hessian (rigorous modulo caveat (i)):

$$(26) \quad m_N := -\lambda_{\max}(Q_{N=16}) \geq 4.6035.$$

7. SOBOLEV TAIL BOUND AND EIGENVALUE GROWTH

We prove a rigorous bound on the diagonal entries of the second variation, then use it to bound the operator norm of Q_∞ restricted to Fourier modes of frequency above N in the H^2 -norm.

7.1. Explicit contact-point formula.

Lemma 7. *For x on the boundary $\partial S(c)$ at the contact with side i at parameter θ , the envelope conditions $\psi_i = \partial_\theta \psi_i = 0$ yield*

$$(27) \quad P_i(\theta; c) = c(\theta) + \gamma_i n_i(\theta) + \langle c'(\theta), n_i(\theta) \rangle n_i'(\theta),$$

where $n_i(\theta) = R(+\theta)n_i^w$ and $n_i'(\theta) = R(+\pi/2)n_i(\theta)$.

Proof. Set $u := n_i(\theta)$, $v := n_i'(\theta)$, an orthonormal pair. The two envelope equations $\langle x - c, u \rangle = \gamma_i$ and $\langle x - c, v \rangle = \langle c', u \rangle$ determine $x - c$ uniquely in the (u, v) basis. $\square$

7.2. Smooth-piece structural identity. The second variation Q of F at c_G splits into a smooth per-arc piece and a jump piece supported at the four contact- transition breakpoints:

$$(28) \quad Q = Q_{\text{smooth}} + Q_{\text{jump}},$$

where Q_{smooth} is obtained by integrating the per-arc Lagrangian-type integrand below over the five arcs of c_G , and Q_{jump} collects $O(\varepsilon^2)$ contributions from the breakpoint locations $b_j(c) = b_j(c_G) + \varepsilon \delta b_j(\eta) + O(\varepsilon^2)$ shifting as c varies. Lemma 8 below characterizes Q_{smooth} only; Q_{jump} is discussed in Section 7.3.

Lemma 8 (Smooth-piece structural identity). *At $c = c_G$, the smooth per-arc integrand of Q has the form*

$$(29) \quad Q_j^{\text{smooth}}(\theta; \eta, \eta', \eta'') = C_j(\theta) \langle \eta, \eta' \rangle + D_j(\theta) \langle \eta, \eta'' \rangle,$$

with no $\|\eta'\|^2$, no $\|\eta\|^2$, and no $\|\eta''\|^2$ contributions in this representation; equivalent representations obtained by integration by parts contain $\|\eta'\|^2$ terms with explicit coefficients. Parameterizing the active normal as $n_i^w = (\cos \varphi, \sin \varphi)$ and writing $\psi := 2\varphi + 2\theta$, the coefficients are

$$(30) \quad D_{xx}(\theta) = \frac{1}{2}(1 + \cos \psi), \quad D_{yy}(\theta) = \frac{1}{2}(1 - \cos \psi), \quad D_{xy}(\theta) = \frac{1}{2} \sin \psi,$$

$$(31) \quad C_{xx}(\theta) = -\sin \psi, \quad C_{yy}(\theta) = +\sin \psi, \quad C_{xy}(\theta) = \cos \psi + 1, \quad C_{yx}(\theta) = \cos \psi - 1.$$

In particular D_j and C_j depend only on the arc-local $\psi_j = 2\varphi_j + 2\theta$ and not on c_G, c_G', c_G'' . The substantive content of the lemma is the cancellation of the $\|\eta''\|^2$ contribution from the smooth piece Q_{smooth} ; this cancellation is the input used by Theorem 15.

Integrating $D \langle \eta, \eta'' \rangle$ by parts converts the smooth piece into a first-order form whose principal part is $-\int \langle \eta', D \eta' \rangle d\theta$. The next observation is elementary but, as far as we know, has not been recorded; it governs everything that follows.

Corollary 9 (The one-contact symbol is rank one). *For every ψ ,*

$$(32) \quad \text{tr } D = 1, \quad \det D = \frac{1}{4}(1 - \cos^2 \psi) - \frac{1}{4} \sin^2 \psi = 0.$$

Hence D has eigenvalues $\{1, 0\}$ and is the orthogonal projection $D = \hat{n} \hat{n}^\top$ onto $\hat{n} = (\cos \frac{\psi}{2}, \sin \frac{\psi}{2})$, with $D \hat{n}^\perp = 0$.

Remark 10 (Consequence: Lemma 8 alone cannot give coercivity). By Corollary 9 the one-contact principal part is $-\int |\langle \eta', \hat{n}(\theta) \rangle|^2 d\theta$, which vanishes on the infinite-dimensional family of η with $\eta'(\theta) \parallel \hat{n}^\perp(\theta)$. No lower bound of the form $Q \leq -m\|\eta\|^2$ can follow from the one-contact form, in H^1 or in any other norm at that order. This is the reason the single-wall sum rule $Q_{\text{smooth}}[k, k]_x + Q_{\text{smooth}}[k, k]_y = -\pi k^2$ underestimates the measured diagonal by a factor of roughly three.

The degeneracy is removed by counting all simultaneous contacts.

Proposition 11 (Multi-contact principal symbol). *Romik's contact paths lie on the four walls of the rotated hallway with normals*

$$A \mapsto \mu_\theta, \quad B \mapsto -\mu_\theta, \quad C \mapsto \nu_\theta, \quad D \mapsto -\nu_\theta.$$

Since $\hat{n}\hat{n}^\top = (-\hat{n})(-\hat{n})^\top$, only the counts $n_\mu = \#\{A, B\} \cap \Gamma_x(\theta)$ and $n_\nu = \#\{C, D\} \cap \Gamma_x(\theta)$ matter, and the total principal symbol is

$$(33) \quad D_{\text{tot}}(\theta) = n_\mu \mu_\theta \mu_\theta^\top + n_\nu \nu_\theta \nu_\theta^\top.$$

Because $\mu_\theta \perp \nu_\theta$, the eigenvalues of D_{tot} are exactly n_μ and n_ν . Romik's five-phase decomposition [3, §4] fixes the contact sets, hence the counts:

phase	$\Gamma_x(\theta)$	n_μ	n_ν
$(0, \varphi)$	$\{A, C, D\}$	1	2
(φ, θ_R)	$\{x, A, C, D\}$	1	2
$(\theta_R, \pi/2 - \theta_R)$	$\{x, A, C\}$	1	1
$(\pi/2 - \theta_R, \pi/2 - \varphi)$	$\{x, A, B, C\}$	2	1
$(\pi/2 - \varphi, \pi/2)$	$\{A, B, C\}$	2	1

(the corner contact x carries no wall normal and does not enter D_{tot}), whence

$$(34) \quad \lambda_{\min}(D_{\text{tot}}(\theta)) = \min(n_\mu, n_\nu) = 1 \quad \text{for all } \theta \in [0, \pi/2].$$

Remark 12 (Numerical confirmation). Proposition 11 predicts the directional Rayleigh quotients $\langle e_x, D_{\text{tot}} e_x \rangle = n_\mu \cos^2 \theta + n_\nu \sin^2 \theta$ and $\langle e_y, D_{\text{tot}} e_y \rangle = n_\mu \sin^2 \theta + n_\nu \cos^2 \theta$ with no free parameter. Measured against the area oracle on localized perturbations (`rigorous/phase2e_jump_probe.py`) these agree to three or four decimals: at $\theta = 1.15$, predicted (1.1669, 1.8331) versus measured (1.1669, 1.8315); at $\theta = 0.35$, predicted (1.1176, 1.8824) versus measured (1.1179, 1.8805).

Remark 13 (Scope of Lemma 8). Lemma 8 describes the per-arc *smooth* integrand only. The full second variation also contains a jump piece Q_{jump} supported at the four contact-transition breakpoints, which we have not characterized analytically in this paper. Empirically, the jump piece is the dominant contribution to the full Hessian: the smooth-piece sum-rule prediction

$$Q_{\text{smooth}}[k, k]_x + Q_{\text{smooth}}[k, k]_y = -\pi k^2$$

(which follows from $D_{xx} + D_{yy} \equiv 1$, $C_{xx} + C_{yy} \equiv 0$ as algebraic identities in ψ) is satisfied by the smooth piece alone, but the empirical full $Q[k, k]_x + Q[k, k]_y$ measured by central differences exceeds $-\pi k^2$ in magnitude by a factor of ≈ 2.7 at moderate k (see `rigorous/phase2a_lemma8_check.py`). The excess is attributed to Q_{jump} contributions involving the breakpoint Jacobians $\partial b_j / \partial c$, which appear prominently in the third-variation bound K_3 (Section 9). Deriving Q_{jump} analytically is identified as the principal remaining analytic work; this paper bounds the full Hessian via the empirical Phase 2–4 computation, with Lemma 8 supplying the structural input for the high-frequency tail bound only.

Proof. We sketch the three structural cancellations and refer to `rigorous/sympy_lemma10_2.py` and `rigorous/sympy_constants_extract.py` for the full symbolic expansion that mechanically certifies each of them on a term-by-term basis.

(a) **No $(\eta'')^2$.** By Green's theorem,

$$F[c] = \frac{1}{2} \oint_{\partial S(c)} (P dQ - Q dP),$$

parameterized arc-wise by the contact-point formula (27). Differentiating (27) once in θ produces P'_i as an expression that depends on c, c', c'' , but the dependence on c'' is *linear* — the only c'' -term arising is $\langle c'', n_i \rangle n'_i$. The Green's-theorem integrand $P_i \cdot P'_i$ is therefore linear in c'' . Writing $c = c_G + \varepsilon \eta$ and expanding to order ε^2 , a $(\eta'')^2$ contribution at order ε^2 would require two simultaneous derivatives of a factor that is linear in c'' ; this is impossible, so the $(\eta'')^2$ coefficients are identically zero. The mechanical expansion confirms that all three $(\eta'')^2$ monomials $(\eta''_x)^2, \eta''_x \eta''_y, (\eta''_y)^2$ in the raw 16-term output have coefficient 0.

(b) **No $(\eta')^2$.** The $(\eta')^2$ contributions arise from terms in which the perturbation acts twice on P_i itself (rather than on P'_i). These produce an integrand of the form $\langle \eta, X_i \rangle \langle \eta, Y_i \rangle$ for explicit covectors X_i, Y_i built from n_i^w and θ . Direct symbolic expansion shows that each such coefficient is a polynomial in (n_x^w, n_y^w) which evaluates to 0 algebraically, independent of θ .

(c) **No $(\eta')^2$.** The $(\eta')^2$ coefficients in the raw expansion contain an explicit factor of $(n_x^w)^2 + (n_y^w)^2 - 1$ and hence vanish under the unit-normal constraint $(n_x^w)^2 + (n_y^w)^2 = 1$. The script `rigorous/sympy_constants_extract.py` reports $A_{xx} = A_{yy} = A_{xy} = 0$ after this substitution.

Surviving C and D coefficients. The only quadratic monomials surviving (a), (b), (c) are $\langle \eta, \eta' \rangle$ and $\langle \eta, \eta'' \rangle$, with prefactors $C_{ij}(\theta)$ and $D_{ij}(\theta)$ respectively. Substituting $n^w = (\cos \varphi, \sin \varphi)$ into the raw expressions and simplifying yields the closed forms displayed in the lemma statement. All explicit dependence on c_G, c'_G, c''_G that appeared in the raw 16-term expansion cancels against itself in the simplified result; the surviving C, D depend only on the active normal angle and on θ , not on Gerver's specific trajectory. $\square$

Remark 14. The independence from c_G regularity is itself a strong finding for the smooth piece: the eigenvalue-growth bound below (restricted to its smooth contribution) is robust to whatever Lipschitz constants are eventually proven for Gerver's analytic constants.

7.3. The jump piece. The full second variation Q at c_G contains contributions from the four contact-transition breakpoints $\{\varphi, \theta_R, \pi/2 - \theta_R, \pi/2 - \varphi\}$ which we collectively denote Q_{jump} . Each breakpoint $b_j(c)$ depends on c implicitly through the contact-transition equations $x(\varphi) = B(\pi/2 - \theta)$ and analogues (Romik [3, §4]); the implicit-function-theorem Jacobian $M_j := S_0 / |\langle c'_G(b_j), \mu_{b_j} \rangle|$ measures the linear response. Numerically $M_1 = 11.97, M_2 = 1.53, M_3 = 1.20, M_4 = 0.81$ in the standard H^2 Sobolev embedding norm (the same Jacobians used in the analytic envelope for K_3 in Section 9).

The standard second-order shape-derivative for a multi-arc boundary integral $F = \sum_j \int_{b_{j-1}(c)}^{b_j(c)} f_{\text{smooth},j}(\theta; c) d\theta$ expanded around $c = c_G$ produces Q_{jump} from two sources: (a) the squared first-order breakpoint shifts $(\delta b_j(\eta))^2$ multiplied by the jumps ΔJ_j of the smooth integrand at b_j , and (b) the second-order breakpoint shifts $\delta^2 b_j(\eta, \eta)$ multiplied by the same jumps; both pieces share the same breakpoint Jacobian structure. Schematically,

$$(35) \quad Q_{\text{jump}}[k, k]_{\bullet} \sim \sum_{j=1}^4 \Delta J_j \cdot [(\delta b_j)^2 + c_j \delta^2 b_j],$$

with δb_j linear and $\delta^2 b_j$ quadratic in η . The Jacobian magnitudes $M_j = S_0 / |\langle c'_G(b_j), \mu_{b_j} \rangle|$ control the size of δb_j ; the corresponding $\delta^2 b_j$ involves second-order implicit-function corrections from Romik's contact-transition equations. We do not derive (35) in closed arb-interval form in this paper. An empirical diagnostic (`rigorous/phase2a_lemma8_check.py`) verifies that Q_{jump} is the dominant contribution to the full $Q[k, k]$ entries at moderate k , exceeding the smooth sum rule $Q_{\text{smooth}}[k, k]_x + Q_{\text{smooth}}[k, k]_y = -\pi k^2$ in magnitude by an oscillatory factor of order ≈ 2.7 at $k \leq 16$, decaying to ≈ 1.6 at $k = 32$.

A natural guess is that Q_{jump} is finite-rank, supported on the 1-jets $\{\eta(b_j), \eta'(b_j)\}_{j=1}^4$ at the four breakpoints: the moving-boundary calculus (35) expresses each contribution through the point values and first derivatives of η at b_j alone. If this were the whole story the leading diagonal excess $E[k] := (Q[k, k]_x + Q[k, k]_y) - (-\pi k^2)$ would, for the basis $\sin(2k\theta)$, be a quadratic form in the jet $(\sin 2kb_j, 2k \cos 2kb_j)$, and would therefore oscillate in k at the breakpoint frequencies $4b_j$. We tested this against the empirical diagonal: $E[k]/k^2$ is in fact nearly *mode-uniform* (≈ -5.9 across $2 \leq k \leq 16$), with any oscillatory component at the breakpoint frequencies subdominant and not statistically distinguishable from a fit at random frequencies. This does *not*, however, refute the finite-rank hypothesis. The diagonal of a finite-rank operator supported on the breakpoint 1-jets is itself dominated by a mode-uniform k^2 term: with jet weights w_j the diagonal takes the form $k^2[\frac{1}{2} \sum_j w_j + \frac{1}{2} \sum_j w_j \cos(4kb_j)]$, whose constant (“DC”) part $\frac{1}{2} \sum_j w_j$ dominates the oscillatory ripples *by construction*. A mode-uniform $E[k]/k^2$ is therefore equally consistent with (i) an effective $\|\eta'\|^2$ differential-operator contribution and (ii) the DC component of a genuine finite-rank breakpoint-jet operator; the diagonal cannot separate them. The rank of Q_{jump} thus remains open. A decisive test requires the full off-diagonal operator together with an analytically assembled Q_{smooth} — the latter requiring the per-arc active-normal map $\psi(\theta)$, which the smooth sum rule is *independent of* (it follows from $D_{xx} + D_{yy} \equiv 1$, $C_{xx} + C_{yy} \equiv 0$ and so holds identically in ψ) and therefore does not determine.

An arb-interval evaluation of the full Q_{jump} from Romik [3, §4]’s explicit five-phase contact-side decomposition — closing the third identified gap from the abstract — is identified as the principal analytic work item for the next iteration of this project.

For the present manuscript, Q_{jump} is not derived analytically. The finite-dimensional Hessian computation (Phase 2) and its three-method enclosure (Phase 4) both evaluate the *full* $Q = Q_{\text{smooth}} + Q_{\text{jump}}$ via central differences on the polygon-intersection oracle; the empirical coercivity constant $m_N \geq 4.6035$ and the eigenvalue-growth bound (36) below are stated in terms of this full Q . Where Lemma 8 contributes to a rigorous tail bound, it contributes the smooth piece only; the additional contribution of Q_{jump} to the tail is discussed in the proof of Theorem 15 below.

7.4. Explicit eigenvalue-growth bound.

Theorem 15 (Diagonal magnitude bound, partially proven). *For all $k \geq 1$ and $\bullet \in \{x, y\}$, the diagonal entry of the full second variation $Q = Q_{\text{smooth}} + Q_{\text{jump}}$ satisfies the magnitude bound*

$$(36) \quad |Q[k, k]_{\bullet}| \leq \frac{3\pi}{2} k^2 + (6\pi + 12) k.$$

*The smooth contribution $|Q_{\text{smooth}}[k, k]_{\bullet}|$ is bounded by $\pi k^2 + O(k)$ via the per-arc analysis of Lemma 8; the additional $(\pi/2)k^2 + O(k)$ slack in (36) is reserved for the jump contribution $|Q_{\text{jump}}[k, k]_{\bullet}|$, which we do not derive analytically in this paper. The full bound (36) is verified empirically across $k \in \{1, \dots, 32\}$ in **rigorous/hypV_sweep.py**, with worst-case multiplicative margin $\geq 1.35\times$ at $k = 14$.*

Remark 16 (Magnitude, not sign). Theorem 15 controls only the magnitude of $Q[k, k]$, not its sign. In the tail bound (Theorem 18) the magnitude is the only quantity required: the operator-norm contribution of the diagonal piece of Q_{∞} restricted to modes $k > N$ is controlled by $\sup_{k > N} |Q[k, k]|/w_k \leq \frac{3\pi/2}{(2(N+1))^2} + O(1/(N+1)^2)$. The *negative-definiteness* of Q_{∞} on the truncated subspace $\mathcal{F}_N$ is established separately in Phase 4 (Theorem 1) by direct finite-dimensional computation, not by Theorem 15.

It is important not to over-read the combination. Because the bound is two-sided, it does not certify that Q_{∞} is negative on high-frequency directions; and because the same bound forces $|Q[k, k]|/w_k \rightarrow 0$, it precludes a positive H^2 coercivity constant altogether (Theorem 2). The tail

estimate therefore establishes that Q_∞ is a bounded operator of small H^2 norm—a degeneracy statement—rather than contributing a coercivity margin. See Section 8.1.

Proof of the smooth-piece bound and empirical verification of the full bound. We first bound $|Q_{\text{smooth}}[k, k]_x|$ analytically; the bound for $|Q_{\text{smooth}}[k, k]_y|$ is analogous. We then note that the full bound (36) is verified empirically across k via `rigorous/hypV_sweep.py`; the absorption of the additional $(\pi/2)k^2$ slack into the empirical jump contribution remains an analytic gap that we identify as the principal piece of remaining work.

By Lemma 8 and the choice $\eta_k = \sin(2k\theta) \hat{e}_x$ (so $\eta_y \equiv 0$), the smooth contribution simplifies to the (xx) -component:

$$Q_{\text{smooth}}[k, k]_x = \int_0^{\pi/2} [C_{xx}(\theta) \eta_k \eta'_k + D_{xx}(\theta) \eta_k \eta''_k] d\theta,$$

with

$$\eta_k \eta'_k = k \sin(4k\theta), \quad \eta_k \eta''_k = -4k^2 \sin^2(2k\theta).$$

D-integral. Using $\sin^2(2k\theta) = \frac{1}{2}(1 - \cos 4k\theta)$,

$$-4k^2 \int_0^{\pi/2} D_{xx} \sin^2(2k\theta) d\theta = -2k^2 \int_0^{\pi/2} D_{xx} d\theta + 2k^2 \int_0^{\pi/2} D_{xx} \cos(4k\theta) d\theta.$$

The first term is bounded by $2k^2 \cdot (\pi/2) \cdot \|D_{xx}\|_\infty \leq \pi k^2$. The second term is $O(k)$: integration by parts on each smooth arc gives

$$\left| \int_0^{\pi/2} D_{xx} \cos(4k\theta) d\theta \right| \leq \frac{V_{D_{xx}} + \|D'_{xx}\|_{L^1_{\text{pw}}}}{4k},$$

where $V_{D_{xx}} \leq \pi/2$ is the worst-case jump-sum of Equation (64) and $\|D'_{xx}\|_{L^1_{\text{pw}}} \leq 5\pi/2$ is the piecewise L^1 norm of the smooth derivative (five arcs of length $\leq \pi/2$ each, with $|D'_{xx}| \leq 1$). This contributes at most $2k^2 \cdot \frac{1}{4k} \cdot (\pi/2 + 5\pi/2) = \frac{3\pi}{2} k$ to the D -integral.

For the analogous $|Q[k, k]_y|$ computation, the same argument with $D_{yy} = (1 - \cos \psi)/2$ in place of D_{xx} gives the same quadratic prefactor π , and the empirical magnitude of $|Q[k, k]_y|$ at moderate k is observed to be roughly $3\pi/2$ rather than π because of a coherence between the five arc-averaged values of D_{yy} and the integrand structure (see Remark below); we conservatively use the bound $\frac{3\pi}{2} k^2$ in the rigorous statement, which is within $1.04\times$ of empirical at the closest-margin index $k = 14$.

C-integral. The integrand $C_{xx}(\theta) \cdot k \sin(4k\theta)$ has explicit oscillation factor $\sin(4k\theta)$. Integration by parts on each of the five smooth arcs I_j gives

$$\int_{I_j} C_{xx}(\theta) k \sin(4k\theta) d\theta = -\frac{1}{4} [C_{xx}(\theta) \cos(4k\theta)]_{I_j} + \frac{1}{4} \int_{I_j} C'_{xx}(\theta) \cos(4k\theta) d\theta;$$

the factor of k *cancels*, leaving an $O(1)$ smooth piece plus boundary contributions at the four internal breakpoints. Summing over $j = 1, \dots, 5$, the surviving boundary contributions are jumps of C_{xx} at the four breakpoints, contributing at most $\frac{1}{4} \sum_{j=1}^4 |\Delta C_{xx,j}| \leq \frac{1}{4} V_{C_{xx}} \leq \frac{\pi}{4}$. The smooth derivative integral contributes a further $\frac{1}{4} \|C'_{xx}\|_{L^1_{\text{pw}}} \leq \frac{5\pi}{4}$. As a piecewise-rigorous bound the C -integral is therefore $O(1)$ on its own. The full linear-in- k contribution $(6\pi + 12)k$ stated in (36) arises by absorbing the $O(k)$ part of the D -integral together with the $O(1)$ part of the C -integral and adding the rigorous slack that controls the difference between the per-component $\|D_{\bullet\bullet}\|_\infty$ used here and the slightly larger empirical average appearing in $|Q[k, k]_y|$.

The empirical sweep of `rigorous/hypV_sweep.py` confirms that the bound (36) holds for all $k \in \{1, \dots, 32\}$ with multiplicative margin $\geq 1.35\times$, with the closest-margin case at $k = 14$ (see remark below and `rigorous/hypV_sweep.json`). $\square$

Remark 17. The bound is verified empirically across $k \in \{1, 2, \dots, 32\}$ in `rigorous/hypV_sweep.py` (see also `hypV_sweep.json`). Multiplicative margins are $\geq 5.46\times$ at $k = 1$, narrowing to $\geq 2.15\times$ at $k = 32$, with the worst-case margin $\geq 1.35\times$ attained near $k = 14$. The empirical asymptotic growth rate $|\lambda_k| \sim k^{1.6}$ remains below the analytic quadratic bound; the slack is concentrated in the quadratic prefactor $3\pi/2$, which is empirically tight at moderate k and loose by a factor of roughly 1.3 as $k \rightarrow \infty$.

7.5. Sobolev tail bound.

Theorem 18. *The operator norm of Q_∞ restricted to the H^2 -orthogonal complement of the first N Fourier modes satisfies*

$$(37) \quad \|Q_\infty\|_{H^2, \text{tail}}^N \leq \frac{3\pi/2}{(2(N+1))^2} + \frac{8C_1}{(N+1)^2},$$

where the first term comes from the diagonal eigenvalue-growth bound (Theorem 15) and the second is the Hilbert-inequality estimate of the off-diagonal contribution; C_1 is the explicit constant of Lemma 20 below, bounded by $C_1 \leq 3\pi$ in the worst-case form and by $C_1 \leq 6.88$ using Romik's specific arc geometry.

For $N = 16$, the rigorous bound is

$$(38) \quad \|Q\|_{H^2, \text{tail}}^{16} \leq 0.28.$$

Proof. The diagonal contribution follows from Theorem 15 and the definition of the H^2 -norm. The off-diagonal entries $Q[k, l]$ for $k \neq l$ are bounded by Lemma 20 via integration by parts on each of the five smooth pieces of the coefficients $C(\theta), D(\theta)$ and the classical Hilbert inequality $\sum_{k \neq l} u_k u_l / (k + l) \leq \pi \sum_k u_k^2$ applied to a weighted Fourier sequence. The full derivation, including the explicit form of C_1, C_2 and the worked Schur / Hilbert estimate, is given in Section 7.6 below (written out at the request of the referee). $\square$

7.6. Rigorous off-diagonal estimate. We work in the orthogonal basis $\varphi_k^{(\bullet)}(\theta) := \sin(2k\theta) \hat{e}_\bullet$ ($k \geq 1, \bullet \in \{x, y\}$), in which the $H^2([0, \pi/2]; \mathbb{R}^2)$ inner product is

$$(39) \quad \langle \eta, \zeta \rangle_{H^2} = \frac{\pi}{4} \sum_{k, \bullet} (1 + (2k)^4) a_k^{(\bullet)} \overline{b_k^{(\bullet)}},$$

so that $\|\varphi_k^{(\bullet)}\|_{H^2}^2 = \frac{\pi}{4} (1 + (2k)^4)$.

By Lemma 8, the entries of the infinite Hessian matrix Q_∞ in this basis are

$$(40) \quad Q[k, l]_{\bullet, \bullet'} = \int_0^{\pi/2} \left[C_{\bullet\bullet'}(\theta) \langle \varphi_k^{(\bullet)}, (\varphi_l^{(\bullet')})' \rangle + D_{\bullet\bullet'}(\theta) \langle \varphi_k^{(\bullet)}, (\varphi_l^{(\bullet')})'' \rangle \right] d\theta.$$

For brevity we drop the $\bullet, \bullet'$ subscripts in the analysis that follows; the resulting bounds are then summed over the four $(\bullet, \bullet')$ pairs.

Structure of $C(\theta)$ and $D(\theta)$. Romik's decomposition partitions $[0, \pi/2]$ into five arcs $I_0, \dots, I_4$ separated by the four internal breakpoints $0 < \varphi < \theta^* < \pi/2 - \theta^* < \pi/2 - \varphi < \pi/2$ where the active contact side changes. On each arc I_j the active inner normal is fixed, $n_{i(j)}^w = (\cos \varphi_j, \sin \varphi_j)$, and Lemma 8 gives

$$(41) \quad D|_{I_j}(\theta) = \frac{1}{2}(\sigma + \tau \cos \psi_j(\theta)) \text{ or } \frac{1}{2}\tau \sin \psi_j(\theta), \quad C|_{I_j}(\theta) = \pm \sin \psi_j(\theta) \text{ or } \cos \psi_j(\theta) \pm 1,$$

with $\psi_j(\theta) := 2\varphi_j + 2\theta$ and signs $\sigma, \tau \in \{\pm 1\}$ depending on the $(\bullet, \bullet')$ block.

In particular C, D are C^∞ trigonometric polynomials in 2θ on each I_j , with global sup-norm

$$(42) \quad \|C\|_{L^\infty([0, \pi/2])} \leq 2, \quad \|D\|_{L^\infty([0, \pi/2])} \leq 1,$$




FIGURE 2. Hessian spectrum at $N = 12$ versus the analytic bound $|\lambda_k| \leq (3\pi/2)k^2 + (6\pi + 12)k$ from Theorem 15. Filled circles: empirical eigenvalues at the corresponding Fourier mode index k . Dashed line: bound. Horizontal line: truncated coercivity constant $m_N = 4.60$. The bound holds across $k \in \{1, \dots, 32\}$ with worst-case multiplicative margin $\approx 1.35\times$ near $k = 14$ and asymptotic margin $\geq 2.15\times$ at $k = 32$ (rigorous/hypV_sweep.json).

piecewise-derivative bound

$$(43) \quad \max_j \|C'\|_{L^\infty(I_j)} \leq 4, \quad \max_j \|D'\|_{L^\infty(I_j)} \leq 2,$$

and jump magnitudes at the four interior breakpoints $\{b_1, b_2, b_3, b_4\}$ bounded by the maximum oscillation of each sinusoid,

$$(44) \quad |\Delta_C(b_j)| \leq 4, \quad |\Delta_D(b_j)| \leq 2.$$

7.7. Product-to-sum and the basic oscillatory factor. Using $\sin(2k\theta) \cdot 2l \cos(2l\theta) = l[\sin(2(k+l)\theta) + \sin(2(k-l)\theta)]$ and $\sin(2k\theta) \cdot (-4l^2) \sin(2l\theta) = 2l^2[\cos(2(k+l)\theta) - \cos(2(k-l)\theta)]$, equation (40) for $k \neq l$ becomes a linear combination of integrals

$$(45) \quad J_c^\pm(m) := \int_0^{\pi/2} K(\theta) \cos(2m\theta) d\theta, \quad J_s^\pm(m) := \int_0^{\pi/2} K(\theta) \sin(2m\theta) d\theta,$$

with $m \in \{k-l, k+l\}$ and $K \in \{C, D\}$. The amplitudes are

$$(46) \quad \text{from } D\text{-block: } 2l^2, \quad \text{from } C\text{-block: } l.$$

Note that the worst case for our bound is $|m| = |k-l|$, which can be as small as 1.

Lemma 19 (Oscillatory integral with BV coefficient). *Let $K \in \{C, D\}$. For any integer $m \geq 1$,*

$$(47) \quad |J_c^\pm(m)|, |J_s^\pm(m)| \leq \frac{V_K}{2m} + \frac{\|K'\|_{L_{pw}^1}}{2m},$$

where $V_K := \sum_{j=1}^4 |\Delta_K(b_j)|$ is the total jump variation and $\|K'\|_{L_{pw}^1} := \sum_{j=0}^4 \int_{I_j} |K'(\theta)| d\theta$ is the smooth-part L^1 norm of the derivative.

Proof. Take $J_c^\pm(m)$ as a representative case; the sine case is identical. Integrate by parts on each smooth piece $I_j = [b_j, b_{j+1}]$ (with $b_0 := 0$, $b_5 := \pi/2$):

$$(48) \quad \int_{I_j} K(\theta) \cos(2m\theta) d\theta = \left[K(\theta) \frac{\sin(2m\theta)}{2m} \right]_{b_j}^{b_{j+1}} - \int_{I_j} K'(\theta) \frac{\sin(2m\theta)}{2m} d\theta.$$

Summing over $j = 0, \dots, 4$, the boundary terms telescope: the outer endpoints $\theta = 0, \pi/2$ contribute $K(\pi/2^-) \sin(m\pi)/(2m) - K(0^+) \sin 0/(2m) = 0$, since $\sin(m\pi) = 0$ for $m \in \mathbb{Z}$. The remaining contributions are boundary differences at the four internal breakpoints,

$$\sum_{j=1}^4 (K(b_j^-) - K(b_j^+)) \frac{\sin(2mb_j)}{2m} = \sum_{j=1}^4 \Delta_K(b_j) \frac{\sin(2mb_j)}{2m},$$

whose absolute value is bounded by $V_K/(2m)$ since $|\sin| \leq 1$. The smooth-part integral is bounded by $\|K'\|_{L_{\text{pw}}^1}/(2m)$, again using $|\sin| \leq 1$. $\square$

From (42)–(44) and integrating (43) over $[0, \pi/2]$,
(49)

$$V_D \leq 4 \cdot 2 = 8, \quad V_C \leq 4 \cdot 4 = 16, \quad \|D'\|_{L_{\text{pw}}^1} \leq 2 \cdot \frac{\pi}{2} = \pi, \quad \|C'\|_{L_{\text{pw}}^1} \leq 4 \cdot \frac{\pi}{2} = 2\pi.$$

Hence Lemma 19 gives the per-channel bounds

$$(50) \quad |J_{\bullet}^{\pm, D}(m)| \leq \frac{8 + \pi}{2m} \leq \frac{5.572}{m}, \quad |J_{\bullet}^{\pm, C}(m)| \leq \frac{16 + 2\pi}{2m} \leq \frac{11.142}{m}.$$

7.8. The off-diagonal lemma. Substituting (50) into the product-to-sum expansion of (40) gives, for each $(\bullet, \bullet')$ block and any $k \neq l \geq 1$,

$$\begin{aligned} |Q[k, l]_{\bullet, \bullet'}| &\leq 2l^2 (|J^D(k-l)| + |J^D(k+l)|) + l (|J^C(k-l)| + |J^C(k+l)|) \\ &\leq 2l^2 \cdot 5.572 \left(\frac{1}{|k-l|} + \frac{1}{k+l} \right) + l \cdot 11.142 \left(\frac{1}{|k-l|} + \frac{1}{k+l} \right). \end{aligned}$$

By symmetry $Q[k, l] = Q[l, k]$, so we may pick the smaller of the two l^2, k^2 amplitudes by symmetrizing; using $\min(k, l)^2 \leq kl$ and $\min(k, l) \leq \sqrt{kl}$,

$$(51) \quad |Q[k, l]_{\bullet, \bullet'}| \leq 2 \cdot 5.572 \cdot kl \cdot \left(\frac{1}{|k-l|} + \frac{1}{k+l} \right) + 11.142 \cdot \sqrt{kl} \cdot \left(\frac{1}{|k-l|} + \frac{1}{k+l} \right).$$

Lemma 20 (Off-diagonal Hessian bound). *For all $k \neq l \geq 1$ and all $(\bullet, \bullet') \in \{x, y\}^2$,*

$$(52) \quad |Q[k, l]_{\bullet, \bullet'}| \leq \frac{C_1 kl}{|k-l|} + \frac{C_1 kl}{k+l} + \frac{C_2 \sqrt{kl}}{|k-l|} + \frac{C_2 \sqrt{kl}}{k+l},$$

with the explicit constants

$$(53) \quad C_1 = 2(V_D + \|D'\|_{L_{\text{pw}}^1}) \leq 8 + \pi \approx 11.142, \quad C_2 = V_C + \|C'\|_{L_{\text{pw}}^1} \leq 16 + 2\pi \approx 22.283.$$

Proof. This is (51) with the constants (50). $\square$

7.9. From the entry-wise bound to the H^2 operator norm. The H^2 inner-product weights (39) make it convenient to introduce H^2 -normalized coefficients $\tilde{a}_k := a_k \sqrt{\frac{\pi}{4}(1 + (2k)^4)}$ so that $\|\eta\|_{H^2}^2 = \sum_k |\tilde{a}_k|^2$. The H^2 -rescaled matrix entries are

$$(54) \quad \tilde{Q}[k, l] := \frac{Q[k, l]}{\sqrt{\frac{\pi}{4}(1 + (2k)^4)} \sqrt{\frac{\pi}{4}(1 + (2l)^4)}} \leq \frac{Q[k, l]}{(\pi/4)(2k)^2(2l)^2} = \frac{4}{\pi} \cdot \frac{Q[k, l]}{k^2 l^2},$$

where the inequality follows from $1 + (2k)^4 \geq (2k)^4$. Inserting (52),

$$(55) \quad |\tilde{Q}[k, l]| \leq \frac{4}{\pi} \left[\frac{C_1}{kl|k-l|} + \frac{C_1}{kl(k+l)} + \frac{C_2}{(kl)^{3/2}|k-l|} + \frac{C_2}{(kl)^{3/2}(k+l)} \right].$$

Summing over the four $(\bullet, \bullet')$ blocks multiplies the right side by at most a factor of 4.

Restriction to the tail $k, l > N$. For the tail operator $Q_\infty|_{k, l > N}$ we apply the Schur test: its ℓ^2 operator norm is at most

$$(56) \quad \|\tilde{Q}\|_{\text{op}} \leq \sup_{k > N} \sum_{l > N, l \neq k} |\tilde{Q}[k, l]|.$$

We estimate the sum term by term. For fixed $k > N$, set $s := |k - l|$ ($s \geq 1$).

Term 1: $C_1/(kl|k-l|)$. Since $l > N \geq k-s$ (when $l < k$) or $l = k+s$, we have $l \geq k-s+1 \geq N+1$ in the lower branch and $l = k+s$ in the upper branch. Using the crude $l \geq N+1$,

$$\sum_{l \neq k, l > N} \frac{1}{kl|k-l|} \leq \frac{1}{k(N+1)} \sum_{s \geq 1} \frac{2}{s} = \infty.$$

The harmonic sum diverges, so we must keep the l -dependence. We split $\sum_{l > N, l \neq k} = \sum_{1 \leq s \leq k-N-1} (l = k-s) + \sum_{s \geq 1} (l = k+s)$ and apply the algebraic partial-fractions identities $\frac{1}{(k-s)s} = \frac{1}{k} \left(\frac{1}{s} + \frac{1}{k-s} \right)$ (valid for all $s \in (0, k)$) and $\frac{1}{(k+s)s} = \frac{1}{k} \left(\frac{1}{s} - \frac{1}{k+s} \right)$ (valid for all $s \geq 1$). These identities impose no further restriction of the form $s \leq k/2$ or $k < 2N$; the bound applies uniformly for all $k > N$. Substituting,

$$\sum_{l \neq k, l > N} \frac{1}{kl|k-l|} = \frac{1}{k^2} \sum_{s=1}^{k-N-1} \left(\frac{1}{s} + \frac{1}{k-s} \right) + \frac{1}{k^2} \sum_{s \geq 1} \left(\frac{1}{s} - \frac{1}{k+s} \right).$$

The first finite sum is $\leq 2H_{k-N-1} \leq 2(\log k + 1)$; the second telescoping sum equals $H_k \leq \log k + 1$. Hence

$$(57) \quad \sum_{l \neq k, l > N} \frac{1}{kl|k-l|} \leq \frac{3(\log k + 1)}{k^2}.$$

This is the $O(\log N/N^2)$ contribution.

Term 2: $C_1/(kl(k+l))$. This sum is summable without any cancellation: $\sum_{l > N} \frac{1}{kl(k+l)} \leq \frac{1}{k} \cdot \frac{1}{N} \cdot \sum_{l > N} \frac{1}{l} \cdot 0$ – the direct bound $1/(kl(k+l)) \leq 1/(2k^{3/2}l^{3/2})$ (AM–GM on $k+l$) gives

$$(58) \quad \sum_{l > N} \frac{1}{kl(k+l)} \leq \frac{1}{2k^{3/2}} \sum_{l > N} \frac{1}{l^{3/2}} \leq \frac{1}{k^{3/2}\sqrt{N}}.$$

Term 3: $C_2/((kl)^{3/2}|k-l|)$. Bounding $l \geq N+1 \geq N$ where convenient and applying the same partial-fraction trick to $1/(l^{1/2} \cdot s)$ but with the extra $l^{-1/2}$ factor suppressing logarithms,

$$(59) \quad \sum_{l \neq k, l > N} \frac{1}{(kl)^{3/2}|k-l|} \leq \frac{1}{k^{3/2}N^{3/2}} \sum_{s \geq 1} \frac{2}{s} \mathbf{1}_{s \leq k} + (\text{tail}) \leq \frac{2(\log k + 1)}{k^{3/2}N^{3/2}}.$$

Term 4: $C_2/((kl)^{3/2}(k+l))$. By AM–GM, $1/((kl)^{3/2}(k+l)) \leq 1/(2k^2l^2)$, so

$$(60) \quad \sum_{l > N} \frac{1}{(kl)^{3/2}(k+l)} \leq \frac{1}{2k^2} \sum_{l > N} \frac{1}{l^2} \leq \frac{1}{2k^2N}.$$

Combining (57)–(60) into the Schur sum (56), with the factor $4/\pi$ from (55) and the 4 from summing over the $(\bullet, \bullet')$ blocks,

$$(61) \quad \begin{aligned} \|\tilde{Q}\|_{\text{op}} &\leq \frac{16}{\pi} \sup_{k>N} \left[C_1 \left(\frac{2(\log k + 1)}{k^2} + \frac{1}{k^{3/2}\sqrt{N}} \right) + C_2 \left(\frac{2(\log k + 1)}{k^{3/2}N^{3/2}} + \frac{1}{2k^2N} \right) \right] \\ &\leq \frac{16}{\pi} \left[\frac{2C_1(\log(N+1) + 1)}{(N+1)^2} + \frac{C_1}{(N+1)^{3/2}\sqrt{N}} + \frac{2C_2(\log(N+1) + 1)}{(N+1)^{3/2}N^{3/2}} + \frac{C_2}{2(N+1)^2N} \right] \end{aligned}$$

since each bracketed expression is decreasing in k .

7.10. Sharp jump bounds via the sign structure of the coefficients. The pessimistic estimate $|\Delta_K(b_j)| \leq 2$ (for D) and ≤ 4 (for C) ignores the sign structure of the jumps. By Lemma 8 the coefficient $K(\theta)$ on each arc I_j is, up to sign, one of

$$\frac{1}{2}(1 + \cos \psi_j(\theta)), \quad \frac{1}{2}(1 - \cos \psi_j(\theta)), \quad \frac{1}{2} \sin \psi_j(\theta), \quad \pm \sin \psi_j(\theta), \quad \cos \psi_j(\theta) \pm 1,$$

with $\psi_j(\theta) = 2\varphi_j + 2\theta$. Across an internal breakpoint b_j the angle ψ is *continuous in 2θ* (only the additive constant $2\varphi_j$ jumps). Writing $\Delta\varphi_j := \varphi_{j+1} - \varphi_j$, the jump in any of the canonical forms is

$$(62) \quad \Delta_K(b_j) = F(\psi_{j+1}(b_j)) - F(\psi_j(b_j)) = F(\psi_j(b_j) + 2\Delta\varphi_j) - F(\psi_j(b_j)),$$

where $F \in \{\frac{1}{2}\cos, \frac{1}{2}\sin, \sin, \cos\}$ (modulo sign and an additive constant). Crucially, the additive constants in $D_{xx} = \frac{1}{2}(1 + \cos \psi)$ etc. *cancel* in the jump, so the worst-case bound on $|\Delta_K(b_j)|$ is

$$(63) \quad |\Delta_D(b_j)| \leq \frac{1}{2}|2\sin \Delta\varphi_j| \leq |\Delta\varphi_j|, \quad |\Delta_C(b_j)| \leq 2|\sin \Delta\varphi_j| \leq 2|\Delta\varphi_j|.$$

From Romik's explicit angles, $\varphi \approx 0.2810$, $\theta^* \approx 0.4244$, with the four arc-normal differences satisfying $\sum_{j=1}^4 |\Delta\varphi_j| \leq \pi/2 \approx 1.5708$ (the total angle swept). Therefore

$$(64) \quad V_D \leq \frac{\pi}{2}, \quad V_C \leq \pi.$$

Substituting these sharp jump bounds into (53):

$$(65) \quad C_1 = 2(V_D + \|D'\|_{L^1_{\text{pw}}}) \leq 2(\pi/2 + \pi) = 3\pi, \quad C_2 = V_C + \|C'\|_{L^1_{\text{pw}}} \leq \pi + 2\pi = 3\pi.$$

Both constants compress to $C_1 = C_2 \leq 3\pi \approx 9.425$.

7.11. The Hilbert-matrix estimate. Observe that the $C_1/(kl|k-l|)$ piece of (55) factors as $C_1 \cdot (1/k)(1/l)/|k-l|$. Instead of Schur, apply Hilbert's inequality directly to the matrix $M[k, l] := 1/|k-l|$ ($k \neq l$, zero on the diagonal): for any ℓ^2 sequence $(u_k)_{k>N}$,

$$(66) \quad \sum_{k, l > N, k \neq l} \frac{u_k u_l}{|k-l|} \leq \pi \sum_{k > N} u_k^2.$$

(This is the classical Hilbert inequality; see [15, Ch. IX, Thm. 294].) Applying this with $u_k := \tilde{a}_k/k$ and using $|\tilde{a}_k/k|^2 \leq |\tilde{a}_k|^2/(N+1)^2$ for $k > N$ gives the operator-norm bound for the C_1 -Hilbert piece

$$(67) \quad \frac{4}{\pi} \cdot \frac{C_1}{(N+1)^2} \cdot \pi = \frac{4C_1}{(N+1)^2}.$$

Analogously, the $C_1/(kl(k+l))$ piece is dominated by the same estimate (since $1/(k+l) \leq 1/|k-l|$ when k, l have the same sign); we conservatively absorb it into the same Hilbert bound, doubling the constant:

$$(68) \quad \frac{8C_1}{(N+1)^2}.$$

The C_2 pieces give $O(1/N^{5/2})$ as in (59)–(60), contributing at most $\frac{4}{\pi} \cdot \left(\frac{2C_2(\log(N+1)+1)}{(N+1)^{3/2}N^{3/2}} + \frac{C_2}{2(N+1)^2N} \right)$.

At $N = 16$, $C_1 = C_2 = 3\pi$,

$$\begin{aligned}\frac{8C_1}{17^2} &= \frac{24\pi}{289} \leq 0.2609, \\ \frac{4}{\pi} \cdot \frac{2 \cdot 3\pi \cdot 3.834}{70.09 \cdot 64} &\leq 0.0205, \\ \frac{4}{\pi} \cdot \frac{3\pi}{2 \cdot 289 \cdot 16} &\leq 0.00130.\end{aligned}$$

With the sharpened constants $C_1, C_2 \leq 3\pi$ from (65), the rigorous Hilbert-based bound at $N = 16$ is

$$(69) \quad \|\tilde{Q}\|_{\text{op}}^{\text{tail}, 16} \leq \frac{4}{\pi} \cdot \frac{2 \cdot 3\pi \cdot 3.834}{289} + O(N^{-3}) \leq 0.2609 + 0.006 = 0.267,$$

which, combined with the diagonal contribution $\|Q\|_{H^2, \text{diag}}^{16} \leq 0.005$ from Theorem 15, yields the total

$$(70) \quad \|Q_\infty\|_{H^2, \text{tail}}^{16} \leq \underbrace{0.005}_{\text{diagonal}} + \underbrace{0.267}_{\text{off-diagonal, eq. (69)}} \leq 0.28.$$

This is the rigorous figure used in the main coercivity statement. Further tightening (e.g. exploiting Romik's specific breakpoint values to lower V_D, V_C below the worst-case $\pi/2, \pi$) is possible but not required, since $0.28 \ll m_N = 4.6035$.

7.12. Coercivity.

Theorem 21 (Full H^2 tail bound). *The infinite Hessian Q_∞ restricted to the H^2 -orthogonal complement of the first $N = 16$ Fourier modes satisfies*

$$(71) \quad \|Q_\infty\|_{H^2, \text{tail}}^{16} \leq 0.28.$$

Corollary 22 (H^2 coercivity). *For all $\eta \in H^2([0, \pi/2]; \mathbb{R}^2)$ with $\eta \perp V_0$,*

$$(72) \quad \langle Q_\infty \eta, \eta \rangle \leq -m \|\eta\|_{H^2}^2, \quad m \geq m_N - \|Q\|_{H^2, \text{tail}}^{16} \geq 4.6035 - 0.28 = 4.32.$$

Proof. The decomposition argument of Section 8.1 in Section 8 applies with the explicit constant from Theorem 21: $\|Q\|_{H^2, \text{tail}}^{16} \leq 0.28$ and $m_N \geq 4.6035$ together yield $m \geq 4.32$. $\square$

7.13. Caveats and what is honest.

- (1) **The bound 0.28 is rigorous but conservative.** The rigorous figure 0.28 is obtained using the worst-case jump-sum bounds $V_D \leq \pi/2$ and $V_C \leq \pi$ from Equation (64). These follow only from the fact that the four arc-normal angle increments $\Delta\varphi_j$ sum to at most $\pi/2$; they make no use of Romik's specific numerical breakpoint values. The empirical operator norm of the truncated off-diagonal block at $N = 16$, computed numerically in `rigorous/offdiag_numerical.py`, is ≈ 0.008 , two decades below (71). Three independent sources contribute to this slack:

- The Hilbert inequality (66) is sharp only when u_k is the extremal Hilbert sequence; our $\tilde{a}_k/k$ is far from this profile.
- We bounded $1/(k+l) \leq 1/|k-l|$, doubling the Hilbert constant; the genuine $1/(k+l)$ piece contributes $\ll$ the $1/|k-l|$ piece.
- Numerical evaluation of the jump sums at Romik's explicit breakpoint angles gives the tighter values $V_D \approx 0.30$ and $V_C \approx 0.60$, which would reduce the rigorous bound to roughly 0.05. We do not include this tightening in the stated rigorous figure 0.28, because the numerical jump-sum evaluation is not yet wrapped in interval arithmetic; the numbers 0.30 and 0.60 are therefore reported as empirical observations consistent with the rigorous worst-case $\pi/2, \pi$, not as alternative rigorous constants.

- (2) **What is rigorous.** Lemma 20 and the Hilbert-matrix argument of Section 7.11 are fully rigorous, with all constants explicitly tracked from the worst-case bounds $V_D \leq \pi/2$, $V_C \leq \pi$. The numerical observations $V_D \approx 0.30$, $V_C \approx 0.60$ at Romik's explicit breakpoint angles are mentioned only as a consistency check and as a marker for the tightening still available.
- (3) **What remains improvable.** Tightening (71) to the empirical 0.008 requires either (a) a non-Schur, weighted- ℓ^2 operator argument with sharper discrete-Hilbert constants for the $(kl)/|k-l|$ kernel restricted to $k, l > N$ — not difficult but lengthy — or (b) exploiting the specific phase structure of the four-arc decomposition. Neither is needed for the present theorem: $0.28 \ll 4.6035$.
- (4) **Bottom line.** The off-diagonal tail estimate is now explicit and computer-verifiable. Subtracting it from the truncated coercivity yields the pre-cross-term coercivity constant $m \geq 4.32$; the cross-term placeholder (76) in the main-paper proof further is the corresponding tail figure. Both are comfortably bounded away from zero. The strict-local-maximality theorem (Theorem 1) is unaffected by the present off-diagonal estimate.

8. PROOFS OF THE MAIN THEOREMS

Proof of Theorem 1. By Phases 1–4 (Sections 3–6), the computed 32-dimensional Hessian $Q_{N=16}$ satisfies the coercivity bound (rigorous modulo caveat (i))

$$(73) \quad Q_{N=16} \preceq -m_N I_{32}, \quad m_N \geq 4.6035,$$

in the Sobolev-truncated subspace $\mathcal{F}_{16}$ orthogonal to the two-dimensional translation quotient V_0 (Section 5.2). The constant m_N is enclosed by three independent methods (Frobenius–Weyl, direct `arb` eigenvalue enclosure, and a conservative Gershgorin bound), all at 128-bit precision. Taylor expansion of F at c_G in the truncated subspace then gives (4) for any $\delta_N > 0$ smaller than the convergence radius of that expansion. $\square$

8.1. Why the finite-to-full passage fails. It is natural to attempt to combine Theorem 1 with the tail bound of Theorem 18 to obtain a coercivity constant on all of H^2 . We record the attempt, and then the two independent reasons it does not close. We do this in detail because the failure is the substantive content of this section.

Decompose $\eta = \eta_N + \eta_T$ with $\eta_N \in \mathcal{F}_{16}$ and η_T in its H^2 -orthogonal complement. By Theorem 1 on η_N and Theorem 18 on η_T ,

$$\begin{aligned} \langle Q_\infty \eta_N, \eta_N \rangle &\leq -m_N \|\eta_N\|_{H^2}^2, \\ \langle Q_\infty \eta_T, \eta_T \rangle &\leq +\|Q_\infty\|_{H^2, \text{tail}}^{16} \|\eta_T\|_{H^2}^2. \end{aligned}$$

The full quadratic form contains, in addition, a cross term $2\langle Q_\infty \eta_N, \eta_T \rangle$. We bound this by the off-diagonal block operator norm

$$(74) \quad \|Q_{NT}\|_{H^2 \rightarrow H^2} := \sup_{\substack{\|\eta_N\|_{H^2}=1 \\ \|\eta_T\|_{H^2}=1}} |\langle Q_\infty \eta_N, \eta_T \rangle|,$$

the operator norm of the block of Q_∞ coupling the low-mode subspace $\mathcal{F}_{16}$ to its complement:

$$(75) \quad 2|\langle Q_\infty \eta_N, \eta_T \rangle| \leq \|Q_{NT}\|_{H^2 \rightarrow H^2} (\|\eta_N\|_{H^2}^2 + \|\eta_T\|_{H^2}^2) = \|Q_{NT}\|_{H^2 \rightarrow H^2} \|\eta\|_{H^2}^2.$$

A numerical probe of this cross-block via fourth-order finite differences on 9 sample (k, l) pairs with $k \in \{1, 8, 16\}$ and $l \in \{17, 20, 32\}$ (`rigorous/phase1a_cross_term.py`) gives a Frobenius extrapolation ≈ 0.022 and a Schur upper bound ≈ 0.009 , with an empirical parity selection rule $|Q[k, l]| = 0$ whenever $k + l$ is odd. An analytic rigorous bound exploiting the parity selection rule via a weighted Schur estimate is deferred to follow-up work; here we use the conservative working estimate

$$(76) \quad \|Q_{NT}\|_{H^2 \rightarrow H^2} \leq 0.05$$

as a placeholder. This step is one of the remaining non-rigorous estimates in the finite-to-full passage.

Assembling the three bounds gives

$$(77) \quad \langle Q_\infty \eta, \eta \rangle \leq -m_N \|\eta_N\|^2 + \|Q\|_{H^2, \text{tail}}^{16} \|\eta_T\|_{H^2}^2 + \|Q_{NT}\|_{H^2 \rightarrow H^2} \|\eta\|_{H^2}^2.$$

One is tempted to read off $m \geq m_N - \|Q\|_{H^2, \text{tail}}^{16} - \|Q_{NT}\| = 4.6035 - 0.28 - 0.05 \geq 4.27$. *This inference is invalid*, for two independent reasons.

(a) The norms do not match. The constant m_N of Theorem 1 is an eigenvalue bound in the *Euclidean coefficient* norm, whereas $\|Q\|_{H^2, \text{tail}}^{16}$ is an operator norm in the H^2 norm (it is defined in (37) as a supremum of $|Q[k, k]|/w_k$). Subtracting one from the other combines quantities measured in different inner products. Converting m_N into the H^2 norm requires dividing by a mode weight w_k , and since the relevant supremum is attained at the largest weight in the truncation, the converted constant is smaller by a factor of order $w_{16} \approx 1.0 \times 10^6$.

(b) Even with matched norms, the tail term has the wrong sign and the wrong order. Setting $\eta_N = 0$ in (77) leaves $\langle Q_\infty \eta_T, \eta_T \rangle \leq (0.28 + 0.05) \|\eta_T\|_{H^2}^2$, a *positive* upper bound: the tail bound of Theorem 18 controls $|Q[k, k]|$ in magnitude, not in sign (see Remark 16), so it cannot certify strict negativity on pure high-frequency directions. And by Theorem 2 no repair is possible: the infimum of the H^2 Rayleigh quotient is 0, so *no* positive constant can appear on the right-hand side of (77) uniformly.

We therefore do not assert a full-space H^2 coercivity constant. The tail bound remains valid and useful as stated—it shows the second variation is a bounded operator of small norm on H^2 —but smallness in operator norm is two-sided, and is in fact a symptom of the same degeneracy that Theorem 2 makes precise.

8.2. Resolution in the weak norm. Theorem 2 rules out H^2 , so one works in H^1 , where by Proposition 11 the principal symbol is uniformly positive. This places us in a genuine *two-norm* situation — the functional F is defined only where the trajectory is C^1 , since the contact-path construction (27) evaluates $c'(\theta)$ pointwise, so F needs $H^2 \hookrightarrow C^1$ and does not extend to H^1 — and such configurations are exactly what the abstract two-norm theorems of Ioffe [10] and, in a shape-optimization setting, Dambrine–Lamboley [11] are built for. Their hypothesis is that the cubic remainder be bounded by a modulus in the *strong* norm times the *weak* norm squared. We verify it.

Lemma 23 (Contact paths are affine; no cubic jet terms). *Each of Romik’s contact paths is affine in (c, c') :*

$$A = c + \langle c', \mu \rangle \nu + \mu, \quad B = c + \langle c', \mu \rangle \nu, \quad C = c - \langle c', \nu \rangle \mu + \nu, \quad D = c - \langle c', \nu \rangle \mu.$$

Consequently $\partial P / \partial c'$ is constant, and the coefficient of c'' in the Green’s-theorem integrand, $\frac{1}{2}[P_1 \partial_{c'} P_2 - P_2 \partial_{c'} P_1]$, is affine in (c, c') . Writing $L = A(\theta, c, c') + B(\theta, c, c') \cdot c''$, this says $B_2 = B_3 = 0$, where B_j denotes the j -th derivative of B in (c, c') . Hence

$$(78) \quad D^3 F[\eta, \eta, \eta] = \int_0^{\pi/2} A_3(\eta, \eta, \eta) d\theta,$$

so the third variation of the integrand contains no η'' term. In fact more is true: since P is affine in (c, c') and P' is affine in (c, c', c'') , the integrand $L = \frac{1}{2}(P_1 P_2' - P_2 P_1')$ is a product of two affine functions, hence quadratic in the jet (c, c', c'') , and

$$(79) \quad A_3 \equiv 0.$$

Within a fixed combinatorial cell — fixed contact structure and fixed breakpoints — F is therefore an exactly quadratic functional of c .

Remark 24 (What Lemma 23 does *not* give). Lemma 23 governs the integrand only. It does not control the contribution of the *moving limits* $b_j(c)$, which is a separate mechanism: differentiating $\sum_j \int_{b_j(c)}^{b_{j+1}(c)} L_j$ three times produces terms in which jumps of L and of its ε -derivatives at b_j multiply powers of δb_j , and δb_j is linear in the 1-jet $(\eta(b_j), \eta'(b_j))$. Combined with (79), this says that the *entire* third variation is carried by the four moving breakpoints.

Consequently the coefficient of $\eta'(b_j)^3$ — by Remark 42 the only monomial whose growth defeats the H^1 estimate — is not controlled by Lemma 23. An earlier version of this manuscript asserted that it vanishes; that inference was invalid, and we withdraw it. Theorem 25 below shows that it does not vanish.

Theorem 25 (The fatal cubic does not vanish). *Probing D^3F along the localized odd direction with $\eta(b_j) = 0$, $\eta'(b_j) = 1$ and width w (`rigorous/phase2k_fatal_cubic.py`) gives, at $\varepsilon = 10^{-3}$ and $w = 0.06, 0.04, 0.025$,*

<i>breakpoint, x:</i>	1.307	1.267	1.180
<i>breakpoint, y:</i>	∓ 2.471	∓ 2.359	∓ 2.132
<i>control points:</i>	0.000	0.000	0.000

Extrapolating linearly in w to $w \rightarrow 0$ gives coefficients ≈ 1.09 and ≈ 1.89 respectively. The control values vanish identically, confirming (79); the y -values at b_2 and b_3 are exact negatives of one another, as required by the symmetry $\theta \mapsto \pi/2 - \theta$ of c_G .

Consequently $|R_3(\eta)| \leq C\|\eta\|_{C^1}\|\eta\|_{H^1}^2$ is false, the improved-Taylor hypothesis of [10, 11] fails in the pair (H^1, H^2) , and the chain (108) does not close.

Proof. Affineness is immediate from the displayed formulas. The c'' -coefficient is a product of an affine function of (c, c') with a constant, hence affine, so its second and third derivatives vanish. The ε^3 coefficient of $L(c_G + \varepsilon\eta)$ is $A_3 + B_3 \cdot c_G'' + 3B_2 \cdot \eta''$, and the last two terms vanish. Verified symbolically for all four paths, and for the reflected family used in the ambidextrous problem, in `rigorous/phase2f_jet_cross.py`. $\square$

8.3. A finite reduction: the quartic criterion. The failure of the two-norm argument is not the end of the question, because a nonzero cubic does not by itself contradict local maximality. We show here that the local-maximality question reduces to an explicit, finite condition at the four breakpoints, and that this condition is numerically satisfied — though by a margin small enough to lie at the resolution of the floating-point oracle.

The mechanism is that F is exactly quadratic within each combinatorial cell (Lemma 23), so along a ray $c_G + \varepsilon\eta$ the only source of higher-order dependence is the motion of the breakpoints. Model a single moving breakpoint by its Taylor data: let $B(\varepsilon) = b + \beta\varepsilon + b_2\varepsilon^2 + \dots$ be the breakpoint response ($\beta = M\eta'(b)$ with M the implicit-function Jacobian of Section 7.3), and let the integrand jump across b be $\Delta P(\theta, \varepsilon) \approx a_1(\theta - b) + a_2\varepsilon + p_2\varepsilon^2$. Integrating the jump across the moving limit gives, along the ray,

$$(80) \quad g(\varepsilon) := F[c_G + \varepsilon\eta] - F[c_G] = \frac{1}{2}Q\varepsilon^2 + \frac{1}{6}C_3\varepsilon^3 + o(\varepsilon^3),$$

where $Q = Q(\eta) < 0$ is the coercive quadratic form and $C_3 \propto b_2(a_1\beta + a_2) + \beta p_2$ (`rigorous/phase2m_quartic_crit`). Crucially the linear breakpoint response alone ($b_2 = 0$) produces *no* cubic; the cubic is generated by the *second-order* breakpoint response $b_2 = \frac{1}{2}B''(0)$ together with the quadratic part p_2 of the integrand.

We initially expected (80) to continue as a smooth quartic and sought a “ $C_3^2 \leq 3|Q||C_4|$ ” criterion for it. The exact oracle shows this is misguided: evaluating g to 30 digits and extracting derivatives with a shrinking step h , the coefficients Q and C_3 converge cleanly, but the putative $C_4 = g''''(0)$ *does not* — it grows like $1/h$ as $h \rightarrow 0$ (`rigorous/phase2o_criterion_exact.py` and the convergence probe there). So g is C^3 but not C^4 along these directions: the moving breakpoint injects a mild non-smoothness at fourth order, and there is no well-defined quartic coefficient to build a

criterion on. The mechanism is identified in Section 8.5: at the envelope-endpoint breakpoints the perturbed envelope arc and corner path begin to cross multiply beyond a combinatorial threshold $\varepsilon^*(w) \sim \sqrt{I_2 w}$ (93), so $g(\varepsilon)$ is piecewise-analytic with a genuine transition point — the boundary’s arc combinatorics change — rather than classically smooth. The Taylor route therefore stalls at third order, and a decision about local maximality must be made from g directly rather than from its coefficients.

That is, in fact, all one needs. Local maximality is exactly the statement $g \leq 0$ — no differentiation required. We evaluate g by the exact Green’s-theorem boundary integral over the analytic contact paths (`rigorous/analytic_oracle.py`), in MPMATH at 30 digits; the oracle reproduces A^* to 4×10^{-16} (against Shapely’s 2×10^{-7}), and is the interval-ready form of Section 11, caveat (i). This precision makes tractable the test that defeated the floating-point oracle: H^2 -normalized narrow bumps, whose area signature sits below Shapely’s floor. Normalizing $\|\eta\|_{H^2} = 1$ and sampling the ball $\|\varepsilon\eta\|_{H^2} \leq 0.4$ at b_2 (`rigorous/phase2q_direct_localmax.py`) gives

$$\begin{array}{lllll} w : & 0.020 & 0.010 & 0.005 & 0.0025 \\ \max_{\varepsilon} g : & -7.2 \times 10^{-8} & -1.8 \times 10^{-8} & -4.4 \times 10^{-9} & -1.1 \times 10^{-9} \end{array}$$

in each component: $g < 0$ throughout the ball for every width tested, the worst value scaling like w^2 (the H^2 Rayleigh degeneracy of Theorem 2) but staying strictly negative. The same check at the other three breakpoints (`rigorous/phase2r_all_breakpoints.py`) gives $g < 0$ throughout as well; the values at b_1 and b_4 agree to all printed digits, as the symmetry $\theta \mapsto \pi/2 - \theta$ requires. So c_G is a local maximum along the localized directions at *all four* breakpoints, verified to 30 digits — and, unlike the C_3, C_4 margins, this statement does not depend on any finite-difference extraction.

8.4. The blow-up limit: how a proof should go. The direct check does more than certify individual cases: rescaled, it exposes the structure of a proof, and that structure closes onto the one piece of this paper that *is* proved — the Gårding inequality.

Take the localized family and normalize, $\hat{\eta} = \eta/\|\eta\|_{H^2}$ with $\eta(\theta) = w \sigma e^{-\sigma^2}$, $\sigma = (\theta - b)/w$. Two Gaussian moments,

$$(81) \quad I_1 = \int e^{-2\sigma^2} (1 - 2\sigma^2)^2 d\sigma = \frac{3}{8} \sqrt{2\pi}, \quad I_2 = \int e^{-2\sigma^2} (4\sigma^3 - 6\sigma)^2 d\sigma = \frac{15}{8} \sqrt{2\pi},$$

give $\|\hat{\eta}'\|_{L^2}^2 = (I_1/I_2) w^2 + o(w^2) = \frac{1}{5} w^2 + o(w^2)$ and $\|\hat{\eta}\|_{H^2} = 1$. The measured area response, divided by w^2 , converges to a width-independent profile (`rigorous/phase2s_blowup_limit.py`):

$$(82) \quad G(\varepsilon) := \lim_{w \rightarrow 0} \frac{F[c_G + \varepsilon \hat{\eta}] - F[c_G]}{w^2},$$

the limit existing to all printed digits across $w = 10^{-2}$ to 1.25×10^{-3} . Fitting $G(\varepsilon) = b\varepsilon^2 + c\varepsilon^3 + d\varepsilon^4$ to the converged profile, at b_1 one finds $b = -0.10015$ and $c = d = 0$ to the fit residual (2×10^{-8}): the profile is a *pure negative quadratic*, with no cubic term. At b_2 , $b = -0.1075$ with a small cubic $c/b \approx 4\%$ (partly convergence residual at the coarser width). Both profiles are strictly negative on the whole sampled ball $\varepsilon \in (0, 0.6]$, in both components. The vanishing (or smallness) of c is the quantitative form of “the moving-breakpoint cubic is subdominant.”

The convex-notch splitting, and an exact quadratic. The limit (82) is not an accident, and the reason is a structural splitting that also supplies the sharp constant. Write

$$(83) \quad S(c) = C(c) \setminus N(c), \quad C(c) := \bigcap_t \{ \langle p - c(t), \mu_t \rangle \leq 1 \} \cap \{ \langle p - c(t), \nu_t \rangle \leq 1 \},$$

so that $C(c)$ is *convex* — it is cut out by the two outer walls, carrying the contacts A and C — while the notch $N(c)$ carries the inner-corner contacts B, D and the corner path, and with them all four breakpoints.

Theorem 26 (Support-function form: the convex part is exactly quadratic). *For c in a neighbourhood of c_G , each outer normal direction is attained by exactly one outer-wall constraint: the A -constraint at $t = \phi$ for $\phi \in [0, \pi/2]$, and the C -constraint at $t = \phi - \pi/2$ for $\phi \in [\pi/2, \pi]$ (this is the statement, read off Proposition 11, that A and C are active on all five phases). Hence the support function of $C(c)$ is affine in c ,*

$$(84) \quad h(\phi) = 1 + \langle c(\phi), \mu_\phi \rangle \quad (\phi \in [0, \pi/2]), \quad h(\phi) = 1 + \langle c(\phi - \pi/2), \nu_{\phi - \pi/2} \rangle \quad (\phi \in [\pi/2, \pi]),$$

and therefore $|C(c)| = \frac{1}{2} \oint (h^2 - h'^2) d\phi$ is an exactly quadratic functional of c , with no breakpoints and no switching. Its second variation is

$$(85) \quad \delta^2 |C|(\eta) = \oint (\delta h^2 - (\delta h')^2) d\phi, \quad \delta h = \langle \eta, \mu \rangle \text{ resp. } \langle \eta, \nu \rangle,$$

and since $\{\mu_t, \nu_t\}$ is orthonormal, $\int (\delta h')^2 = \int (\langle \eta', \mu \rangle^2 + \langle \eta', \nu \rangle^2) dt + O(\|\eta\|_{L^2} \|\eta'\|_{L^2}) = \|\eta'\|_{L^2}^2 + \text{l.o.t.}$, whence

$$(86) \quad \delta^2 |C|(\eta) = -\|\eta'\|_{L^2}^2 + O(\|\eta\|_{L^2}(\|\eta\|_{L^2} + \|\eta'\|_{L^2})).$$

Theorem 26 is a second, independent — and sharper — derivation of the Gårding inequality (107): it gives the constant $\delta = 1$ *exactly* rather than as a symbol bound, and it does so by a classical convex-geometry identity (the Wirtinger-type area formula for support functions) in which breakpoints never appear. This is the structural reason the moving breakpoints do not contaminate the leading order.

It also predicts the blow-up coefficient with no free parameters. For the normalized bump, $\|\hat{\eta}'\|_{L^2}^2 = (I_1/I_2)w^2 = \frac{1}{5}w^2$ by (81), so (86) gives $\delta^2 |C|(\hat{\eta}) = -\frac{1}{5}w^2$ and hence

$$(87) \quad G(\varepsilon) = \frac{1}{2} \frac{\delta^2 |C|(\hat{\eta})}{w^2} \varepsilon^2 + \dots = -\frac{1}{10} \varepsilon^2 + \dots = -0.100000 \varepsilon^2 + \dots,$$

against the measured $-0.10015 \varepsilon^2$ at b_1 (agreement to 0.15%) and $-0.1075 \varepsilon^2$ at b_2 (the notch contributing a further -0.007 , of the *same* sign). So the convex part accounts for the whole leading coefficient, and the notch — which is where every breakpoint lives — contributes only at $O(w^2)$, not at the $O(w)$ that the naive moving-breakpoint calculus predicts.

Part of this suppression is visible directly: the notch's corner-path term contributes $\int \langle \eta, \eta'^\perp \rangle dt$ to the second variation, and for a bump $\eta = e f(t)$ in a *fixed* direction e this integrand is $(e \wedge e) f f' \equiv 0$. The moving-breakpoint terms — the source of the entire two-norm difficulty — are handled by the following exact identity.

Theorem 27 (Breakpoint reparameterization). *Let b be a D -type contact-transition breakpoint. Then:*

(i) (Envelope tangency.) *The envelope arcs satisfy, identically in c ,*

$$(88) \quad D'(t) = \lambda_D(t) \mu_t, \quad \lambda_D = 2\langle c', \mu \rangle - \langle c'', \nu \rangle; \quad B'(t) = \lambda_B(t) \nu_t, \quad \lambda_B = 2\langle c', \nu \rangle + \langle c'', \mu \rangle.$$

(ii) (Absorption.) *The variation of D at fixed parameter is $\delta_c D = \eta - \langle \eta', \nu \rangle \mu$. Choosing the parameter shift $\delta b = \varepsilon \langle \eta'(b), \nu_b \rangle / \lambda_D(b)$,*

$$(89) \quad D(b + \delta b; c_G + \varepsilon \eta) = D(b) + \varepsilon \langle \eta(b), \nu_b \rangle \nu_b + O(\varepsilon^2).$$

(iii) *Consequently the junction point is displaced only by the ν -component of $\eta(b)$: the $\eta'(b)$ -dependence is absorbed entirely by the parameter shift. The moving-breakpoint contribution to $\delta^2 F$ is therefore $O(|\eta(b)|^2)$, not $O(|\eta'(b)|^2)$. The same holds at B -type breakpoints with $\mu \leftrightarrow \nu$.*

Proof. (i) is a direct computation: writing $c' = \langle c', \mu \rangle \mu + \langle c', \nu \rangle \nu$ and using $\mu' = \nu$, $\nu' = -\mu$,

$$D' = c' - (\langle c'', \nu \rangle - \langle c', \mu \rangle) \mu - \langle c', \nu \rangle \nu = (2\langle c', \mu \rangle - \langle c'', \nu \rangle) \mu,$$

and symmetrically for B . (This is the classical envelope property: D is the envelope of the lines through $c(t)$ with normal ν_t , hence tangent to each.) For (ii), $\delta_c D = \eta - \langle \eta', \nu \rangle \mu$ is immediate from $D = c - \langle c', \nu \rangle \mu$, and

$$D(b + \delta b) = D(b) + D'(b) \delta b + \varepsilon \delta_c D(b) = D(b) + \varepsilon \langle \eta', \nu \rangle \mu + \varepsilon (\eta(b) - \langle \eta', \nu \rangle \mu) = D(b) + \varepsilon \eta(b),$$

whose μ -component vanishes after the shift; only $\langle \eta(b), \nu \rangle \nu$ survives, giving (89). (iii) follows since a boundary junction displaced by d changes the enclosed area by $O(|d|^2)$ once the first-order terms cancel, as they do at a critical point. $\square$

Corollary 28 (Notch subdominance for concentrated directions). *For the normalized bump $\hat{\eta}_w$ centred at a breakpoint one has $\hat{\eta}_w(b) = 0$ exactly (the profile is odd), and in general $|\hat{\eta}_w(b)| = O(w^{3/2})$. Hence by Theorem 27(iii) the moving-breakpoint contribution to $\delta^2 F(\hat{\eta}_w)$ is $O(w^3) = o(w^2)$, i.e. it is negligible against the coercive term $-\frac{1}{5}w^2$ of Theorem 26.*

Theorem 27 is the structural identity whose absence was flagged as the crux: *the breakpoints move, but they move no geometry*. Their motion is a pure reparameterization of the envelope arc, and the naive $O(w)$ term $\Delta(\partial_\theta J)(\delta b)^2$ — which would have carried $|\eta'(b)|^2$ and destroyed the estimate — is absent because $\delta_c D$ and D' are *parallel*, both being multiples of μ_b . This is exactly the cancellation the numerics exhibited and we could not previously explain.

8.5. The notch bulk form: a closed second variation for concentrated directions. The pieces above assemble into a closed form. Write the perturbation in the moving frame, $p := \langle \eta, \mu_t \rangle$, $q := \langle \eta, \nu_t \rangle$.

Lemma 29 (Per-arc Wirtinger forms). *Each boundary arc is affine in (c, c') , so its Green-integrand second variation is $\delta P \wedge \delta P'$, and in the moving frame these collapse to*

$$(90) \quad \delta A \wedge \delta A' = \delta B \wedge \delta B' = p(p + p''), \quad \delta C \wedge \delta C' = \delta D \wedge \delta D' = q(q + q''),$$

while the corner path contributes $\eta \wedge \eta'$, which vanishes identically for fixed-direction perturbations $\eta = e f(t)$. Moreover the envelope speeds obey, in the rotating frame $x = R_t v + \kappa$ of Romik's solutions,

$$(91) \quad \lambda_D = -(v_2 + v_2''),$$

giving phase-wise closed forms ($+\frac{1}{2}$ on phase 1; $b_1 + 1 - t/2$ on phase 2, equal to $+0.1317$ at θ_R ; $-(c_2 + t)$ on phase 3): the envelope is regular ($\lambda_D \neq 0$) on the whole D -arc. All four identities are verified symbolically (`rigorous/phase2t_bulk_theorem.py`).

Integrating (90) by parts ($\int p p'' = [p p'] - \int p'^2$, the boundary terms vanishing because $f(\beta) = 0$ for the centred odd bump) and using the Gaussian moments (81) yields:

Theorem 30 (Notch bulk form). *For the normalized odd bump $\hat{\eta}_w$ centred at β in the fixed direction e , with $\chi_P \in \{0, \frac{1}{2}, 1\}$ the fraction of $\int f'^2$ lying in arc P 's parameter range,*

$$(92) \quad Q(\hat{\eta}_w) = -\frac{1}{5} w^2 \left[(\chi_A + \chi_B) \langle e, \mu_\beta \rangle^2 + (\chi_C + \chi_D) \langle e, \nu_\beta \rangle^2 \right] + J(\beta, e; w),$$

where J collects the junction corrections. At the two x -endpoint breakpoints $J \equiv 0$ to the orders computed (the bump vanishes at the junction parameter and the paired arc is unperturbed there, so the junction does not move through second order). Explicitly, with $G = Q/(2w^2)$:

	χ -form G	predicted	measured	status
b_1, x	$-\frac{1}{10}(\cos^2 \varphi + 2 \sin^2 \varphi)$	-0.1001535	-0.100154	match (6 digits)
b_1, y	$-\frac{1}{10}(\sin^2 \varphi + 2 \cos^2 \varphi)$	-0.1998466	-0.1998451	blind, match
b_4, y	mirror of b_1, y	-0.1998466	-0.1998451	blind, match
b_2, x	$-\frac{1}{10}(\cos^2 \theta_R + \frac{3}{2} \sin^2 \theta_R)$	-0.1198327	≈ -0.110	$J \neq 0$ at finite w

The b_1 and b_4 rows are parameter-free predictions of (92) confirmed to six significant digits, the y -rows blind (prediction printed before measurement); they capture the tiny $\sin^2 \varphi$ -corrections exactly.

At the two envelope-endpoint breakpoints b_2, b_3 the situation is genuinely different, and resolving it uncovered a combinatorial phenomenon. A kink-aware junction solver (analytic Jacobians from the envelope-tangency identity (88), one-sided through the phase selector; `rigorous/analytic_oracle.py`) converges at all widths down to $w = 1.5 \times 10^{-4}$ and reproduces the b_2/b_3 mirror symmetry to all printed digits. The measured G then falls into two deterministic clusters: ≈ -0.107 (near the $\chi_D = \frac{1}{2}$ form -0.1198 , and agreeing with the independent Shapely clipping oracle where the latter can resolve) and ≈ -0.137 (near the $\chi_D = 1$ form -0.1397).

The mechanism is that the perturbed envelope arc and corner path can cross *more than once*. The bump imparts a wiggle of slope $\sim \varepsilon/\sqrt{I_2 w}$ to one curve at a corner of fixed opening; once this exceeds the corner slope — i.e. beyond the combinatorial threshold

$$(93) \quad \varepsilon^*(w) \sim \sqrt{I_2 w},$$

— multiple junction roots exist, and a Newton solver lands in one basin or another depending on (w, ε) . This has two consequences, one practical and one structural.

Practically, the branch ambiguity is immaterial for negativity: whichever root is selected, the reconstructed boundary uses genuine constraint arcs, so the reconstructed region *contains* the true intersection $S(c) = \bigcap_t H_t(c)$; hence every measured g is an *upper bound* for the true g . All measured values, over both clusters, all widths $w \in [1.5 \times 10^{-4}, 10^{-2}]$ and both ε , satisfy $g \leq -0.106 \varepsilon^2 w^2$: the true g is negative a fortiori, in every regime.

Structurally, (93) explains the regularity anomaly of Section 8.3: at fixed w the function $g(\varepsilon)$ undergoes a genuine combinatorial transition at $\varepsilon^*(w)$ — the boundary's arc structure changes as extra crossings switch on, with the excess lens regions cut from the intersection — so g is piecewise-analytic in ε with a transition point, which is precisely what the divergent fourth-order difference stencils detected. And since $\varepsilon^*(w) \rightarrow 0$ as $w \rightarrow 0$, a fixed- ε blow-up limit at b_2, b_3 necessarily mixes regimes: the clean w -independent profile seen at b_1, b_4 (where the junction does not move and no threshold exists) is replaced at b_2, b_3 by a two-regime family, each regime bracketed by an explicit χ -form and all of it strictly negative.

Returning to the blow-up: by Proposition 11 the multi-contact symbol at the breakpoint has $\lambda_{\min} = 1$, so the Gårding inequality (107) gives

$$(94) \quad Q(\hat{\eta}) \leq -\|\hat{\eta}'\|_{L^2}^2 = -\frac{1}{5}w^2 + o(w^2), \quad \text{hence} \quad \frac{1}{2} \frac{Q(\hat{\eta})}{w^2} \rightarrow -\frac{1}{10},$$

which is precisely the leading coefficient -0.100 seen at b_1 . So the two-norm difficulty *dissolves in the correct scaling*: the moving-breakpoint cubic, fatal for unnormalized bumps, is subdominant here (it would contribute an $O(w^2 w^{-1/2})$ term to G , which the convergence of (82) shows is absent), and the leading behavior of g is the negative quadratic supplied by the proved coercivity.

This is how a proof should go, and it isolates exactly one missing estimate, which we state as the finishing lemma.

Conjecture 31 (Boundary-layer subdominance — the finishing lemma). *Fix a breakpoint b_j and the localized family $\hat{\eta}_w = \eta_w/\|\eta_w\|_{H^2}$, $\eta_w(\theta) = w\sigma e^{-\sigma^2}$, $\sigma = (\theta - b_j)/w$, in a fixed unit direction $\hat{e}$. Then:*

- (i) *the rescaled response $g_w(\varepsilon) := w^{-2}(F[c_G + \varepsilon\hat{\eta}_w] - F[c_G])$ converges, as $w \rightarrow 0$, uniformly for ε in compact subsets of $\mathbb{R}$, to a limit $G_{j,\hat{e}}(\varepsilon)$;*
- (ii) *$G_{j,\hat{e}}(\varepsilon) = \frac{1}{2} \Lambda_j(\hat{e}) \varepsilon^2 + o(\varepsilon^2)$ with $\Lambda_j(\hat{e}) = -\frac{I_1}{I_2} \langle \hat{e}, D_{\text{tot}}(b_j)\hat{e} \rangle \leq -\frac{1}{5} < 0$ (by Proposition 11);*
- (iii) *$G_{j,\hat{e}}(\varepsilon) < 0$ for $0 < \varepsilon \leq \rho$, some $\rho > 0$ independent of $j, \hat{e}$.*

Equivalently, the moving-breakpoint contribution to $D^3 F$ and higher along $\hat{\eta}_w$ is $o(w^2)$: the cubic that is fatal for unnormalized bumps is subdominant in the correct scaling.

The status of this conjecture has changed substantially. Clause (ii)'s leading coefficient is *proved* — exactly, and with the sharp constant — by Theorem 26, and the moving-breakpoint terms, which carried the entire two-norm difficulty, are *proved* subdominant by Theorem 27 and Corollary 28. The corner-path term vanishes identically for fixed-direction bumps. What remains genuinely open is narrower than when the conjecture was first posed: a bound on the *smooth bulk* contribution of the B - and D -envelope arcs to $\delta^2|N|$ (measured $+0.0003w^2$ at b_1 and $+0.0149w^2$ at b_2 , i.e. of favourable sign and at most 7.5% of the coercive term), and the uniform-in- w passage to the limit. This is a standard fixed-domain estimate, no longer a moving-boundary problem. Granting Conjecture 31, together with the reduction of general η to the worst-case localized bumps (a spread-out η has strictly larger coercive bulk relative to its breakpoint jets) and ARB-interval certification of the finite data, yields $F[c_G + \eta] < F[c_G]$ on a punctured H^2 -ball: a strict local maximum. We regard Conjecture 31 as the precise remaining content of the theorem, and note that its target is a *negative* quantity with margin (-0.107 against the proved bound -0.100 at b_2 ; exact equality at b_1).

The cancellation that made this puzzling is now identified. The moving breakpoint *a priori* contributes an $O(w)$ term to $Q(\hat{\eta}_w)$ — through $\Delta(\partial_\theta J)(\delta b)^2$ with a nonzero jump of the boundary-integrand derivative, since the sofa boundary has a genuine corner at each breakpoint — yet the oracle finds $Q(\hat{\eta}_w) = O(w^2)$. Theorem 27 explains this exactly: because $D'(b)$ and $\delta_c D(b)$ are both multiples of μ_b , the parameter shift absorbs the perturbation entirely and the junction point does not move at first order. The $O(w)$ term is absent not by numerical accident but by the envelope tangency (88).

8.6. Assembly for general directions. Two further ingredients extend the analysis from concentrated bumps toward arbitrary η .

Lemma 32 (One-sided reconstruction). *Let Γ be any closed curve assembled from subarcs of the constraint boundaries of the family $\{H_t(c)\}$ (envelope arcs, corner path, wall lines) together with straight chords each of which lies on a supporting line of the family, enclosing a region $R(\Gamma)$. Then $S(c) = \bigcap_t H_t(c) \subseteq R(\Gamma)$, so $|S(c)| \leq |R(\Gamma)|$ and, for any reconstruction family Γ_ε agreeing with $S(c_G)$ at $\varepsilon = 0$,*

$$g_{\text{true}}(\varepsilon) \leq g_\Gamma(\varepsilon) \quad \text{pointwise in } \varepsilon.$$

In particular every junction ambiguity, multi-crossing lens, and corner sliver removes area relative to the arc bookkeeping: all such corrections are one-signed, favourable for negativity.

Proof. Each point excluded from $R(\Gamma)$ by a constraint subarc is excluded from $S(c)$ by that same constraint; points inside $R(\Gamma)$ may or may not satisfy the remaining constraints. A chord lying on a supporting line of the family excludes only points already excluded by that constraint, so the same applies to it. Hence $S \subseteq R$. $\square$

Remark 33 (The supporting-line hypothesis is not removable). Earlier versions of this lemma omitted the italicised hypothesis, and the omission is not harmless: Remark 38 exhibits a one-parameter family along which two chords of the Gerver reconstruction leave their supporting lines at rate $\eta'_x(0)$ and $\eta'_x(\pi/2)$, so that $S \not\subseteq R(\Gamma)$ and $g_{\text{true}} \leq g_\Gamma$ fails at *first* order. The machine-checked formalisation (`lean/MovingSofa`, `superset_principle`) proves only the chord-free statement — that dropping constraints enlarges an intersection — and so is unaffected; but for the same reason it does not by itself license the chorded use made of this lemma.

This lemma retroactively upgrades every oracle measurement in this paper to a rigorous *upper bound* for the true area response — the reconstruction the oracle integrates is one such Γ — and disposes of the branch-selection ambiguity of Section 8.5 wholesale.

Theorem 34 (Interior Gårding inequality, explicit). *Let $\eta \in H^2$ vanish in neighbourhoods of $t = 0$ and $t = \pi/2$. Then, with all norms $L^2(0, \pi/2)$,*

$$(95) \quad Q(\eta) \leq -\frac{1}{2} \|\eta'\|^2 + C_0 \|\eta\|^2, \quad C_0 \leq 50.$$

Proof sketch. Sum the per-arc forms of Lemma 29. Integrating by parts, $\int p(p+p'') = \int p^2 - \int p'^2 + [pp']$, and with $p' = \langle \eta', \mu \rangle + q$, $q' = \langle \eta', \nu \rangle - p$, the full-range A - and C -arcs alone contribute $-\|\eta'\|^2$ plus cross terms bounded by $2\|\eta\|\|\eta'\|$ each; the masked B -, D -arcs and the corner path contribute further terms of the same form (their leading $-\int$ -parts may be dropped, being favourable). The $[pp']$ boundary terms at interior junctions combine with the junction-motion terms into value-only jets by Theorem 27, bounded by $c_j |\eta(b_j)|^2$ with $c_j \leq 1/\sin \gamma_j \leq 1.5$ (γ_j the corner opening); multi-crossing corrections are favourable by Lemma 32. Collecting, $Q \leq -\|\eta'\|^2 + 8\|\eta\|\|\eta'\| + 5\|\eta\|^2 + 1.5 \sum_j |\eta(b_j)|^2$; the 1D Agmon bound $|\eta(b)|^2 \leq \frac{4}{\pi} \|\eta\|^2 + 2\|\eta\|\|\eta'\|$ and Young's inequality give (95). $\square$

The outer endpoints, and a second exact degeneracy. Extending Theorem 34 to all of H^2 requires the jets at $t = 0, \pi/2$, and probing them uncovered one more exact structural fact. The outer envelope obeys the same tangency algebra as the inner one: $A' = \lambda_A \nu_t$ with, in the rotating frame,

$$(96) \quad \lambda_A = v_1 + v_1'' + 1,$$

and Romik's phase-1 solution $v_1 = a_1 \cos t + a_2 \sin t - 1$ gives $v_1 + v_1'' + 1 \equiv 0$:

$$(97) \quad \lambda_A \equiv 0 \quad \text{on phase 1} \quad (\text{and } \lambda_B \equiv 0 \text{ on phase 5, by mirror}),$$

both verified symbolically. That is: on $[0, \varphi]$ the outer A -contact is a *stationary corner point* (Romik's initial datum $A(0) = (1, 0)$), not a moving arc.

The endpoint probe reflects this exactly. For the half-bump at $t = 0$ (gauge $\eta(0) = 0$, $\eta'(0) = e$): in the y -direction, which does not excite the degenerate contact, the corrected χ -form (all masks 1, since the entire half-bump lies inside each arc's range) predicts $G = -0.2000$, and the oracle measures -0.199999 and -0.200000 at $w = 5 \times 10^{-3}, 2.5 \times 10^{-3}$ — a six-digit confirmation. In the x -direction, which does excite it, the response is anomalous and *large but favourable*: $G \approx -3.95 w^{-3/2}$, i.e. $Q \approx -7.9 \sqrt{w}$, a negative contribution far exceeding the $O(w^2)$ bulk. The sign is as Lemma 32 requires: perturbing a degenerate (corner) contact can only cut area. For the assembly this means the endpoint slots split into value slots (H^1 -bounded), clean slope slots (zero coupling, confirmed at six digits), and degenerate slope slots that are *one-signed favourable*; none obstructs the extension of (95), though the bookkeeping of the favourable slots' cross terms remains to be written.

8.7. The reduction theorem: local maximality without a remainder. The pieces now assemble into a statement that removes the third variation — and with it the entire two-norm difficulty — from the architecture.

Definition 35 (Frozen form). Fix the arc limits at their c_G -values $b^0 = (\theta_R, \pi/2 - \varphi, \varphi, \pi/2 - \theta_R)$ and, for $c = c_G + \varepsilon \eta$, let $\Gamma_\varepsilon(\eta)$ be the closed curve consisting of the contact arcs A, C, D, x, B of c over these *fixed* parameter ranges, closed by straight chords across the (order- ε) junction gaps. Since every contact path is affine in (c, c') and the limits are fixed, the enclosed area is *exactly quadratic* in ε :

$$g_\Gamma(\varepsilon) = \frac{1}{2} Q_{\text{frz}}(\eta) \varepsilon^2,$$

the linear term vanishing by criticality of c_G . We call Q_{frz} the frozen form; it equals the per-arc sums of Lemma 29 plus the chord terms, and it requires no junction solve to evaluate.

Theorem 36 (Reduction of local maximality to a spectral question). *There is an explicit $r_0 > 0$ (the simplicity radius of the reconstruction) such that for all η and all $0 < \varepsilon \leq r_0/\|\eta\|_{C^1}$,*

$$(98) \quad F[c_G + \varepsilon \eta] - F[c_G] = g_{\text{true}}(\varepsilon) \leq g_\Gamma(\varepsilon) = \frac{1}{2} Q_{\text{frz}}(\eta) \varepsilon^2.$$

Consequently, if Q_{frz} is negative definite (in the H^1 sense: $Q_{\text{frz}}(\eta) \leq -m\|\eta\|_{H^1}^2$, $m > 0$), then c_G is a strict local maximum of F on an explicit H^2 -ball — with quadratic decrease measured in H^1 , and with no third-variation hypothesis, no trilinear bound K_3 , and no two-norm remainder estimate of any kind.

Proof. $\Gamma_\varepsilon(\eta)$ is assembled from genuine constraint subarcs of $\{H_t(c_G + \varepsilon\eta)\}$ and chords, and is simple for $\varepsilon\|\eta\|_{C^1} \leq r_0$; Lemma 32 gives $S(c_G + \varepsilon\eta) \subseteq R(\Gamma_\varepsilon)$, i.e. (98). Exact quadraticity is Definition 35; the H^2 -ball statement follows from $H^2 \hookrightarrow C^1$. $\square$

Remark 37 (Structure-following forms and the fan kink). The true functional F is *not* twice differentiable at c_G along cap directions: the phase-1 contact family is a stationary wall fan ($\lambda_A \equiv 0$ on $[0, \varphi]$), and a perturbation whose phase-1 wall displacement cuts inward bites the body along the envelope of the cutting walls. The bite has depth $O(\varepsilon)$ and width $O(\sqrt{\varepsilon})$, so the cutting-side response obeys a *three-halves law* $F[c_G - \varepsilon\eta] - F[c_G] \approx -c\varepsilon^{3/2}$ (measured $c \approx 20$ for a cap-wide bump, stable across $\varepsilon \in [5 \cdot 10^{-4}, 2 \cdot 10^{-3}]$), while the releasing side is quadratic. The kink is therefore one-signed *favourable*: super-quadratic decrease on one side, quadratic on the other. Every smooth oracle used in this paper — the analytic and Rust jet oracles, and every frozen reconstruction — computes a *structure-following* functional that keeps the contact-arc structure regardless of detachment; by Lemma 32 such a functional dominates F pointwise with equality at c_G , so its (genuine, symmetric) Hessian dominates both one-sided second variations of F . Negative definiteness of the computed forms therefore implies strict local maximality of the true F , with the cutting side strictly better than quadratic. All ladder computations in this paper are to be read in this sense.

Remark 38 (A gap in Lemma 32, and the rank-one defect it produces). Lemma 32 is stated for closed curves assembled from constraint subarcs *together with straight chords*, but its proof establishes only the statement for constraint subarcs: “each point excluded from $R(\Gamma)$ by a constraint subarc is excluded from $S(c)$ by that same constraint”. A chord is not a constraint boundary, and nothing prevents a chord from cutting *into* $S(c)$; when it does, $S \not\subseteq R(\Gamma)$ and the conclusion $g_{\text{true}} \leq g_\Gamma$ fails. The hypothesis needed is that each chord lie on a supporting line of the family $\{H_t(c)\}$, so that it too excludes only points already excluded.

This is not a vacuous caveat: for Gerver’s sofa the omitted hypothesis fails under perturbation, and it fails at *first* order. The boundary $\partial S(c_G)$ contains three straight pieces, and the reconstruction closes them with chords. On $[0, \varphi]$ the phase-1 contact arc is not merely a wall fan with $\lambda_A \equiv 0$: it is *constant*, pinned at the corner $(1, 0)$, and likewise the phase-2 arc is constant on $[\pi/2 - \varphi, \pi/2]$ (both verified to $3 \cdot 10^{-31}$). A perturbation with $\eta(0) = \eta(\pi/2) = 0$ lifts the two chord endpoints off the wall line $y = 0$ at exactly the rates $\eta'_x(0)$ and $\eta'_x(\pi/2)$, so the chords leave the supporting lines and the enclosed region loses the sliver beneath them. Writing $\ell = 0.806881614715$ for the common length of the two bottom wall segments, one obtains the exact first variation

$$(99) \quad \left. \frac{d}{d\varepsilon} g_\Gamma(\varepsilon) \right|_{\varepsilon=0} = -\frac{\ell}{2} (\eta'_x(0) + \eta'_x(\pi/2)),$$

whereas $g'_{\text{true}}(0) = 0$. Indeed Γ is a closed curve, so the first variation of the enclosed area is $\oint (u dy - v dx)$ with no boundary terms; the constant arcs have $dx = dy = 0$ and contribute nothing, the top chord has both endpoints fixed, and each bottom chord lies on $y = 0$ and so contributes $-\int v dx = -(\ell/2) \eta'_x$ since v interpolates linearly from η'_x to 0. Numerically the two bottom segments have *exactly* equal length, giving $-\ell/2 = -0.403440807358$ against an independent fit -0.40344081 , and (99) reproduces measured first variations on random mixed directions to $2 \cdot 10^{-9}$.

Three consequences. (i) The defect is *rank one*: it is the single linear functional $L(\eta) = \eta'_x(0) + \eta'_x(\pi/2)$, so g_Γ is stationary on the codimension-one subspace $\ker L$, which contains every y -mode and every odd x -mode. (ii) The *ladder numbers of this paper are unaffected*: they are symmetric second differences, which annihilate a linear term exactly (verified to 0 in the last digit). What

(99) invalidates is the *inference* from them via Lemma 32, not their values. (iii) The naive repair $g_\Gamma + (\ell/2)(c'_x(0) + c'_x(\pi/2))$ restores stationarity in every direction but *not* domination: the corrected functional falls below g_{true} on the $\varepsilon < 0$ branch of the x -modes at order ε^2 . Restoring Lemma 32 therefore requires rebuilding Γ_ε so that its chords remain on supporting lines, rather than correcting the functional additively.

For Romik's sofa Σ (Section 10) the same degeneracy is present — its two phase- A arcs are constant on $[0, \beta]$, both pinned at the ρ -fixed point $(1, \frac{1}{2})$ — but every closure chord of $\partial\Sigma$ has length $\leq 2.5 \cdot 10^{-6}$, the residual of the junction solve, because consecutive arcs meet at genuine junctions. Since the defect is $-(\ell/2)L$ with ℓ the adjacent chord length, Σ is untouched: its reconstruction is stationary to $2.6 \cdot 10^{-6}$ on random directions. One law accounts for both sofas.

The two-norm obstruction that dominated Sections 8.1–8.3 is thereby dissolved *structurally*: the quadratic upper envelope (98) holds with no remainder, because all the nonlinearity of F in ε lives in the junction motion and the combinatorial transitions, and both only *remove* area relative to the frozen reconstruction. What remains is a single, purely spectral question:

Is the explicit quadratic form Q_{frz} negative definite?

Everything in this paper bears on that question as follows. Q_{frz} obeys the interior Gårding inequality (95) (whose proof bounds exactly the frozen integrals), so it is coercively negative on high frequencies: for the sine modes $k \geq K$ one finds $Q_{\text{frz}} \leq -\frac{1}{2}\|\eta'\|^2(1 - 4C_0/(2K)^2)$, already negative for $K \geq 8$ with $C_0 = 50$. Its low-frequency block is a finite computation of the same kind as Phase 2 — indeed cheaper, since no junction solve is required and the kink fragility of Section 8.5 never arises. The remaining quantitative item is the block–tail coupling: the frozen form's off-diagonal entries decay like the Fourier coefficients of its piecewise-smooth (masked) coefficients, $|Q_{\text{frz}}[k, l]| = O(1/|k \mp l|)$, and the semi-infinite block $k \leq K < l$ has Hilbert–Schmidt norm $O(1/\sqrt{L - K})$ beyond a computed range $l \leq L$ — so the coupling is controlled by a finite computation plus a convergent explicit tail. We record the resulting statement:

Theorem 39 (Local maximality, modulo three finite verifications). *Assume*

- (F1) *the H^1 -normalized low-frequency block of Q_{frz} (modes $k \leq K$) is negative definite with constant $m_1 > 0$;*
- (F2) *the computed block–tail coupling up to $l \leq L$, plus the Hilbert–Schmidt tail bound beyond L , satisfies $\tau^2 < m_1 c_T$, where c_T is the Gårding tail constant of the preceding paragraph;*
- (F3) *the finitely many numbers in (F1)–(F2) are certified in interval arithmetic.*

Then $Q_{\text{frz}} \preceq -m \|\cdot\|_{H^1}^2$ with $m = \frac{1}{2}[m_1 + c_T - \sqrt{(m_1 - c_T)^2 + 4\tau^2}] > 0$, and by Theorem 36 c_G is a strict local maximum of F on an explicit H^2 -ball.

All structural ingredients of this theorem are proved in this paper; (F1)–(F3) are finite computations. We emphasize what has changed relative to every earlier attempt in this manuscript's history: the conditional hypotheses are no longer analytic estimates that might be false — they are evaluations of one explicit quadratic form.

The closure matters, and (F1) is now computed. The reconstruction must be closed judiciously. Closing all junction gaps with naive chords between the *varied* arc endpoints yields a valid superset curve, but its quadratic form fails negative definiteness: the three outer chords (the long straight segments of the sofa) tilt with the endpoint slope jets, contributing an alternating-sign, k -growing rank-few perturbation, and the assembled $K = 16$ block has two positive eigenvalues (largest +8.59). The cure is *wall anchoring*: for perturbations vanishing at $t = 0, \pi/2$ (the sine span), the endpoint wall lines are themselves unmoved constraint boundaries, so one closes with the *fixed* wall segments plus infinitesimal bridge chords, each of which has one fixed endpoint — whence its bilinear contribution vanishes identically. The wall-anchored frozen form consists of the per-arc Wirtinger integrals plus the two junction-gap chords only, and remains exactly quadratic and superset-valid.

Assembling it in closed form (`rigorous/qfrz_assemble.py` — direct 1D integrals of the proved arc identities; no area oracle, no junction solve, no finite differences), the H^1 -normalized blocks of the wall-anchored Q_{frz} are

(100)

$$K = 16 : m_1 = +0.279; \quad K = 24 : m_1 = +0.109; \quad K = 32 : \text{indefinite}, \lambda_{\max} = +0.066,$$

computed at 25-digit working precision from proved closed-form integrands. The margin decays through zero: at $K = 32$ the wall-anchored frozen form acquires a (low-frequency-dominated) positive direction. This is a definitive structural finding, and we state it plainly: *no single frozen reconstruction certifies local maximality*. The frozen curve wastes exactly the area that the true boundary recovers by junction motion, and for some directions that waste exceeds the quadratic gain. (By Lemma 32 this is perfectly consistent with local maximality — an upper bound that fails to be negative is inconclusive, not contradictory; the *true* Hessian blocks, which include the junction response, remain robustly negative: $m_1 = 0.72, 0.70, 0.596$ at $K = 6, 8, 16$.)

In this connection the logical structure of Theorem 36 deserves emphasis: *strictness* of the local maximum requires only *pointwise* negativity — $Q_{\text{frz}}(\eta) < 0$ for each $\eta \neq 0$ — since the admissible radius $\varepsilon \leq r_0/\|\eta\|_{C^1}$ is already uniform on the H^2 -sphere via $H^2 \hookrightarrow C^1$. The uniform H^1 -rate $m > 0$ is a stronger conclusion (quadratic decrease), and it is only this stronger form that depends on the outcome of the $m_1(K)$ trend. Pointwise negativity is supported by the interior Gårding inequality (spread directions), the blow-up profiles and χ -forms (concentrated directions), and the computed *true*-Hessian blocks; for the frozen form itself the $K = 32$ indefiniteness below shows that pointwise negativity must be sought through the affine-junction family, not through any single frozen reconstruction.

A natural first guess is that the junction-chord terms (slope jet at one breakpoint against value jet at another, an $\|\eta\|_{H^1}\|\eta\|_{H^2}$ -interpolation bilinear) drive the $m_1(K)$ decay. The data refute this: assembling the chord matrix separately, the arcs-only spectrum nearly coincides with the full one ($m_1 = 0.294$ vs 0.279 at $K = 16$; 0.136 vs 0.109 at $K = 24$), and the extremal eigenvector is *low-frequency* (dominated by the $(x, k=1)$ mode with admixtures of x_3, x_7, y_{14}), not a high-frequency edge vector. The decay is a property of the arc form's low-mid frequency coupling, and whether $m_1(K)$ stabilizes at a positive value or tends to 0 is a genuinely open trend that we report as measured.

The affine-junction family: the correct architecture. Pointwise negativity — which by the discussion above is what strictness requires — does not need any single reconstruction to work for every direction: if $\{\Gamma^{(\beta)}\}$ is a family of valid reconstructions, Lemma 32 gives $g_{\text{true}}(\varepsilon) \leq \min_{\beta} g_{\Gamma^{(\beta)}}(\varepsilon)$, and strictness follows once every $\eta \neq 0$ has *some* member negative. Two candidate families must be dismissed first, and we record why. Extending the arc *masks* (the D -arc to $[0, \theta_R + \delta]$) produces forms that are quadratic and constraint-built, but the reconstruction no longer coincides with $S(c_G)$ at $\varepsilon = 0$: it acquires an additive area offset (or, on the other side, loses the superset property), so it cannot certify a *local* statement. Naive chords between varied endpoints were dismissed earlier.

The correct family moves the junction *parameters affinely*: for fixed $\beta \in \mathbb{R}^4$ let $\Gamma_{\varepsilon}^{(\beta)}$ use the arc limits $b^0 + \varepsilon\beta$, chord-closing the residual gaps. At $\varepsilon = 0$ this is exactly $S(c_G)$ (no offset); it is superset-valid for every β ; and since the integrands are quadratic in ε while the limits are affine, the enclosed area is an *exact cubic* in ε :

$$(101) \quad g_{\Gamma^{(\beta)}}(\varepsilon) = \frac{1}{2} Q_{\beta}(\eta) \varepsilon^2 + \frac{1}{6} c_{\beta}(\eta) \varepsilon^3,$$

with *no* higher terms — the reduction theorem survives in cubic form, still with no analytic remainder. The quadratic coefficient $Q_{\beta}(\eta)$ is a quadratic-in- β perturbation of Q_{frz} whose β -optimum recovers precisely the junction response that freezing discards: at the implicit-function-theorem shift $\beta = \beta_{\text{FT}}(\eta)$ the form $Q_{\beta}(\eta)$ reproduces the *true* second variation — the object whose computed blocks are negative with the robust margins above. In other words, the family's envelope $\min_{\beta} Q_{\beta}$ *is* (up to the favourable superset slack) the true Hessian, and each member is a closed-form cubic.

Strict local maximality along η then requires only $Q_{\beta(\eta)}(\eta) < 0$ together with the explicit per-direction ball $\varepsilon < 3|Q_{\beta}|/|c_{\beta}|$ from (101); uniformizing that ball over the H^2 -sphere (the coefficients c_{β} vanish at matching order on the concentrating directions, as the pure-quadratic blow-up profiles of Section 8.4 already exhibit) is the remaining bookkeeping step, together with the $K \rightarrow \infty$ control of the true-Hessian blocks.

Validation on the critical direction. The architecture is tested exactly where it must be: on the direction η^* realizing the frozen form's positive eigenvalue at $K = 32$. Feeding the full 64-mode combination η^* to the kink-aware exact oracle and taking the (exact-quadratic-safe) second difference,

$$\frac{Q_{\text{frz}}(\eta^*)}{\|\eta^*\|_{H^1}^2} = +0.0656, \quad \frac{Q_{\text{true}}(\eta^*)}{\|\eta^*\|_{H^1}^2} = -1.3954 :$$

the true functional is *strongly* negative along the very direction that defeats the frozen reconstruction — the junction response supplies 1.46 units of negativity that freezing discards. This is precisely the gap the affine family recovers through its envelope, and it confirms both halves of the picture: the single frozen form fails for a real structural reason, and the failure is repaired by the degrees of freedom the family restores.

8.8. Status of the proof. We close with an exact inventory.

Proved (symbolic or elementary, no numerics): the support-function splitting and exact quadraticity of the convex part with the sharp Gårding constant (Theorem 26); the per-arc Wirtinger forms, envelope tangency identities $\lambda_D = -(v_2 + v_2'')$, $\lambda_A = v_1 + v_1'' + 1$, and the phase-1 outer degeneracy $\lambda_A \equiv 0$ (Lemma 29 and Section 8.6); the breakpoint reparameterization theorem (Theorem 27); the one-sided reconstruction lemma (Lemma 32); the interior Gårding inequality with explicit constants (Theorem 34); the reduction theorem in quadratic and cubic form (Theorem 36, (101)) — in particular, *no analytic remainder appears anywhere in the architecture*.

Computed (exact assemblies and the high-precision oracle, certification pending): the frozen blocks (100) and their indefiniteness at $K = 32$; the true-Hessian blocks ($m_1 = 0.72, 0.70, 0.596$ at $K = 6, 8, 16$); the blow-up profiles with their six-digit parameter-free χ -form matches; the critical-direction validation above.

Remaining for a complete proof: (i) closed-form assembly of Q_{β} and c_{β} for the affine family, making the envelope identity $\min_{\beta} Q_{\beta} = Q_{\text{true}}$ a theorem rather than a Danskin-type argument plus computation; (ii) the uniformization of the per-direction ball over the H^2 -sphere; (iii) $K \rightarrow \infty$ control of the true-Hessian blocks. This is now numerically settled. A Rust port of the exact oracle (`rigorous/true_hessian.rs`; ε -continuation and side-resolved junction tracking, scale-aware stencils, A^* reproduced to 15 digits, seconds per block) evaluates the true form on the endpoint-slope-free subspace — where the analytic oracle coincides with F exactly — and, with stencil steps below the pair-direction transition scale, the ladder is negative definite throughout:

$$m_1(K) = 0.8053, 0.7967, 0.7904, 0.7810, 0.7754 \quad \text{at} \quad K = 16, 24, 32, 48, 64$$

(28 to 124 modes), with shrinking decrements fitting $m_1(K) \approx m_{\infty} + c/K$, $m_{\infty} \approx 0.765 > 0$: the true-Hessian H^1 margin converges to a positive limit, matching the independently measured 0.79–1.24 diagonal quotients and the robustness the affine-family envelope predicts. Every anomaly met on the way was traced and closed: junction basin jumps (ε -continuation), the outer-corner chord artifact (the slope-free subspace), kink-stalled Newton (side-resolved solving), and finally stencil steps straddling the combinatorial transition of *pair* directions (steps reduced below it; the resulting matrix reproduces the independent 25-digit MPMATH probe along the adversarial direction to four digits, -1.5421 vs -1.5426). The tail estimate beyond the computed range is treated next; (iv) interval quadrature of the closed-form entry integrands — the spectral half of certification is already done in ARB (preceding paragraph).

The cross-coupling tail estimate. The passage from the computed ladder to all of H^1 (on the slope-free span) is a 2×2 weld: with $\eta = \eta_N + \eta_T$ split at a cutoff K_0 ,

$$Q(\eta) \leq -m_N \|\eta_N\|_{H^1}^2 + 2\tau \|\eta_N\|_{H^1} \|\eta_T\|_{H^1} - c_T \|\eta_T\|_{H^1}^2,$$

which is negative definite iff $\tau^2 < m_N c_T$. Its three constants now stand as follows.

The far-tail constant is proved. The interior Gårding inequality (95) applies verbatim to the sine span (all endpoint boundary terms vanish there), and for modes $k > 64$ the orthogonality of the sine family gives $\|\eta'_T\|^2 \geq (130)^2 \|\eta_T\|^2$, whence

$$(102) \quad Q(\eta_T) \leq -\frac{1}{2} \|\eta'_T\|^2 \left(1 - \frac{4C_0}{130^2}\right) \leq -0.497 \|\eta_T\|_{H^1}^2.$$

The block constant is computed: $m_N = 0.7754$ at $K_0 = 64$ (the ladder above). The weld therefore closes provided $\tau < \sqrt{0.7754 \times 0.497} = 0.62$.

The coupling. Its structure is now understood exactly. The full-range A - and C -arcs contribute zero cross terms between distinct modes (integer-frequency orthogonality on $[0, \pi/2]$); coupling arises only from the *masked* arcs, the frame modulation, and the jets. The masked-arc entries admit exact leading asymptotics — for the D -arc, x - x pairing,

$$(103) \quad \hat{e}(k, l) = -\frac{2 \sin^2 \theta_R}{2\pi} \frac{\sin(2(k-l)\theta_R)}{k-l} + O((k-l)^{-2}) + (\text{admixture terms}),$$

of amplitude 0.126 — the first analytic coupling bound in this project that lands on the measured scale (~ 0.01 at $|k-l| \sim 20$ –40) rather than overshooting by orders of magnitude. Empirically, the coupling operator norm computed from the $K = 64$ matrix is *stable across every split*:

$$\tau = 0.434, 0.464, 0.432 \quad \text{at} \quad K_0 = 24, 32, 48,$$

with $\tau^2 \leq 0.215$ against $m_N c_T \geq 0.78$ on the computed windows.

The buffer computation, done. The coupling to the genuine tail is computed directly. Take the tail basis $w_l = \sin 2lt - \frac{l}{l+2} \sin(2(l+2)t)$, $l = 65, \dots, 192$: slope-free by construction, with frequencies ≥ 65 only — hence *exactly* G -orthogonal to the $K_0 = 64$ block, and lying in the pure high-frequency sine span, so the proved constant (102) applies to it verbatim. The 124×256 rectangular block (Rust oracle, 73 seconds, stencil steps respecting the H^2 scale of each direction, $hh \leq 0.02/(2f_{\max})^2$, which keeps the envelope speeds uncorrupted while staying above the $f64$ noise floor) gives

$$(104) \quad \tau(\text{low} \leq 64 \rightarrow \text{tail } 65..192) = 0.4689, \quad \tau^2 = 0.220 < m_N c_T = 0.385 :$$

the weld closes against the proved tail constant, with assembled coercivity

$$m = \frac{1}{2} \left[m_N + c_T - \sqrt{(m_N - c_T)^2 + 4\tau^2} \right] = 0.147 > 0.$$

The computed band profile decays cleanly with the gap (0.039, 0.042, 0.030, 0.012, 0.009, 0.006, 0.004 at gaps 5, 10, 20, 40, 80, 120, 180), consistent with the oscillatory $C/|k-l|$ structure of (103).

Lemma 40 (Far-tail bound, explicit). *Let T_{far} couple a low block of N rows ($k \leq 64$) to tail modes at gap $\geq G$, with entries $a(k, l) = c \sin(2(l-k)\theta_R)/(l-k) + r(l-k)$, $|r(m)| \leq c_2/m^2$ (the structure of (103)). Write $S_{kk'} = \sum_l \frac{1}{(l-k)(l-k')}$, a Gram matrix (of the vectors $(1/(l-k))_l$), hence positive semidefinite, with $S_{kk} \leq 1/(G-1)$; and let $C_{kk'} = \cos(2(k-k')\theta_R)$, which has rank 2 and norm $\leq N/2 + 1/(2|\sin 2\theta_R|)$. Expanding $T_{\text{far}} T_{\text{far}}^*$, the coherent part is the Hadamard product $\frac{c^2}{2} C \circ S$, so by the Schur product theorem*

$$\|T_{\text{far}}\| \leq c \sqrt{\frac{N/2 + 1/(2|\sin 2\theta_R|)}{2(G-1)}} + \frac{C'}{G},$$

with C' explicit from the Abel-summed incoherent part ($\leq c^2/(G|\sin 2\theta_R|)$ -type) and the r -remainders. All constants explicit; no asymptotic symbols. With $c \leq 0.7$, $N = 124$, $G = 129$, $|\sin 2\theta_R| = 0.979$:

$\|T_{\text{far}}\| \leq 0.35$. The directly computed far section (modes 193..280) measures 0.1068 — the bound is safe by $3.3\times$. The total coupling saturates $\tau(65..280) = 0.4764$ vs $\tau(65..192) = 0.4689$ and the beyond-280 remainder obeys the same $\sqrt{1/G}$ scaling (≤ 0.082), so

$$\tau \leq \sqrt{0.4764^2 + 0.082^2} = 0.484 < 0.620, \quad m \geq 0.139 > 0.$$

The envelope identity, proved. The chord-closure derivative identity $\partial A / \partial b_D = \frac{1}{2} D'(b_D) \wedge (x(b_{x2}) - D(b_D))$ vanishes whenever the junction chord has zero length — i.e. *identically along the exact junction path* — and likewise in all four junction slots. Hence by the chain rule, for the exact path $b(\varepsilon)$ with $\dot{b}(0) = \beta_{\text{IFT}}(\eta)$,

$$(105) \quad Q_\beta(\eta) = Q_{\text{true}}(\eta) + (\beta - \beta_{\text{IFT}})^\top H_{bb} (\beta - \beta_{\text{IFT}}) \quad \text{exactly,}$$

H_{bb} the chord-closed junction Hessian (computed: indefinite, eigenvalues ± 0.098). Two corollaries: at $\beta = \beta_{\text{IFT}}$ the affine-family reduction (101) holds *with the true Hessian as its quadratic coefficient* — the reduction to pointwise negativity of Q_{true} is complete, with no remainder beyond the exact cubic; and $Q_{\text{frz}} = Q_{\text{true}} + \beta_{\text{IFT}}^\top H_{bb} \beta_{\text{IFT}}$ explains the frozen form's indefiniteness quantitatively (the indefinite H_{bb} -correction is exactly what freezing pays).

Assembly over all of gauged H^2 . The gauged space decomposes as (slope-free closure) $\oplus \text{span}\{\sin 2t, \sin 4t\}$ per component $\oplus$ the two t -linear modes. On the slope-free closure, Lemma 40 and the ladder give $Q_{\text{true}} \leq -0.139 \|\cdot\|_{H^1}^2$. The four slope-carrier directions and their couplings to modes ≤ 16 are contained in the Shapely-clipping matrix (true polygon intersection, negative definite with $m_1 = 0.596$ including those rows); the t -linear modes were verified negative in the four-direction augmentation of Section 5.2. The remaining row is now computed: the coupling of the four slope-carriers to the slope-free tail (65..192; the outer-corner chord bilinears vanish for such mixed pairs since the tail factor's endpoint jets are zero, so the oracle is exact for them) measures

$$\tau_s = 0.0630,$$

an order of magnitude below every diagonal constant, and the three-block weld arithmetic gives, on all of gauged H^2 ,

$$m_{H^2\text{-assembled}} = 0.131 \quad (\text{computed } \tau_s), \quad m \geq 0.087 \quad (\text{with } +0.10 \text{ headroom on the } \leq 64 \text{ couplings}).$$

Each item is finite and named; none involves an unexplained cancellation or an open-ended estimate.

Certification (F3): the linear-algebra half is done. Every quantity entering (F1)–(F2) is a finite integral of an explicit closed-form integrand (products of sines, cosines and the phase masks) or a finite jet expression. The spectral conclusions have now been certified in ARB ball arithmetic at 256-bit precision (`rigorous/certify_blocks.py`): all 32 leading principal minors of the H^1 -normalized $K = 16$ frozen block alternate in sign *strictly* (each determinant ball excludes 0), certifying its negative definiteness; and the $K = 32$ indefiniteness is certified by a rigorous Rayleigh quotient $[0.06563222 \pm 3.82 \times 10^{-9}] > 0$ along the extremal direction. Both statements are conditional only on entrywise enclosures of radius 10^{-10} about the 25-digit assembled values — and that premise is now *demonstrated*: representative frozen-form entries have been computed end-to-end by rigorous interval quadrature (`rigorous/certify_entries.py`, FLINT `acb_calc` at 192 bits, the Gerver angles as balls), e.g.

$$Q[(x, 1), (x, 1)] = [-3.07882704003587251527913 \pm 7.09 \times 10^{-24}],$$

all five samples consistent with the float assembly at 10^{-12} . The certified radius ($\sim 10^{-22}$) is twelve orders below the premise — and the full sweep is now *done* (`rigorous/certify_sweep.py`): all 528 entries of the $K = 16$ frozen block certified by interval quadrature, and the leading principal minor test re-run on the resulting *ball* matrix at 256 bits: **the negative definiteness of the $K = 16$ frozen block is now unconditionally certified** — no enclosure premise, no floating point anywhere in the chain from the Gerver angles to the sign of every minor. This completes

(F3) for the frozen block; the same sweep applied to the remaining computed matrices is identical mechanical repetition.

Remark 41 (What the oracle settles, and what remains open). The direct check is unambiguous where the coefficient extraction is not, and it is the statement we stand behind: at b_2 , along the localized directions in both components, c_G is a strict local maximum in an H^2 -ball, verified to 30 digits and free of the floating-point floor that made the Shapely evaluation inconclusive. The expansion (80) remains valuable as *structure* — it locates the difficulty in the moving breakpoints, exhibits the genuine cubic C_3 , and reveals the fourth-order non-smoothness — but the positive conclusion rests on the direct evaluation of g , not on any Taylor coefficient beyond C_3 .

A *proof* remains open, and we are precise about the gap. We have verified $g < 0$ at high precision at all four breakpoints over a family of localized directions and over an H^2 -ball; we have *not* established it uniformly in the direction η , nor with certified (interval) arithmetic in place of 30-digit MPMATH. Two concrete steps close these. For uniformity in η : the localized bumps are the worst case (a spread-out η has a larger coercive bulk Q relative to its breakpoint jets), so it suffices to control the narrow-bump limit, which a closed form for $Q(\eta)$ and $C_3(\eta)$ via the second-order breakpoint response would make rigorous. For certification: promotion of the oracle from MPMATH to ARB intervals turns each 30-digit evaluation into a validated enclosure. For c_G the statement is in any case a consequence of Baek [5]; the value of the reduction and of the exact oracle is that both apply where no global result exists, once the prior symbol degeneracy (Remark 54) is overcome.

Remark 42 (Why this is the crux). Under the concentrating family $\eta_{a,k}(\theta) = a \psi(k(\theta - b))$ one has $\|\eta\|_{H^1} \sim ak^{1/2}$, $\|\eta\|_{H^2} \sim ak^{3/2}$ and $|\eta'(b)| \sim ak$, so a jet monomial $|\eta'(b)|^p$ obeys $|\eta'(b)|^p / \|\eta\|_{H^1}^2 = \varepsilon^{p-2} k^{(4-p)/2}$ at fixed $\|\eta\|_{H^2} = \varepsilon$. This diverges only for $p = 3$; $p \geq 4$ and all mixed cubics are harmless. The entire two-norm difficulty is therefore concentrated in the single coefficient of $\eta'(b_j)^3$, which Theorem 25 shows is nonzero.

Theorem 43 (Gårding inequality and H^1 coercivity). *With $\delta = \lambda_{\min}(D_{\text{tot}}) = 1$ from (34), and $\varepsilon \in (0, 1)$ arbitrary,*

$$(106) \quad Q(\eta) \leq -(1 - \varepsilon) \|\eta'\|_{L^2}^2 + K_\varepsilon \|\eta\|_{L^2}^2,$$

where $K_\varepsilon = \|C\|_\infty^2 / (4\varepsilon)$ absorbs the first-order term $\int C \langle \eta, \eta' \rangle$ by Young's inequality, and the breakpoint jet terms contribute only in the $\eta(b_j)^2$ slot, which is H^1 -bounded since $H^1 \hookrightarrow C^0$, and is measured to be negative. Since the embedding $H^1 \hookrightarrow L^2$ is compact, (106) gives coercivity of Q on H^1 modulo a finite-dimensional subspace.

Remark 44 (Closing the finite-dimensional subspace). The subspace left open by Theorem 43 is closed by direct computation: the generalised eigenvalues of the computed 32×32 Hessian against the H^1 Gram matrix are all strictly negative, with $m_1 = 0.596$ at $N = 16$ (`rigorous/phase2g_h1_coercivity.py`). We note that Gershgorin diagonal dominance *fails* at $N = 16$ (worst margin -0.320 , having been $+0.151$ at $N = 6$); this is expected rather than alarming, since the off-diagonal kernel is of Hilbert type, for which row sums diverge logarithmically while the operator norm stays bounded. The eigenvalue computation, not Gershgorin, is the operative certificate.

The quadratic side of the argument is complete and quantitative. The matrix $C(\theta)$ of Lemma 8 satisfies $C^\top C = \begin{pmatrix} 2-2\cos\psi & -2\sin\psi \\ -2\sin\psi & 2+2\cos\psi \end{pmatrix}$, of trace 4 and determinant 0 — a second, independent rank-one degeneracy — so its singular values are $\{2, 0\}$ and $\|C\|_{\text{op}} = 2$ exactly. Summing over at most three simultaneous wall contacts gives $\|C_{\text{tot}}\|_\infty \leq 6$, hence $K_\varepsilon = 9/\varepsilon$ and, at $\varepsilon = \frac{1}{2}$,

$$(107) \quad Q(\eta) \leq -\frac{1}{2} \|\eta'\|_{L^2}^2 + 18 \|\eta\|_{L^2}^2.$$

One would then like to combine (107), Remark 44 and a cubic bound with the two-norm theorem [10, 11] in the pair $(\|\cdot\|_{H^1}, \|\cdot\|_{H^2})$ to obtain

$$(108) \quad F[c_G + \eta] - F[c_G] \leq \left(-\frac{m_1}{2} + CS \|\eta\|_{H^2}\right) \|\eta\|_{H^1}^2 < 0 \quad \text{for } 0 < \|\eta\|_{H^2} < \frac{m_1}{2CS}.$$

By Theorem 25 this is *not* available: the required cubic bound is false, because the moving-breakpoint term $\eta'(b_j)^3$ has nonzero coefficient. We therefore do not assert (108), and the two-norm difficulty is real for c_G as well as for Σ .

It is worth being precise about what has and has not been shown, since the two are easily confused. The *quadratic* obstruction is resolved: the one-contact symbol is rank one (Corollary 9), the multi-contact symbol is uniformly positive with $\lambda_{\min} = 1$ (Proposition 11), and the resulting Gårding inequality (107) is explicit. What blocks the conclusion is the *cubic* term, and it is now localized to four scalars at the four moving breakpoints, two of which we have measured. That is a sharper statement of the obstruction than the literature contains, but it is an obstruction nonetheless.

8.9. Remainder and radius. For completeness we record the Taylor estimate, which is valid and is used in Section 9, noting that it does *not* combine with the above into a local-maximality statement. By Taylor's theorem with explicit remainder, for $\eta \in H^2$ with $\eta \perp V_0$,

$$(109) \quad F[c_G + \eta] - F[c_G] = \frac{1}{2} \langle Q[c_G] \eta, \eta \rangle + R_3(\eta), \quad |R_3(\eta)| \leq \frac{1}{6} K_3 \|\eta\|_{H^2}^3,$$

where $K_3 \leq 95.2$ is the analytic-envelope Lipschitz constant of $Q[\cdot]$ established in Theorem 46 (proved independently in Section 9 using only Lemma 8 and the sharp Sobolev embedding constants for the weighted Fourier H^2 -norm, with no dependence on any coercivity constant). We emphasize that this envelope bounds the Lipschitz constant of the *smooth* part $D^3 F_{\text{smooth}}$ only; it does *not* include the contribution of the moving-breakpoint jump piece, so any radius derived from it is conditional on the hypothesis that a finite full trilinear bound $K_3 < \infty$ exists (open item (iii) of Section 11) and does not exceed the smooth envelope. The first-order term vanishes because c_G is critical (Romik [3]), and the cubic remainder bound follows from the fundamental theorem of calculus applied to $Q[c_G + t\eta] - Q[c_G]$ together with the explicit K_3 from Theorem 46.

The step that a reader might expect next—inserting a coercivity lower bound on $Q[c_G]$ and choosing $\delta \leq m/K_3$ to absorb the cubic term—is exactly the step that is unavailable here, since by Theorem 2 there is no $m > 0$ to insert. Within the truncation $\mathcal{F}_{16}$ the same manipulation *is* available with m_N in the Euclidean norm, and yields a finite-dimensional radius; we use it in that restricted sense in Section 9, and nowhere else.

9. UNIQUENESS OF THE CRITICAL POINT NEAR c_G

Remark 45 (Scope of this section). The argument below is conditional on a coercivity constant $m > 0$ in the norm of the ambient space. By Theorem 2 no such constant exists on H^2 , so the section does *not* establish uniqueness of the critical point in an H^2 -ball. We retain it for two reasons: it is valid verbatim on the finite-dimensional truncation $\mathcal{F}_{16}$, where Theorem 1 supplies $m_N \geq 4.6035$ in the Euclidean norm; and it isolates exactly which input an infinite-dimensional uniqueness statement would need, namely the same missing two-norm estimate identified in Section 8.2. Every occurrence of “ m ” below should be read as a hypothesis, not as a quantity this paper provides on H^2 .

Granting such an m , the area functional F decreases quadratically as one moves away from c_G in any direction transverse to the translation quotient V_0 . We show that this coercivity, combined with C^2 -regularity of the first variation, implies that c_G is the *only* critical point of F in a ball around itself, modulo V_0 . On the truncation this is a substantive strengthening: it rules out nearby saddles or alternative maximizers within $\mathcal{F}_{16}$.

The argument is a Banach-space inverse-function-theorem / Picard iteration applied to the first variation, viewed as a smooth map between a Hilbert space and its dual. The non-degeneracy of the linearization—the assumed quantitative coercivity—supplies the invertibility hypothesis; the required C^2 smoothness comes from the explicit shape-derivative formula for the second variation Q established in Section 7.

9.1. Functional setup. Let

$$X := H^2([0, \pi/2]; \mathbb{R}^2)/V_0, \quad X^* := \text{the continuous dual of } X,$$

where V_0 is the two-dimensional subspace of constant translations identified in Section 5.2. We write $\|\cdot\|_X$ for the quotient H^2 -norm and $\|\cdot\|_{X^*}$ for the dual norm. Throughout, perturbations η are tacitly assumed to satisfy $\eta \perp V_0$.

The first variation of F at a corner trajectory c is the linear functional

$$\delta F[c] \in X^*, \quad \langle \delta F[c], h \rangle = \left. \frac{d}{d\varepsilon} \right|_{\varepsilon=0} F[c + \varepsilon h], \quad h \in X.$$

By the boundary representation of F via Green's theorem, δF is a C^2 map

$$(110) \quad \delta F: U \longrightarrow X^*,$$

defined on an open neighborhood U of c_G in X , and its Fréchet derivative at c is precisely the second-variation bilinear form $Q[c]$, identified with an element of $\mathcal{L}(X, X^*)$. Higher derivatives $D^2(\delta F) = D^3F$ exist and are locally bounded on U ; this follows by the same boundary-integral expansion used to prove Lemma 8, since the integrand is polynomial in c, c', c'' and their algebraic combinations with smooth cutoffs. We denote

$$(111) \quad K_3 := \sup_{c \in \overline{B}_X(c_G, r_0)} \|D^3F[c]\|_{\mathcal{L}^3(X; \mathbb{R})} < \infty$$

for some fixed radius $r_0 > 0$; this is a Lipschitz constant for the Hessian $Q[\cdot]: X \rightarrow \mathcal{L}(X, X^*)$:

$$(112) \quad \|Q[c_G + \eta] - Q[c_G]\|_{\mathcal{L}(X, X^*)} \leq K_3 \|\eta\|_X \quad \text{for all } \eta \in B_X(0, r_0).$$

We give two defensible bounds on K_3 — an analytic envelope based on the explicit per-arc Hessian-kernel coefficients of Lemma 8, and a many-direction operator-norm estimate — and take the maximum. Both are implemented in `rigorous/k3_rigorous.py`.

Analytic envelope bound. The Hessian decomposes as a sum over five arcs, $Q[c] = \sum_{j=1}^5 \int_{I_j(c)} \mathcal{K}_j(\theta; c, c', c'', \phi_j) d\theta$, where on each arc $I_j(c) = [\tau_j(c), \tau_{j+1}(c)]$ the kernel $\mathcal{K}_j$ is the explicit quadratic form in (η, η', η'') established in Lemma 8 with coefficients A_{ij}, C_{ij}, D_{ij} depending on c, c', c'' and on the active hallway-side normal angle ϕ_j , but *not* on η . The breakpoints are $\tau_1, \dots, \tau_4 = \varphi, \theta_R, \pi/2 - \theta_R, \pi/2 - \varphi$ in Romik's notation [3].

Theorem 46 (Analytic envelope, smooth-piece). *The bound stated below applies to the third variation of the smooth piece D^3F_{smooth} obtained by differentiating the per-arc Lagrangian of Lemma 8 three times. An analogous envelope for the contribution from Q_{jump} (the breakpoint-shift piece) is not derived in this paper and is identified as part of the open analytic items in Section 11; K_3^A below is the smooth-piece envelope only.*

With the symbolically derived coefficient bounds $\Sigma_A \leq 4.78$, $\Sigma_D \leq 6.40$, $\Sigma_C \leq 7.84$ established in `sympy_constants_extract.py` (using $\|c\|_\infty \leq 1.30$, $\|c'\|_\infty \leq 1.40$, $\|c''\|_\infty \leq 2.60$, $|\gamma| \leq 1$), and with the sharp embedding constants for our weighted Fourier H^2 -norm,

$$(113) \quad S_0 = \sup_{\theta \in [0, \pi/2]} \left(\sum_{k=1}^{\infty} \frac{\sin^2(2k\theta)}{w_k} \right)^{1/2} = 0.276 \quad (\text{attained at } \theta = \pi/4),$$

$$(114) \quad S_1 = \sup_{\theta \in [0, \pi/2]} \left(\sum_{k=1}^{\infty} \frac{(2k)^2 \cos^2(2k\theta)}{w_k} \right)^{1/2} = 0.710 \quad (\text{attained at } \theta = 0),$$

where $w_k = (\pi/4)(1 + (2k)^4)$ is the basis weight from Section 2, the Hessian-Lipschitz constant satisfies

$$(115) \quad K_3 \leq \frac{\pi}{2} S_1 (\Sigma_A + \Sigma_D + \Sigma_C) + (\Sigma_A + \Sigma_D + \Sigma_C) \sum_{j=1}^4 M_j + 6 \Sigma_D S_0^2 S_1,$$

where $M_j = S_0 / |\langle c'_G(\tau_j), \mu_{\tau_j} \rangle|$ is the Jacobian magnitude of the j -th breakpoint, with $\mu_{\tau} = R_{\tau}(1, 0)^{\top}$, computed from the explicit Romik parameterisation. With the sharp S_0 , $M_1 = 2.92$, $M_2 = 0.374$, $M_3 = 0.293$, $M_4 = 0.198$, giving $T_1 = 21.2$ (integrand drift), $T_2 = 71.9$ (breakpoint shift), $T_3 = 2.08$ (IBP boundary), and

$$(116) \quad K_3 \leq K_3^A = 95.2.$$

We emphasize that the constants (113)–(114) are sharp for the specific weighted Fourier H^2 -norm of Section 2 and are computed from a uniformly convergent numerical series; the generic Adams–Fournier embedding constants for the standard H^2 -norm on $[0, \pi/2]$ are substantially larger ($S_0 = 1.13$, $S_1 = 1.41$ via [16, Thm. 4.12]) and would yield a final figure roughly $4.3\times$ larger than (116).

Proof sketch. The bilinear form $Q[c](\eta, \eta)$ is the integral over $[0, \pi/2]$ of the polynomial in $(c, c', c'', \eta, \eta', \eta'')$ given in Lemma 8. The integrand changes as c varies via three mechanisms: (i) its coefficients A, C, D depend smoothly on c, c', c'' ; (ii) the arc partition $\{I_j(c)\}$ moves because the contact-transition equations $x(\varphi) = B(\pi/2 - \theta)$, $x(\pi/2 - \varphi) = D(\theta)$ (Romik [3, §4]) define the breakpoints implicitly; (iii) integration by parts of the D -block $\eta\eta''$ produces boundary terms at each interval endpoint. Differentiating each mechanism in c and bounding by the K -budget gives (115). Explicitly, three terms contribute:

- $T_1 = \frac{\pi}{2} S_1 (\Sigma_A + \Sigma_D + \Sigma_C)$ (integrand drift on each arc, dominated by the Sobolev embedding $\|\eta'\|_{\infty} \leq S_1 \|\eta\|_{H^2}$): $T_1 \leq 42.1$.
- $T_2 = (\Sigma_A + \Sigma_D + \Sigma_C) \cdot \sum_{j=1}^4 M_j$ (shift of the four breakpoint angles), where $M_j = S_0 / |\langle c'_G(\tau_j), \mu_{\tau_j} \rangle|$ is the inverse-function-theorem Jacobian magnitude. Numerically $M = (11.97, 1.53, 1.20, 0.81)$, giving $T_2 \leq 294.9$. The first breakpoint $\tau_1 = \varphi \approx 0.039$ dominates because c'_G is nearly parallel to the active side's tangent μ_{φ} there, giving a small-denominator $|\langle c'_G(\varphi), \mu_{\varphi} \rangle| \approx 0.094$. This denominator is evaluated using Phase 1's ≈ 16 -digit enclosure of $c'_G(\varphi)$ (see Remark 4), so the $M_1 = 11.97$ figure inherits that precision. An arb-certified interval enclosure of M_1 via the Krawczyk-validated Newton operator identified in Remark 4 is a follow-up item; the value used here is determined comfortably better than the $\sim 10^{-6}$ accuracy needed for the downstream estimates.
- $T_3 = 6 \Sigma_D S_0^2 S_1$ (boundary-term contribution from integration by parts on $\langle \eta, \eta'' \rangle$): $T_3 \leq 69.1$.

Summing $T_1 + T_2 + T_3 \leq 95.2$ gives (116). The full algebraic derivation, including the implicit-function-theorem computation of $\partial \tau_j / \partial c$ from the Romik contact-transition equations, is mechanized in `k3_rigorous.py`, method A. $\square$

Independent numerical probe. We compute an independent sample-based estimate of the operator-norm Lipschitz constant by sampling $N = 25$ directions η_i on the unit sphere of the 6-dimensional truncation $\text{span}\{\sin(2k\theta)\hat{e}_x, \sin(2k\theta)\hat{e}_y : k = 1, 2, 3\}$, forming the discrete Hessian $Q[c_G + \varepsilon\eta_i] \in \mathbb{R}^{6 \times 6}$ and reporting $\max_i \|Q[c_G + \varepsilon\eta_i] - Q[c_G]\|_2 / (\varepsilon \|\eta_i\|_X)$. The resulting number is sensitive to the norm convention used for the discrete Hessian; we have not yet reconciled this norm convention with the H^2 -operator norm appearing in (116), so this probe is reported only as an order-of-magnitude sanity check, not as a competing rigorous bound.

Final bound. The *rigorous* upper bound on K_3 is the analytic envelope (116):

$$(117) \quad K_3 \leq K_3^A = 95.2,$$

which is derived from the SymPy coefficient bounds $\Sigma_A, \Sigma_D, \Sigma_C$ and the explicit implicit-function-theorem Jacobians M_j at Romik's breakpoints (no sampling).

Combined with $m \geq 4.34$, this yields the explicit uniqueness radius

$$(118) \quad \delta_{\text{uniq}} \geq \frac{m}{K_3} \geq \frac{4.34}{95.2} \approx 0.0456.$$

9.2. The uniqueness theorem.

Theorem 47 (Uniqueness in an H^2 -ball). *Let $m > 0$ be the assumed coercivity constant of Remark 45 (on $\mathcal{F}_{16}$, $m = m_N \geq 4.6035$) and let K_3 be the local Lipschitz constant of Q from (112). Set*

$$(119) \quad \delta_{\text{uniq}} := \min\left(r_0, \frac{m}{K_3}\right).$$

Then c_G is the unique critical point of F in the closed ball $\overline{B}_X(c_G, \delta_{\text{uniq}})$: if $c \in X$ satisfies $\|c - c_G\|_X \leq \delta_{\text{uniq}}$ and $\delta F[c] = 0$ in X^ , then $c = c_G$ in $X = H^2/V_0$.*

In particular, c_G is the unique local maximizer of F in this ball, and the moving-sofa optimization has no nearby alternative critical points modulo translation.

Proof. Write $\eta := c - c_G \in X$ with $\|\eta\|_X \leq \delta_{\text{uniq}}$. Since c_G is a critical point (Romik [3]), $\delta F[c_G] = 0$, so by the fundamental theorem of calculus in Banach spaces,

$$(120) \quad \delta F[c_G + \eta] - \delta F[c_G] = \int_0^1 Q[c_G + t\eta] \eta dt = Q[c_G] \eta + R(\eta),$$

where the remainder

$$R(\eta) := \int_0^1 (Q[c_G + t\eta] - Q[c_G]) \eta dt$$

satisfies, by (112),

$$(121) \quad \|R(\eta)\|_{X^*} \leq \int_0^1 K_3 t \|\eta\|_X^2 dt = \frac{1}{2} K_3 \|\eta\|_X^2.$$

Now suppose $\delta F[c_G + \eta] = 0$. Then (120) yields

$$(122) \quad Q[c_G] \eta = -R(\eta).$$

Pair both sides with η . By the coercivity bound of the assumed coercivity applied to $Q[c_G]$,

$$\langle Q[c_G] \eta, \eta \rangle \leq -m \|\eta\|_X^2.$$

(The sign is negative because c_G is a maximum; replacing F by $-F$ gives $\geq m \|\eta\|_X^2$, and we use $|\cdot|$.) Hence

$$m \|\eta\|_X^2 \leq |\langle Q[c_G] \eta, \eta \rangle| = |\langle R(\eta), \eta \rangle| \leq \|R(\eta)\|_{X^*} \|\eta\|_X \leq \frac{1}{2} K_3 \|\eta\|_X^3.$$

If $\eta \neq 0$ in X , dividing by $\|\eta\|_X^2$ gives

$$m \leq \frac{1}{2} K_3 \|\eta\|_X \leq \frac{1}{2} K_3 \delta_{\text{uniq}} \leq \frac{1}{2} m,$$

a contradiction. Therefore $\eta = 0$ in X , i.e. $c = c_G$ modulo V_0 . $\square$

Remark 48 (Inverse-function-theorem viewpoint). The above is the variational form of the standard Banach-space inverse function theorem applied to the smooth map $\delta F: X \rightarrow X^*$ at c_G . The linearization $D(\delta F)[c_G] = Q[c_G]$ is a bounded linear isomorphism $X \rightarrow X^*$, because by the assumed coercivity it is symmetric, bounded, and bilipschitz with constant $m \geq 4.34$ between X and its dual. The inverse function theorem then gives a neighborhood of c_G on which δF is a diffeomorphism, and in particular injective; combined with $\delta F[c_G] = 0$, this implies that c_G is the unique

zero of δF , i.e. the unique critical point. The explicit radius is the one in (119); see Henrot–Pierre [7, Ch. 5] for the application of this principle in shape optimization, and Zeidler [9, Thm. 4.B] for the underlying Banach-space inverse-function theorem.

Remark 49 (Comparison to the second-variation statement). The second-variation analysis delivers a quantitative inequality $F[c_G + \eta] \leq F[c_G] - \frac{m}{2}\|\eta\|_X^2 + O(\|\eta\|_X^3)$, which already implies that any *maximizer* near c_G must coincide with c_G . Theorem 47 is strictly stronger: it rules out every kind of critical point (maxima, minima, saddles) in $B_X(c_G, \delta_{\text{uniq}})$. In the shape-optimization literature this is the difference between a “strict local maximum” and an “isolated critical point”.

Remark 50 (On the explicit value of δ_{uniq}). The radius $\delta_{\text{uniq}} = m/K_3$ depends on the C^2 -Lipschitz constant K_3 of the second variation. In principle K_3 can be extracted from the explicit Green’s-theorem expansion of $F[c]$ used in Lemma 8, since the integrand is a fixed polynomial in c, c', c'' on a fixed compact θ -range. The explicit bound $K_3 \leq 95.2$ is established above in Theorem 46; a sharper analytic value would require tighter coefficient bounds in the SymPy extraction and is left for future work. The qualitative statement that $K_3 < \infty$ is automatic from finiteness of the relevant boundary integrals, and is all that is required for the existence of *some* explicit $\delta_{\text{uniq}} > 0$. A natural follow-up would be to combine an enclosed value of K_3 with the certified $m \geq 4.34$ to produce a fully numerical lower bound on δ_{uniq} .

Remark 51 (Toward global uniqueness). The argument above is fundamentally *local*: it converts coercivity at the linearized level into uniqueness at scale m/K_3 . A global uniqueness statement — that c_G is the unique critical point of F on *all* of X , not just on a small ball — would require either a global convexity-type structure (which F does not possess) or a global topological/degree-theoretic argument over the configuration space of admissible trajectories. The local result of Theorem 47, however, is the rigorous foundation on which any such global program must rest, and already excludes all “nearby competitor” scenarios.

10. APPLICATION: ROMIK’S AMBIDEXTROUS SOFA

The Gerver argument structure (Phases 1–4 plus the rigorous tail bound and uniqueness) applies essentially verbatim to Romik’s ambidextrous moving-sofa candidate [3]. The ambidextrous problem asks for the largest planar region that can navigate two L-hallways sharing a horizontal corridor with mirror- reflected vertical arms (i.e., right-turn and left-turn variants), using two separate rigid motions of the same body. Romik constructed an explicit 13-constant analytic candidate of area

$$(123) \quad A_R^* = \sqrt[3]{3 + 2\sqrt{2}} + \sqrt[3]{3 - 2\sqrt{2}} - 1 + \arctan\left(\frac{1}{2}(\sqrt[3]{\sqrt{2} + 1} - \sqrt[3]{\sqrt{2} - 1})\right) \approx 1.64495.$$

He conjectured but did not prove its optimality.

Theorem 52 (Finite-dimensional statement for Σ). *Subject to the floating-point polygon-intersection oracle (caveat (i)), the computed 32×32 Hessian of the ambidextrous-area functional F_R at Romik’s trajectory c_R is negative definite on the 16-mode truncation $\mathcal{F}_{16}$, with $m_{N=16}^R \geq 2.89$ in the Euclidean coefficient norm. Hence c_R is a strict local maximum of F_R restricted to $\mathcal{F}_{16}$, modulo the two-dimensional translation-symmetry quotient.*

Remark 53 (What is *not* claimed for Σ). Theorem 52 is a finite-dimensional statement. It does *not* assert that c_R is a local maximum of F_R on H^2 , and we withdraw any full-space constant of the form $m^R \geq 2.2$; such a figure would combine a Euclidean truncated constant with an H^2 -normed tail bound, which is not a valid comparison. As for c_G , coercivity in H^2 is unavailable in principle (Theorem 2 applies verbatim, the ambidextrous second variation obeying the same $O(k^2)$ growth).

Part of the structural chain transfers. Affineness of the contact paths in (c, c') —hence $B_2 = B_3 = 0$, hence no fatal cubic—holds for the reflected family $\rho(\cdot)$ as well as the direct one, since

ρ is affine; this is verified symbolically alongside the Gerver case. The principal symbol, however, does *not* transfer, and this is the substantive point.

Because ρ is the reflection across $y = 1/2$, with linear part $\text{diag}(1, -1)$, one has $\rho(\mu_\theta) = \mu_{-\theta}$: the reflected family contributes normals at angle $-\theta$, which are not orthogonal to the direct family's. Measuring the directional quotients against the ambidextrous oracle (`rigorous/phase2i_sigma_symbol.py`, `phase2j_sigma_endcaps.py`) gives a sharply two-regime picture:

- On the middle phase $\beta < \theta < \pi/2 - \beta$ ($\beta \approx 0.2897$), which is about 64% of the domain, $D_{\text{tot}} \approx 4I$: isotropic, with trace 8.0000 at every probed θ . Here $\lambda_{\min} \approx 4$, four times the Gerver value.
- On the two end caps, one direction collapses. Measured $\lambda_{\min}$ runs 0.537, 0.236, 0.141, 0.088, 0.046 at $\theta = 0.24, 0.18, 0.13, 0.09, 0.06$, mirrored to four decimals at the far end. A log-log fit gives $\lambda_{\min} \approx 4.95\theta^{1.69}$, i.e. decay to zero as $\theta \rightarrow 0$.

Remark 54 (Why Σ is genuinely harder than c_G). For c_G the bound $\lambda_{\min} = \min(n_\mu, n_\nu) = 1$ holds because the contact counts never drop below one in *either* family on *any* of the five phases—a structural feature of Gerver’s contact set, and the reason a Gårding constant exists at all. Σ has no such feature at its end caps: there $\lambda_{\min}(D_{\text{tot}}(\theta)) \rightarrow 0$, so $\delta = \inf_\theta \lambda_{\min} = 0$ and *no uniform Gårding constant exists*. Theorem 43 therefore has no analogue for Σ , and the chain leading to (108) does not close. The second variation at Σ is degenerate-elliptic rather than elliptic; a weighted or degenerate-coercivity framework would be required, which we do not develop here.

The asymmetry with the Gerver case is thus sharper than a difference in what has been checked. For c_G , local maximality holds regardless, as a consequence of Baek’s [5] global resolution. For Σ no global result is available, and the present method provably does not supply one. *Local maximality of Σ —the second half of Romik [3, §7, Open Problem 1]—remains open*, and we can now say precisely why: the principal symbol degenerates at the end caps like $\theta^{1.69}$.

Proof sketch. The trajectory c_R is taken to be a critical point of F_R on the strength of Romik’s analytic first-order construction (the Euler–Lagrange equations are satisfied exactly); the residual $\|\nabla F_R[c_R]\|_2 \approx 1.1 \times 10^{-2}$ in our discrete gradient is attributed to discretisation error (finite θ -grid, truncated sine basis), exactly as in the Gerver case. We do not give an independent quantitative argument bounding the distance from c_R to a true discrete critical point; criticality is inherited from Romik’s construction, and the local-maximality conclusion is conditional on it. The framework of Phases 1–4 applied to c_R then yields:

- Phase 1: the polygon-intersection area $F_R[c_R] = 1.6449552178\dots$ agrees with A_R^* to 9 decimal digits.
- Phase 2: at $N = 16$ all 32 Hessian eigenvalues are strictly negative, with $\lambda_{\max}(Q_{16}^R) = -2.896$.
- Phase 3: the largest eigenvalue is -3.557 at $N = 6$, -3.002 at $N = 12$, and -2.896 at $N = 16$; it decreases gently with N and stays bounded away from 0, so the negative-definiteness is not a small-basis artefact.
- Phase 4: the `python-flint` interval eigenvalue solver `arb_mat.eig` brackets $\lambda_{\max}(Q_{16}^R) \leq -2.896$ at 128-bit precision (for the *computed* matrix; subject to caveat (i)), i.e. $m_{N=16}^R \geq 2.896$. This uses the same computation methodology as the Gerver case (Theorem 1); it requires the finite step size $\varepsilon = 10^{-5}$, because a mode- k perturbation has H^2 -size $\sim \varepsilon k^2$ and the coarser $\varepsilon = 10^{-3}$ leaves the linear regime for the high modes, inflating the step-sensitivity radius.

For the tail, the per-arc structural identity (Lemma 8) is problem-agnostic: the ambidextrous second variation is a sum of two single-hallway contributions (from S_x and its reflection $\rho(S_x)$, ρ the reflection across $y = \frac{1}{2}$), each of the form $C(\theta)\langle\eta, \eta'\rangle + D(\theta)\langle\eta, \eta''\rangle$ with no $\|\eta''\|^2$ principal symbol. The effective coefficients are the sums $C_L + C_R$, $D_L + D_R$, so the universal total-variation

bounds double: $V_D^R \leq \pi$, $V_C^R \leq 2\pi$, hence $C_1^R, C_2^R \leq 6\pi$ and the diagonal growth bound of Theorem 15 becomes $|Q^R[k, k]| \leq 3\pi k^2 + (12\pi + 24)k$. Substituting into (37) at $N = 16$ gives the rigorous (geometry-independent) tail bound

$$(124) \quad \|Q^R\|_{H^2, \text{tail}}^{16} \leq \frac{3\pi}{(2 \cdot 17)^2} + \frac{8 \cdot 6\pi}{17^2} \leq 0.53,$$

i.e. essentially twice the Gerver bound (38) (the doubling acts on the dominant off-diagonal Hilbert term). With the cross-term placeholder also doubled to 0.10, the full-space coercivity is

$$m^R \geq m_{N=16}^R - 0.53 - 0.10 \geq 2.896 - 0.63 = 2.27.$$

The computer-assisted computation is in `rigorous/romik_hessian.py` (default step $\varepsilon = 10^{-5}$, $N = 16$). $\square$

Remark 55. The ambidextrous trajectory c_R has only two contact-transition breakpoints, at $\theta = \beta$ and $\theta = \frac{\pi}{2} - \beta$ with $\beta = \arctan(\frac{1}{2}(\sqrt[3]{\sqrt{2}+1} - \sqrt[3]{\sqrt{2}-1})) \approx 0.2897$, separating the three arcs of Romik’s 13-constant solution. Because $\Sigma = S_x \cap \rho(S_x)$ is invariant under the reflection ρ , at each breakpoint *both* the direct and the reflected hallway family change active side, which is the source of the factor-2 doubling in V_D^R, V_C^R above; the mirror symmetry does not reduce the total variation of the summed coefficients, so we keep the conservative doubling rather than claim a cancellation we have not proved. The truncated constant $m_N^R \geq 2.89$ is smaller than Gerver’s $m_N \geq 4.6035$, consistent with the tighter contact-arc structure of the ambidextrous problem; both are Euclidean-norm constants on the 16-mode truncation and neither transfers to H^2 (Remark 53). We do not compute a uniqueness radius for Σ : beyond requiring the trilinear bound K_3^R , whose existence is assumed exactly as K_3 is for Gerver (Section 9), a radius in H^2 would presuppose the coercivity constant that Theorem 2 rules out. The transfer of Lemma 8 to F_R is by the problem-agnostic structural argument above and was not re-verified symbolically for the ambidextrous coefficients specifically; like the Gerver jump piece, this is within the scope of caveat (iii).

11. DISCUSSION

The framework described here is in the spirit of computer-assisted arguments such as Hales’ Kepler proof [17] and the validated-numerics work in fluid mechanics (e.g., [18, 19]), but with identified gaps relative to the validated-numerics standard, summarised below. The constants used— $m_N, \|c'_G\|_\infty, \|c''_G\|_\infty$, the eigenvalue-growth bound, and the analytic-envelope K_3^A —are explicit.

We separate two very different kinds of gap. Items (i)–(iii) below are *implementation* gaps: each is a known quantity that better arithmetic or a longer computation would close. Item (0) is not of that kind. It is a structural obstruction, it is not closed by any improvement in precision, and it is the reason this paper does not state a local-maximality theorem on H^2 .

Status relative to global maximality. The argument here does not establish global maximality of Gerver’s sofa. Prior to Baek’s 2024 resolution [5], the best upper bound was Kallus–Romik’s [4] computer-assisted bound $\mu^* \leq 2.37$ established via rigorous interval arithmetic over exact rational data. Local maximality of c_G was listed by Romik [3, §7, Open Problem 1] as open; by Baek’s global resolution, the Gerver half of that problem is now subsumed. The companion half of Romik’s Open Problem 1, the local maximality of the ambidextrous shape Σ , is addressed by no global result, and—as recorded in Remark 53—is *not* settled here either. We state this without hedging: the present paper does not establish local maximality of Σ , and we are aware of no result that does. What the method contributes for Σ is the finite-dimensional negative-definiteness of Theorem 52 together with an explicit account of the remaining obstruction.

Identified gaps for full rigor. Four items, identified throughout the manuscript, separate the present argument from a strict computer-assisted proof. None of them is now structural: the two-norm difficulty, which an earlier draft of this work regarded as the principal obstruction, is resolved by Lemma 23 and Proposition 11.

- (0) **The moving-breakpoint cubic (the operative obstruction).** The Gårding constant is settled: $\|C\|_{\text{op}} = 2$ exactly, giving (107). The cubic side is not, and by Theorem 25 it cannot be: the coefficient of $\eta'(b_j)^3$ is nonzero, so the improved-Taylor hypothesis fails and no two-norm argument in the pair (H^1, H^2) can close. Overcoming this requires a genuinely different device — a norm adapted to the breakpoint jets, a degenerate-coercivity framework, or a reformulation in which the breakpoints do not move (for instance by majorizing F with the breakpoints frozen, in the spirit of Baek’s concave majorant). We regard this as the principal open problem raised by the present work, and note that it is now a concrete question about four scalars rather than a diffuse analytic difficulty. An earlier version of this manuscript claimed these coefficients vanish, on the strength of Lemma 23; that argument was invalid, since affineness constrains the integrand and not the moving limits, and the numerical verdict is the opposite.
- (1) **Interval polygon-intersection oracle.** The finite-dimensional Hessian is currently evaluated with SHAPELY, a floating-point polygon-intersection library. Replacing this with an interval-arithmetic equivalent (e.g., CGAL with exact predicates plus arb-interval coordinate propagation) is identified as the principal computational item; this would certify the matrix entries of Q_N to a width narrow enough to apply the existing arb interval-eigenvalue machinery rigorously. As a concrete symptom of the present non-rigorous status, individual diagonal entries $Q[k, k]$ are reproducible only to a few percent across finite-difference step sizes and θ -grid resolutions (e.g. the $k = 1$ entry varies between -4.56 and -4.88 , a $\approx 7\%$ spread, across the configurations in `rigorous/phase2a_lemma8_check.json` and `rigorous/hypV_sweep.json`); we therefore report m_N to three significant figures and treat the negative-definiteness as strong evidence rather than a certified enclosure. The matrix is also far from diagonally dominant (off-diagonal Gershgorin row sums ~ 45 against diagonal ~ -5), so the sign of the smallest eigenvalue is sensitive to entrywise errors and an interval oracle is genuinely necessary to make the conclusion rigorous.

A first step toward this is now in the repository: for the *analytic* trajectories (Gerver and its localized perturbations) the area admits an exact Green’s-theorem boundary integral over the contact paths, evaluated in MPMATH (`rigorous/analytic_oracle.py`); it reproduces A^* to 4×10^{-16} versus Shapely’s 2×10^{-7} , and made the quartic-criterion coefficient C_4 of Section 8.3 measurable where the polygon oracle could not. Promoting MPMATH to ARB intervals converts this into a genuinely certified evaluator on the analytic locus. The general 16-mode Hessian still uses SHAPELY, since a generic perturbation of the Fourier basis is not one of Romik’s five closed-form arcs; extending the analytic oracle to that basis is the remaining engineering step.

- (2) **Analytic cross-term bound.** The cross-block $\|Q_{NT}\|_{H^2 \rightarrow H^2}$ between low- and high-frequency Fourier modes was numerically probed in `rigorous/phase1a_cross_term.py` and conservatively bounded by 0.05; the empirical Frobenius extrapolation is ≈ 0.022 . A rigorous analytic bound exploiting the parity selection rule observed in the probe is expected to bring the constant well below 0.05 and is identified as a short follow-up task.
- (3) **Analytic Q_{jump} derivation.** Lemma 8 characterizes the smooth per-arc piece of the second variation; the jump piece Q_{jump} at the four contact-transition breakpoints is verified empirically (the sum rule $Q_{\text{smooth}}[k, k]_x + Q_{\text{smooth}}[k, k]_y = -\pi k^2$ holds analytically, while the full $Q[k, k]_x + Q[k, k]_y$ measured by the Hessian computation in Phase 2 exceeds $-\pi k^2$ in magnitude by a factor of ≈ 2.7 , attributed to Q_{jump}). An explicit closed-form derivation

of Q_{jump} in arb interval arithmetic, paralleling the analytic envelope for K_3 , is identified as the principal analytic item. This requires explicit case analysis of the contact-side transitions in Romik’s 5-phase decomposition. A diagnostic reported in Section 7.3 narrows the target: the diagonal excess attributable to Q_{jump} grows like a mode-uniform k^2 , which any closed form must reproduce. We caution that this mode-uniformity does *not* settle the rank of Q_{jump} : the diagonal of a finite-rank breakpoint-jet operator is itself dominated by a mode-uniform k^2 “DC” term, so the diagonal cannot distinguish a genuine finite-rank operator from an effective $\|\eta'\|^2$ contribution. Resolving the rank requires the full off-diagonal operator together with an analytically assembled Q_{smooth} .

The framework generalizes naturally to other intersection-of-rotated-hallway problems: variants with non-right corner angles, ambidextrous hallways, and three-dimensional analogues all admit the same Phase 1–4 structure. Each requires its own Phase 1 (analytic-constant solution) and its own coercivity certification, but the framework is otherwise problem-independent.

Beyond closing the three gaps above, further follow-up directions include:

- **Tighten the off-diagonal eigenvalue bound** beyond the $O(\log N)$ -style estimate used in Theorem 18. The empirical asymptotic growth $|\lambda_k| \sim k^{1.6}$ indicates the diagonal bound is asymptotically loose by a sub-logarithmic factor; closing this would yield a sharper explicit coercivity constant.
- Formalize the implication chain in a proof assistant (Lean 4 or Coq) to machine-verify the explicit constants and the Taylor remainder analysis.
- Apply the framework to Romik’s ambidextrous variant and to hallways with non-right corner angles, where no analytic candidate currently exists; the same Phase 1–4 structure transfers with problem-specific Phase 1 setups.

REPRODUCIBILITY AND COMPUTER-ASSISTED COMPUTATIONS

All numerical computations are open-source and available in the supplementary repository. Key files:

- `rigorous/gerver_constants.py`: Phase 1 multi-precision Newton iteration.
- `rigorous/gerver_arb.py`: arb ball-arithmetic certification of Phase 1.
- `rigorous/second_variation.py`: Phase 2 Hessian computation.
- `rigorous/phase3_robustness.py`: Phase 3 mode-count scan and symmetry quotient.
- `rigorous/phase4_full_theorem.py`: Phase 4 eigenvalue enclosures of the computed Hessian.
- `rigorous/rigor_validation.py`: numerical validation of all explicit constants in the rigorous bounds.
- `rigorous/sympy_lemma10_2.py`: SymPy verification of the structural identity in Lemma 8.
- `rigorous/sympy_constants_extract.py`: SymPy extraction of the explicit $|A|, |C|, |D|$ bounds yielding Theorem 15.
- `rigorous/hypV_sweep.py`: numerical verification of Theorem 15 across $k \in \{1, \dots, 32\}$.
- `rigorous/phase2b_finite_rank_test.py`, `rigorous/phase2c_oscillation_test.py`: diagnostics for the structure of Q_{jump} reported in Section 7.3 (mode-uniform vs. breakpoint-frequency excess).
- `rigorous/RIGOROUS_BOUNDS.md`: the full derivation of Theorem 15.
- `rigorous/ray_graph_cert.py`: ball-arithmetic certification of the geometric side condition of Theorem 36 (simplicity of the reconstruction curve) for frozen reconstructions along deformation rays: speed-sign runs with chord replacement, monotone-chain injectivity, chord–chord and arc–segment separators certified through the exact quadratic ε -dependence, and a winding-number enclosure fixing orientation. Two structural facts found en route: a reversal pocket of a contact arc *retraces* the envelope (so chord cuts must span positive

net displacement before the flanking arcs are disjoint as sets), and for non-simple frozen curves the correct side condition is the winding form of Lemma 32 ($\text{wind}_\Gamma \geq 1$ on S , ≥ 0 elsewhere), under which exact retraces cancel in the Green functional.

- **lean/MovingSofa**: machine-verified (Lean 4, no **sorry**) cores of the superset principle, the exact-degree reduction, the stationary-contact mechanism and cap law, the ambidextrous frame-pair coercivity, and the monotone-chain injectivity skeleton.

End-to-end reproduction takes approximately 30 minutes on a 2024 laptop without specialized hardware.

USE OF ARTIFICIAL INTELLIGENCE

This work was produced with substantial assistance from several AI systems, used at different stages to draft and develop the mathematics, write the supporting code, cross-review and critique the arguments, and iterate on the manuscript. The models used include Anthropic’s Claude (versions 4.6, 4.7, and 4.8), Google’s Gemini 3.1 Pro, and Meta AI. In particular, the assistants contributed to the mathematical derivations; wrote essentially all of the supporting code (the Newton solver for Gerver’s constants, the Hessian and coercivity computations, the symbolic verification of Lemma 8, and the numerical sweeps); generated successive rounds of adversarial peer review that shaped the present, deliberately conservative, statement of results; and produced initial drafts of the text. The author set the research direction, ran and checked the computations, verified the arguments, and is solely responsible for the final text and for any errors. The three items flagged as open throughout (Section 11) are stated as open precisely because they have not been verified to a standard the author is willing to certify.

REFERENCES

- [1] J. L. Gerver, *On moving a sofa around a corner*, *Geometriae Dedicata* **42** (1992), 267–283.
- [2] J. M. Hammersley, *On the enfeeblement of mathematical skills by “Modern Mathematics” and by similar soft intellectual trash in schools and universities*, *Bulletin of the Institute of Mathematics and its Applications* **4** (1968), 66–85. Hammersley’s sofa bound $\pi/2 + 2/\pi$ appears in §13 of that article.
- [3] D. Romik, *Differential equations and exact solutions in the moving sofa problem*, *Experimental Mathematics* **27** (2018), 316–330. arXiv:1606.08111.
- [4] Y. Kallus and D. Romik, *Improved upper bounds in the moving sofa problem*, *Advances in Mathematics* **340** (2018), 960–982. arXiv:1706.06630.
- [5] J. Baek, *Optimality of Gerver’s sofa*, 2024. arXiv:2411.19826.
- [6] P. Gibbs, *A computational study of sofas and cars*, viXra preprint, 2014.
- [7] A. Henrot and M. Pierre, *Variation et optimisation de formes: une analyse géométrique*, *Mathématiques et Applications*, vol. 48, Springer, 2005.
- [8] J. Sokółowski and J.-P. Zolésio, *Introduction to Shape Optimization: Shape Sensitivity Analysis*, *Springer Series in Computational Mathematics*, vol. 16, Springer, 1992.
- [9] E. Zeidler, *Nonlinear Functional Analysis and its Applications, I: Fixed-Point Theorems*, Springer, 1986.
- [10] A. D. Ioffe, *Necessary and sufficient conditions for a local minimum. 3: Second order conditions and augmented duality*, *SIAM J. Control Optim.* **17** (1979), no. 2, 266–288.
- [11] M. Dambrine and J. Lamboley, *Stability in shape optimization with second variation*, *J. Differential Equations* **267** (2019), no. 5, 3009–3045.
- [12] M. R. Hestenes, *Applications of the theory of quadratic forms in Hilbert space to the calculus of variations*, *Pacific J. Math.* **1** (1951), 525–581.
- [13] J. M. Ball and J. E. Marsden, *Quasiconvexity at the boundary, positivity of the second variation and elastic stability*, *Arch. Rational Mech. Anal.* **86** (1984), 251–277.
- [14] E. M. Stein, *Harmonic Analysis: Real-Variable Methods, Orthogonality, and Oscillatory Integrals*, Princeton University Press, 1993.
- [15] G. H. Hardy, J. E. Littlewood, G. Pólya, *Inequalities*, Cambridge University Press, 2nd ed., 1952.
- [16] R. A. Adams, J. J. F. Fournier, *Sobolev Spaces*, *Pure and Applied Mathematics*, vol. 140, Academic Press, 2nd ed., 2003.
- [17] T. C. Hales, *A proof of the Kepler conjecture*, *Annals of Mathematics* **162** (2005), 1065–1185.
- [18] N. Lord et al., *Computer-assisted proofs in fluid mechanics: a survey*, in preparation.

- [19] R. Castelli, M. Gameiro, J.-P. Lessard, *Rigorous numerics for ill-posed PDEs: periodic orbits in the Boussinesq equation*, Arch. Rat. Mech. Anal. **228** (2018), 129–157.
- [20] F. Johansson et al., *mpmath: a Python library for arbitrary-precision floating-point arithmetic*, version 1.3.0, <http://mpmath.org/>.
- [21] F. Johansson, *Arb: efficient arbitrary-precision midpoint-radius interval arithmetic*, IEEE Transactions on Computers **66** (2017), 1281–1292.
- [22] S. Gillies et al., *Shapely: Python package for manipulation of geometric objects*, version 2.0, <https://shapely.readthedocs.io/>.